

# Introduction to Algorithms

## 10. Balanced Search Trees

---

Hyungsoo Jung

# Balanced search trees

**Balanced search tree:** A search-tree data structure for which a height of  $O(\lg n)$  is guaranteed when implementing a dynamic set of  $n$  items.

**Examples:**

- AVL trees
- 2-3 trees
- 2-3-4 trees
- B-trees
- Red-black trees

# Binary Search Tree - Best Time

---

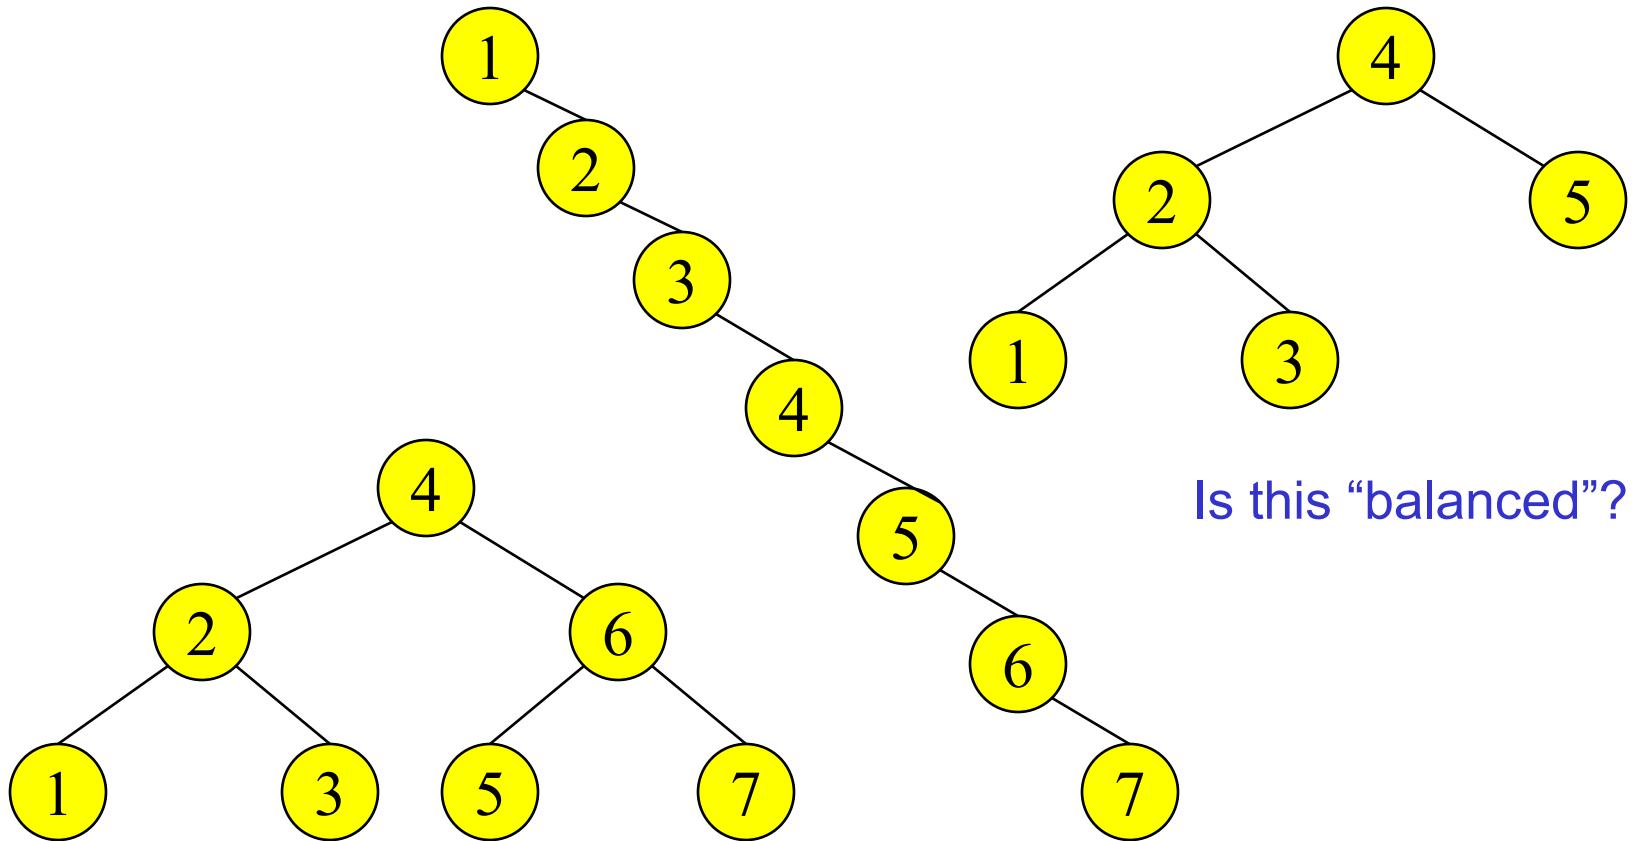
- All BST operations are  $O(d)$ , where  $d$  is tree depth
- minimum  $d$  is  $d = \lfloor \log_2 N \rfloor$  for a binary tree with  $N$  nodes
  - › What is the best case tree?
  - › What is the worst case tree?
- So, best case running time of BST operations is  $O(\log N)$

# Binary Search Tree - Worst Time

---

- Worst case running time is  $O(N)$ 
  - › What happens when you Insert elements in ascending order?
    - Insert: 2, 4, 6, 8, 10, 12 into an empty BST
  - › Problem: Lack of “balance”:
    - compare depths of left and right subtree
  - › Unbalanced degenerate tree

# Balanced and unbalanced BST



# Approaches to balancing trees

---

- Don't balance
  - › May end up with some nodes very deep
- Strict balance
  - › The tree must always be balanced perfectly
- Pretty good balance
  - › Only allow a little out of balance
- Adjust on access
  - › Self-adjusting

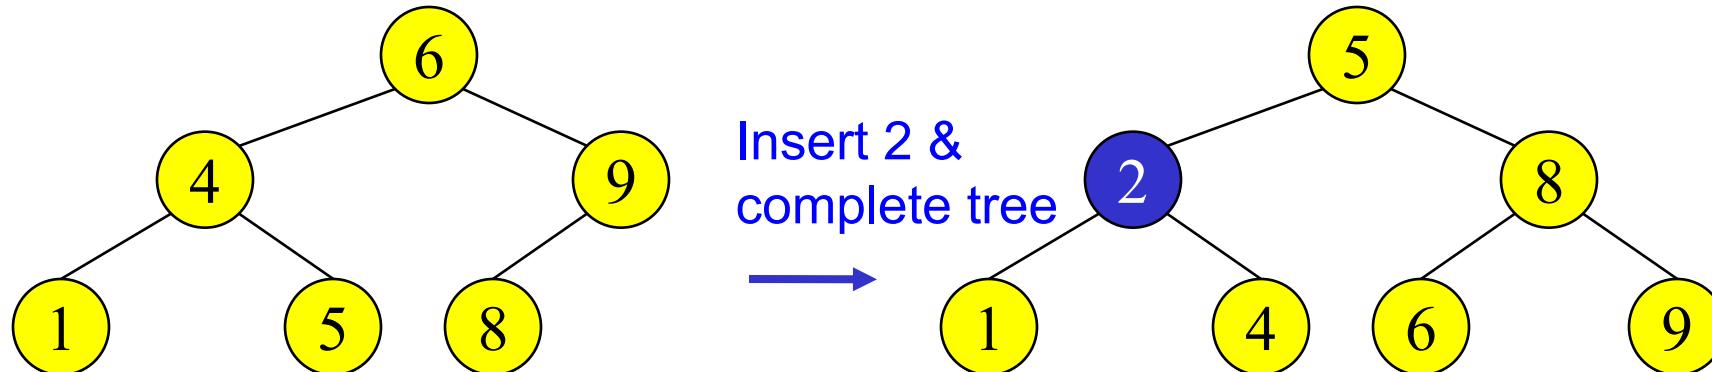
# Balancing Binary Search Trees

---

- Many algorithms exist for keeping binary search trees balanced
  - › Adelson-Velskii and Landis ([AVL](#)) trees (height-balanced trees)
  - › [Splay trees](#) and other self-adjusting trees
  - › [B-trees](#) and other multiway search trees

# Perfect Balance

- Want a **complete tree** after every operation
  - › tree is full except possibly in the lower right
- This is expensive
  - › For example, insert 2 in the tree on the left and then rebuild as a complete tree



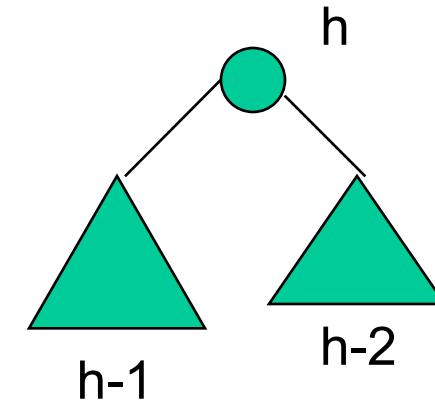
# AVL - Good but not Perfect Balance

---

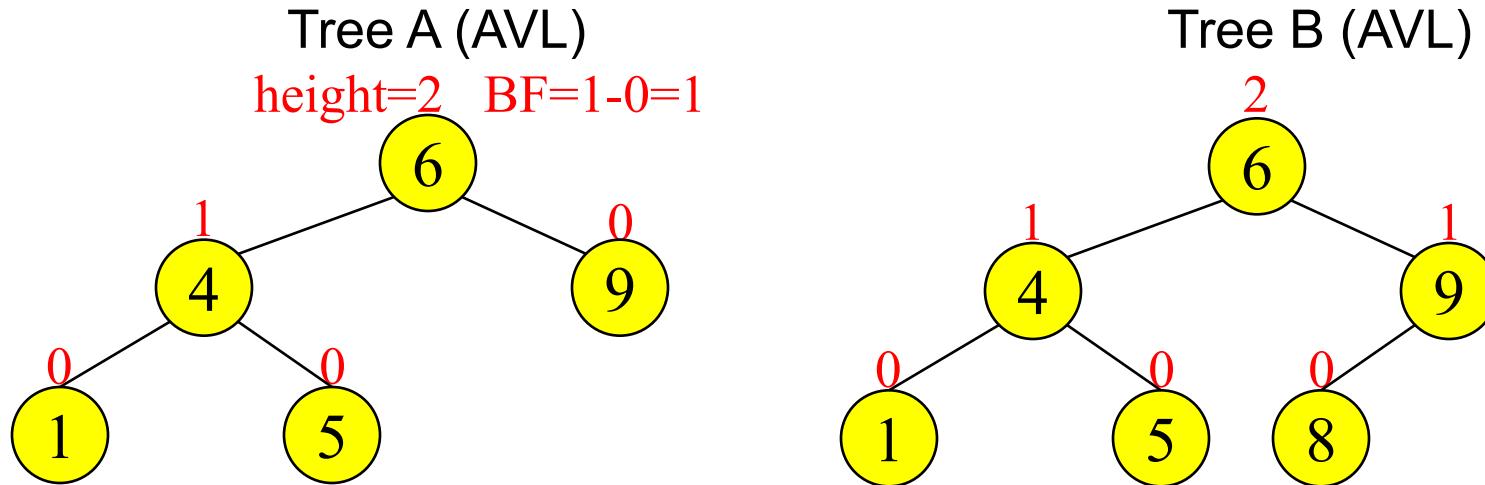
- AVL trees are height-balanced binary search trees
- Balance factor of a node
  - ›  $\text{height}(\text{left subtree}) - \text{height}(\text{right subtree})$
- An AVL tree has balance factor calculated at every node
  - › For every node, heights of left and right subtree can differ by no more than 1
  - › Store current heights in each node

# Height of an AVL Tree

- $N(h)$  = **minimum** number of nodes in an AVL tree of height  $h$ .
- **Basis**
  - ›  $N(0) = 1, N(1) = 2$
- **Induction**
  - ›  $N(h) = N(h-1) + N(h-2) + 1$
- **Solution** (recall Fibonacci analysis)
  - ›  $N(h) \geq \phi^h$  ( $\phi \approx 1.62$ )



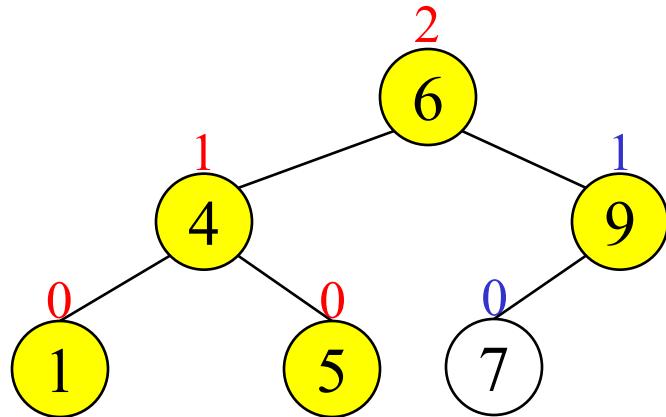
# Node Heights



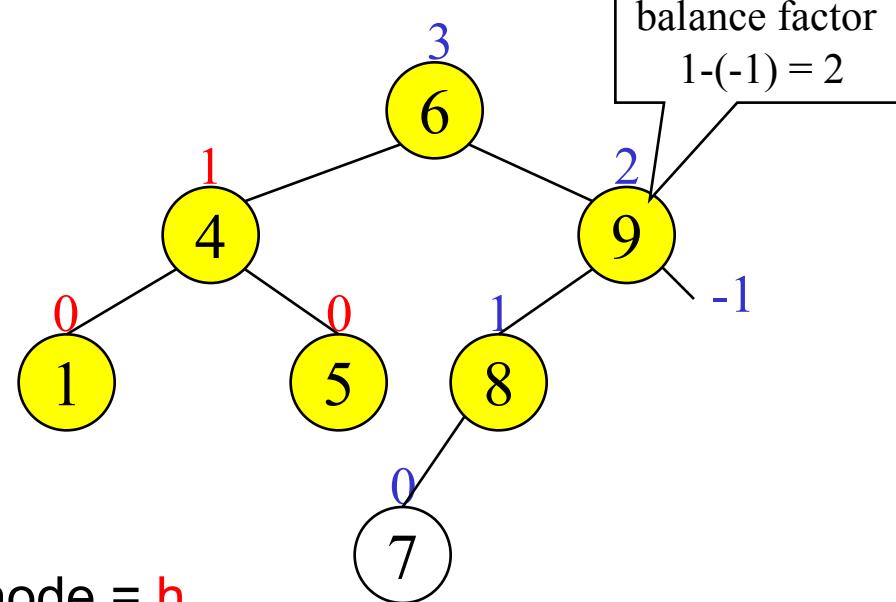
height of node =  $h$   
balance factor =  $h_{\text{left}} - h_{\text{right}}$   
empty height = -1

# Node Heights after Insert 7

Tree A (AVL)



Tree B (not AVL)



height of node =  $h$

balance factor =  $h_{\text{left}} - h_{\text{right}}$

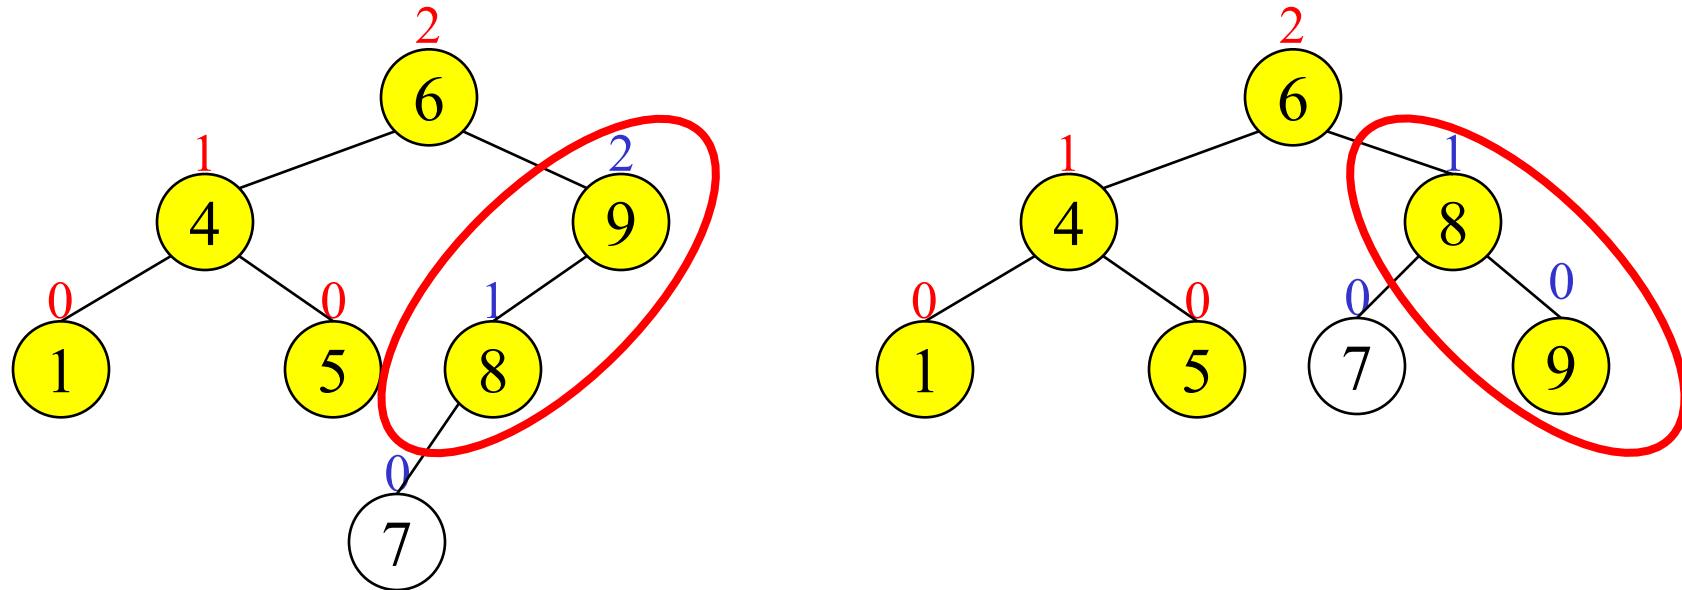
empty height = -1

# Insert and Rotation in AVL Trees

---

- Insert operation may cause balance factor to become 2 or -2 for some node
  - › only nodes on the path from insertion point to root node have possibly changed in height
  - › So after the Insert, go back up to the root node by node, updating heights
  - › If a new balance factor (the difference  $h_{\text{left}} - h_{\text{right}}$ ) is 2 or -2, adjust tree by *rotation* around the node

# Single Rotation in an AVL Tree



# Insertions in AVL Trees

---

Let the node that needs rebalancing be  $\alpha$ .

There are 4 cases:

Outside Cases (require single rotation) :

1. Insertion into **left subtree of left child of  $\alpha$** .
2. Insertion into **right subtree of right child of  $\alpha$** .

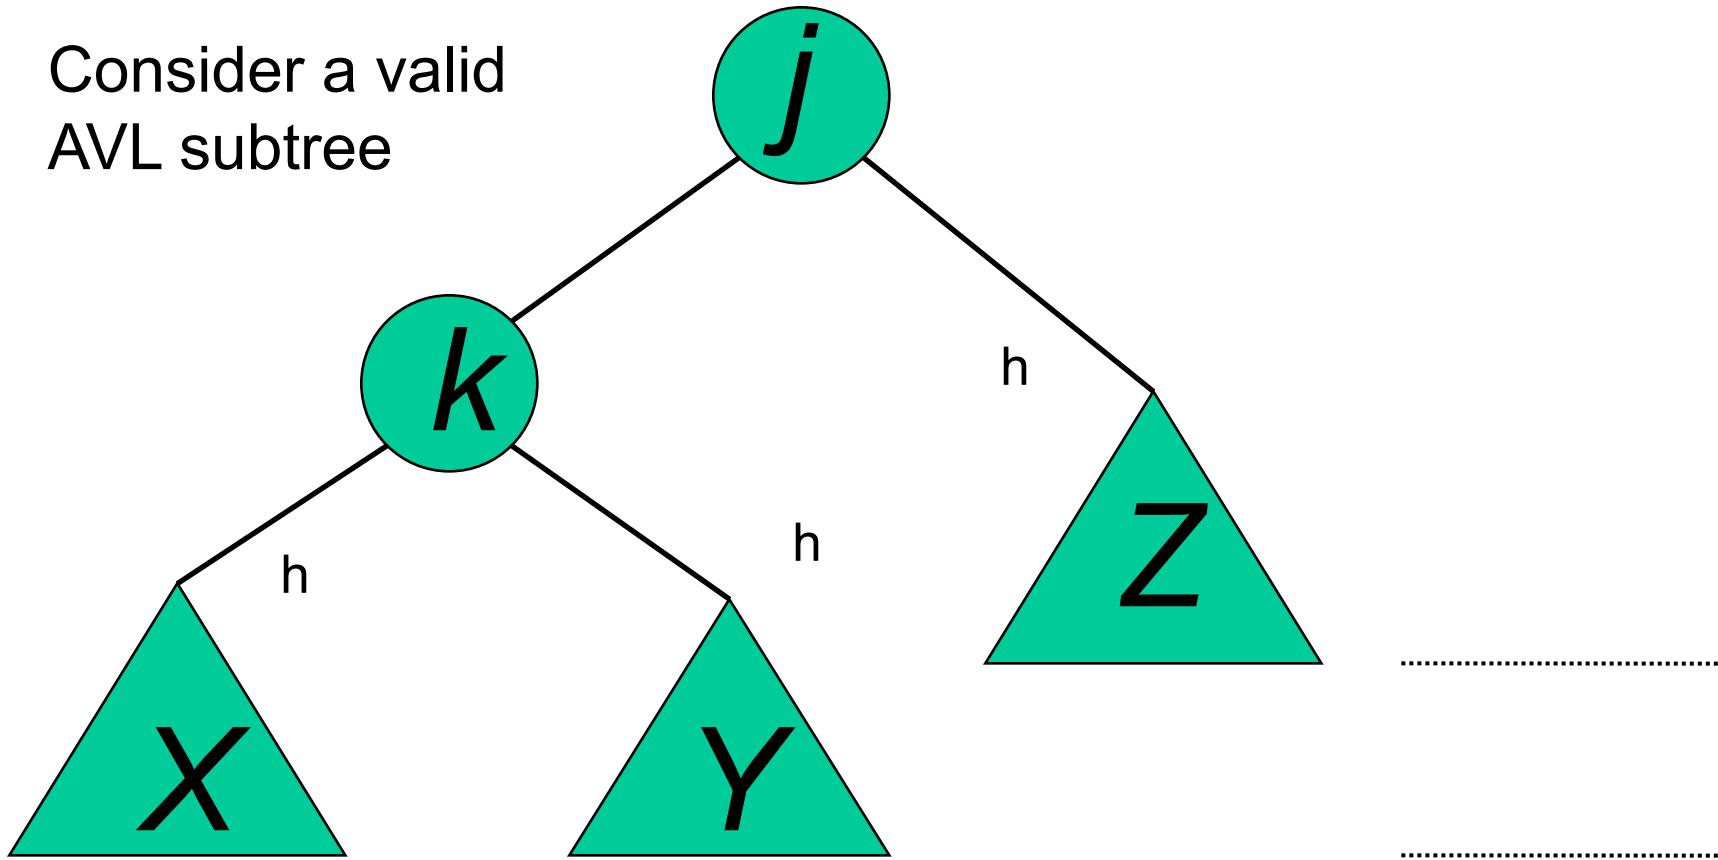
Inside Cases (require double rotation) :

3. Insertion into **right subtree of left child of  $\alpha$** .
4. Insertion into **left subtree of right child of  $\alpha$** .

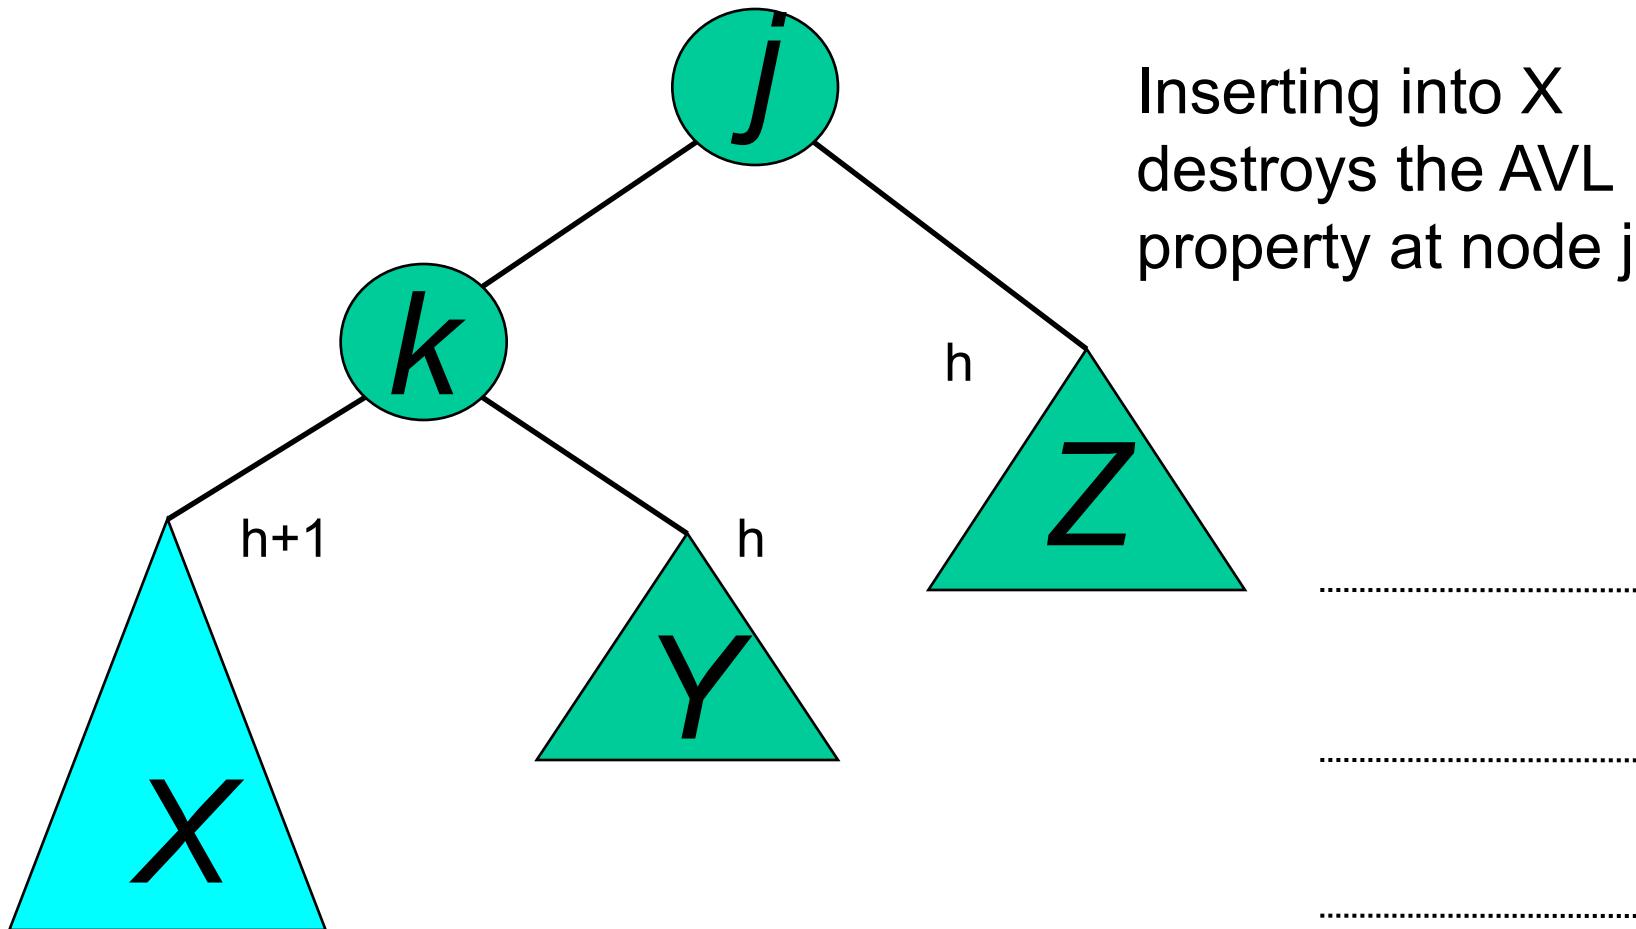
The rebalancing is performed through four separate rotation algorithms.

# AVL Insertion: Outside Case

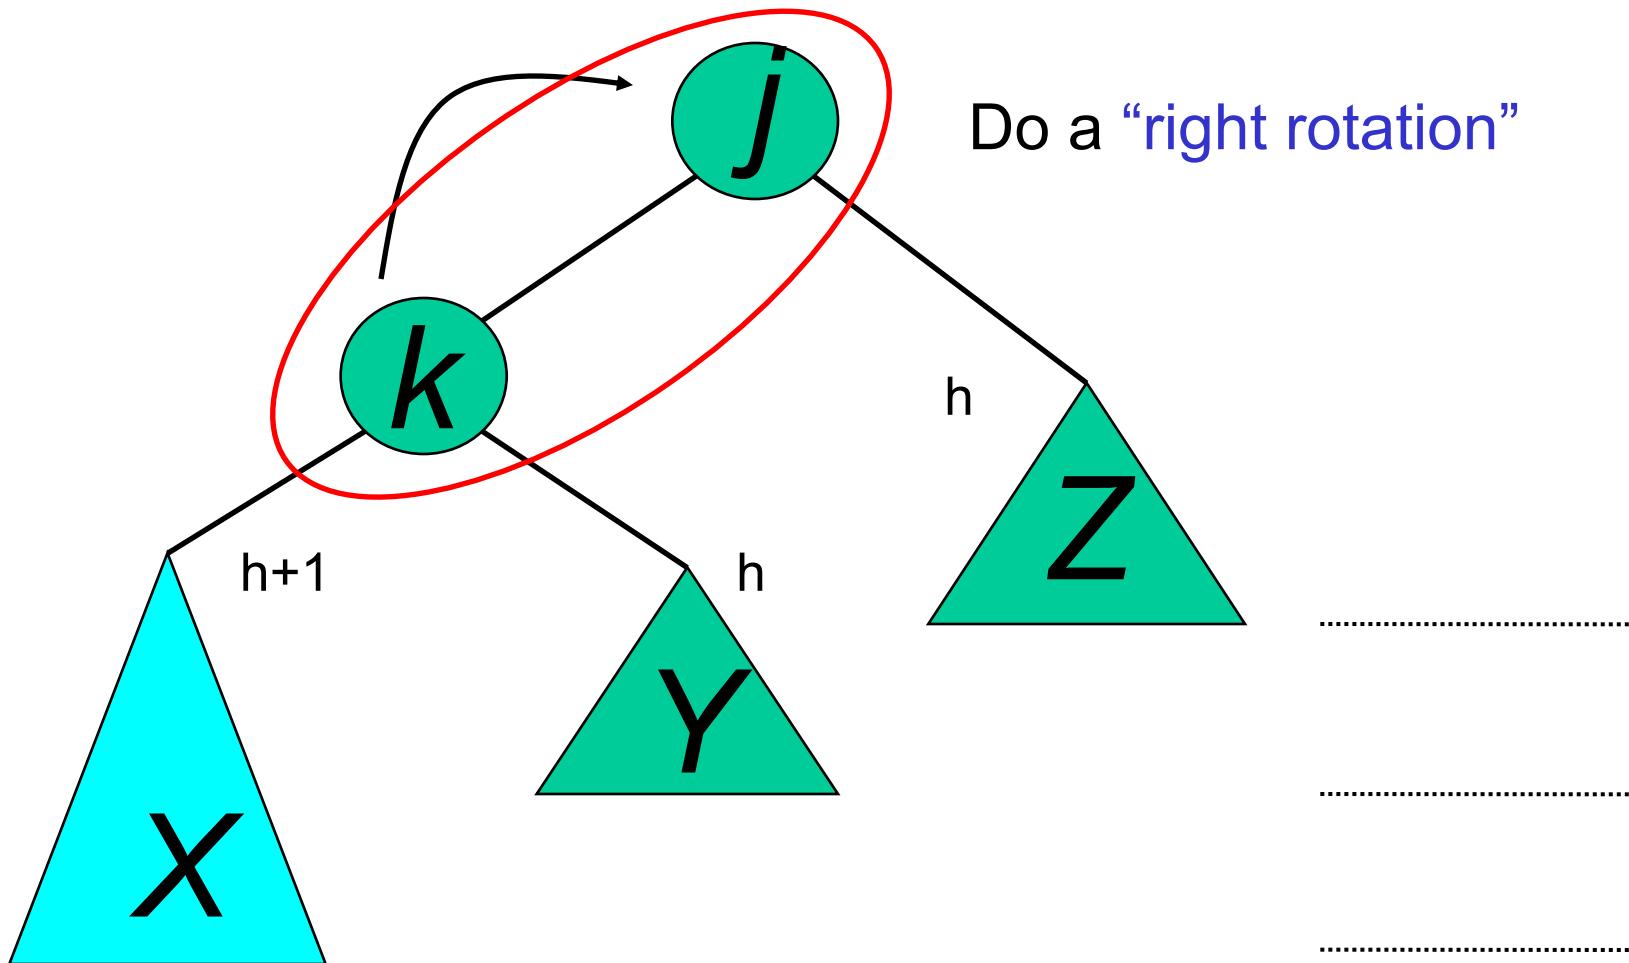
Consider a valid  
AVL subtree



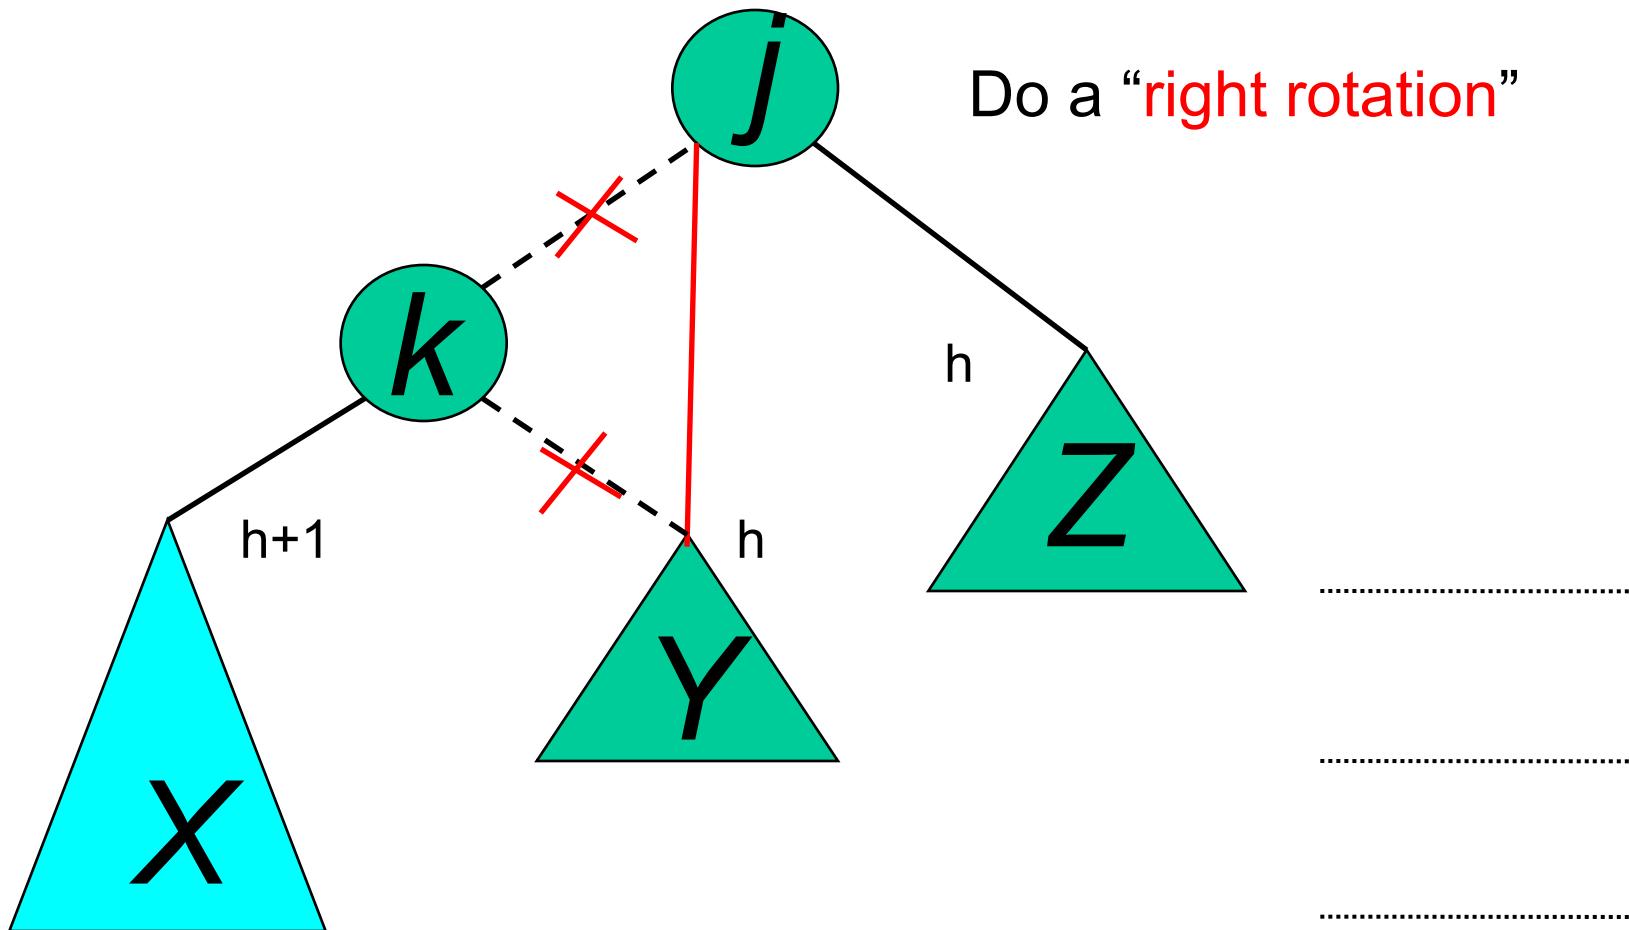
# AVL Insertion: Outside Case



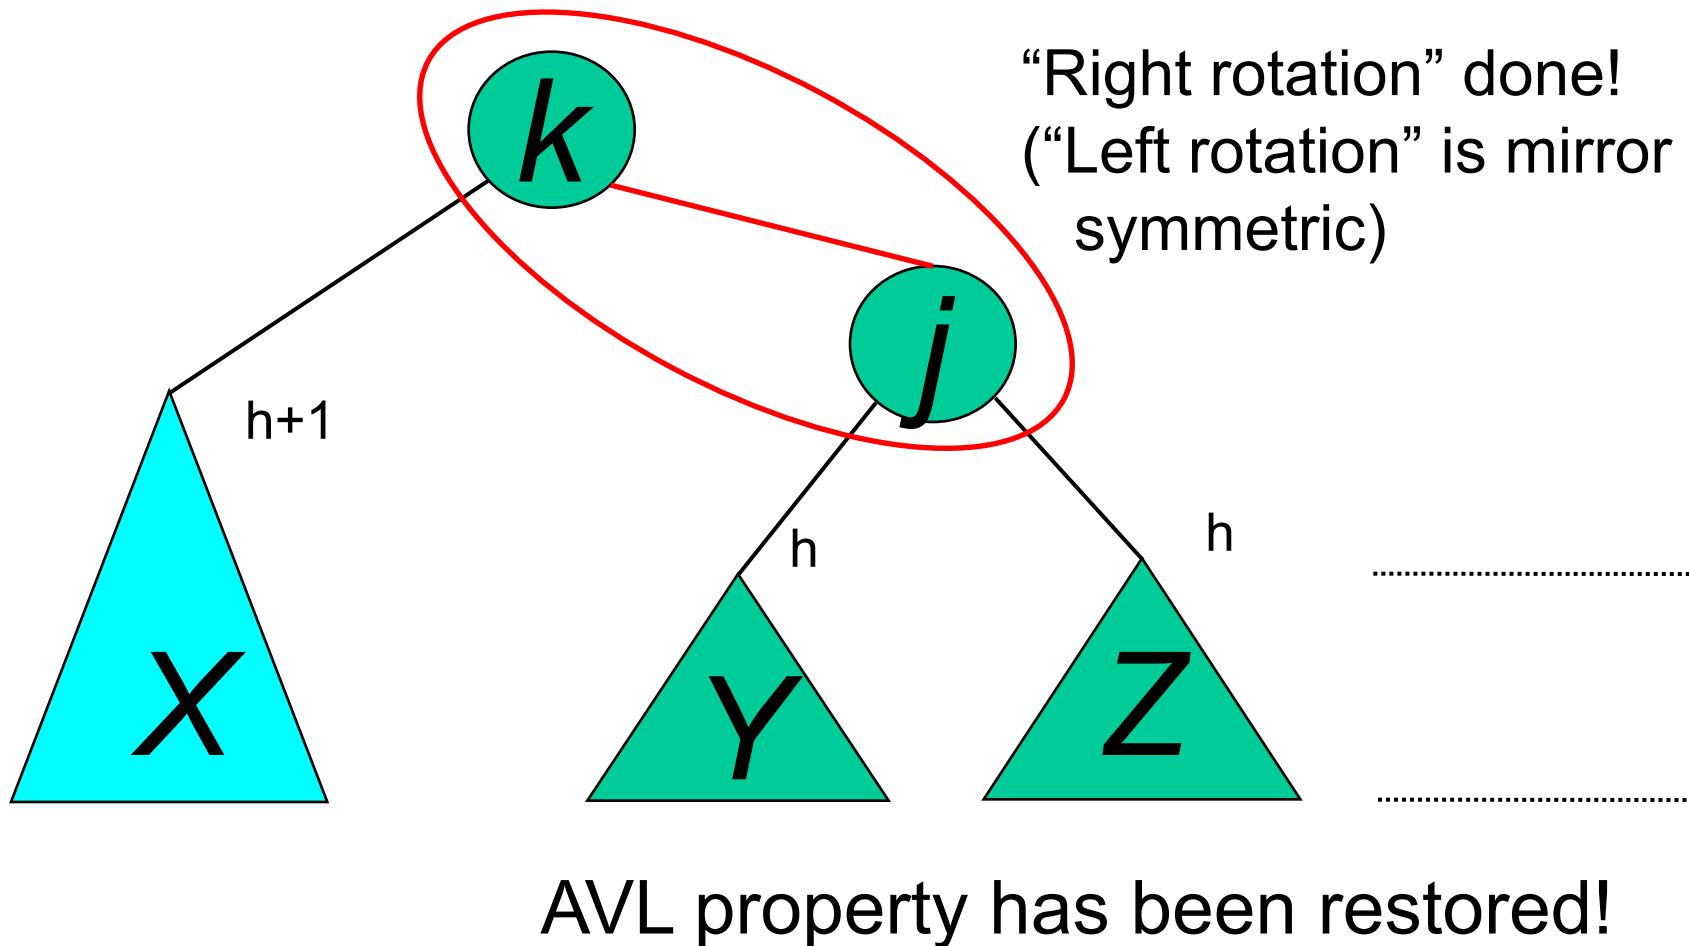
# AVL Insertion: Outside Case



# Single right rotation

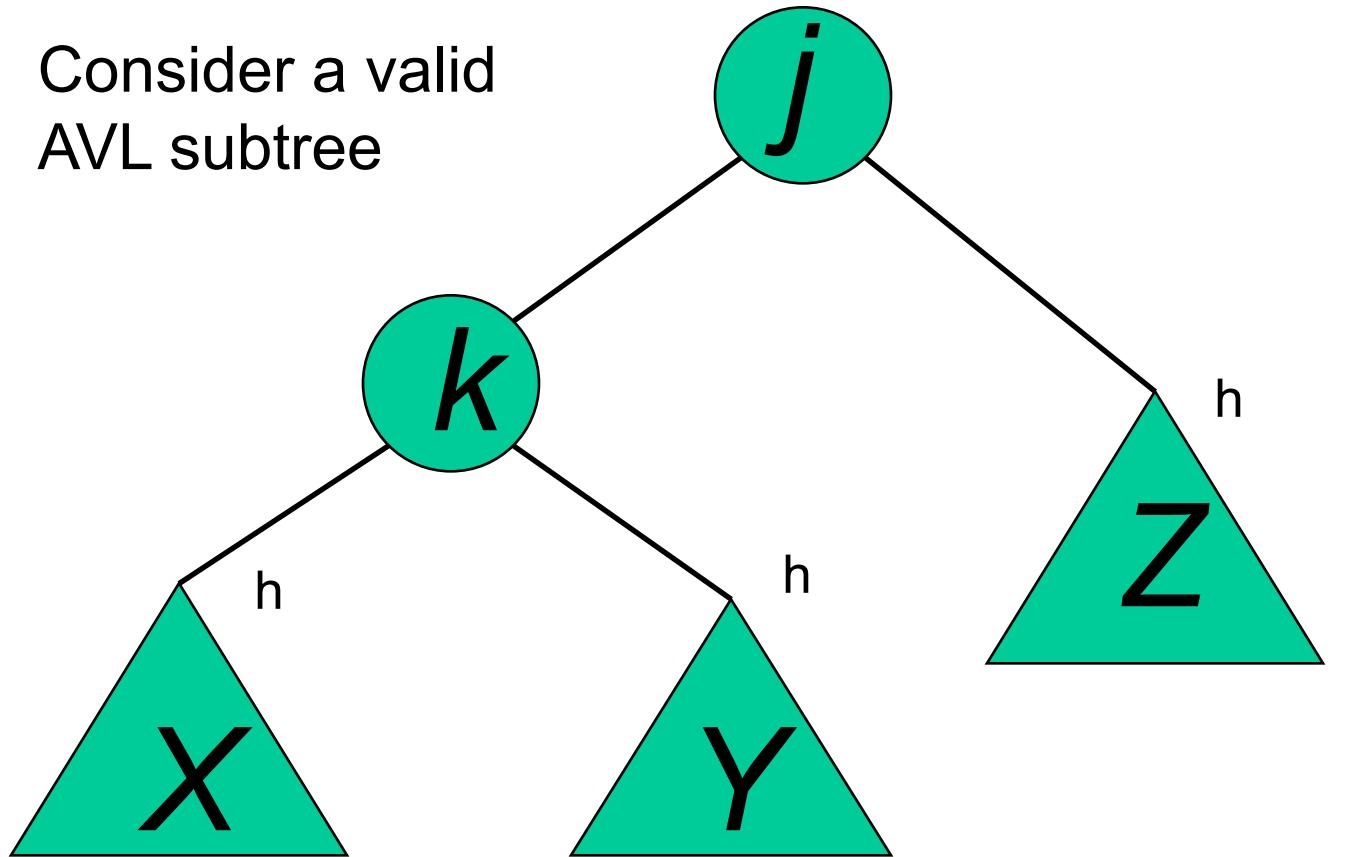


# Outside Case Completed



# AVL Insertion: Inside Case

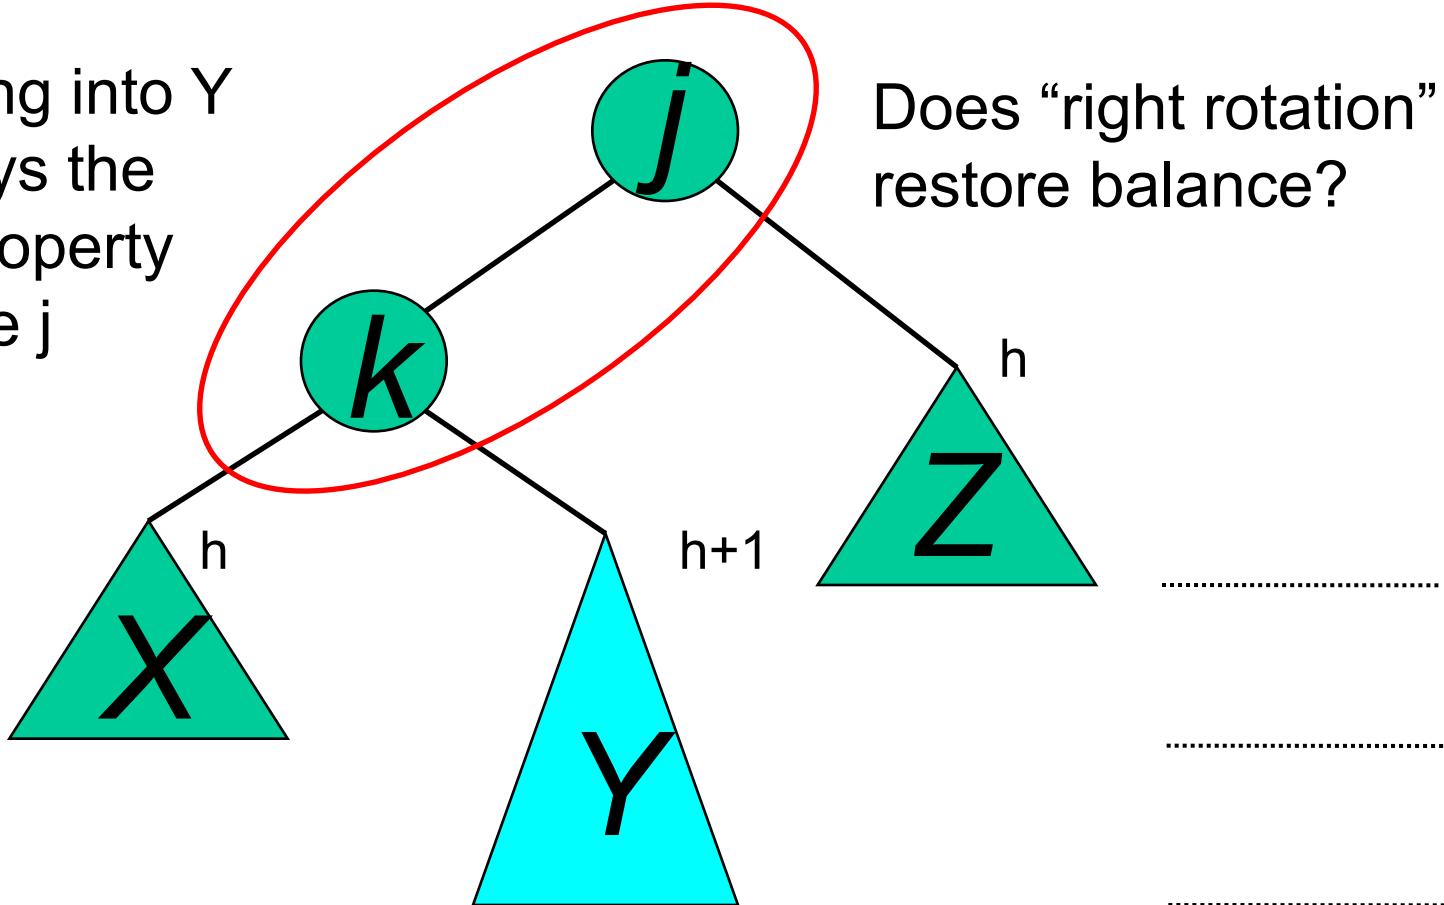
Consider a valid  
AVL subtree



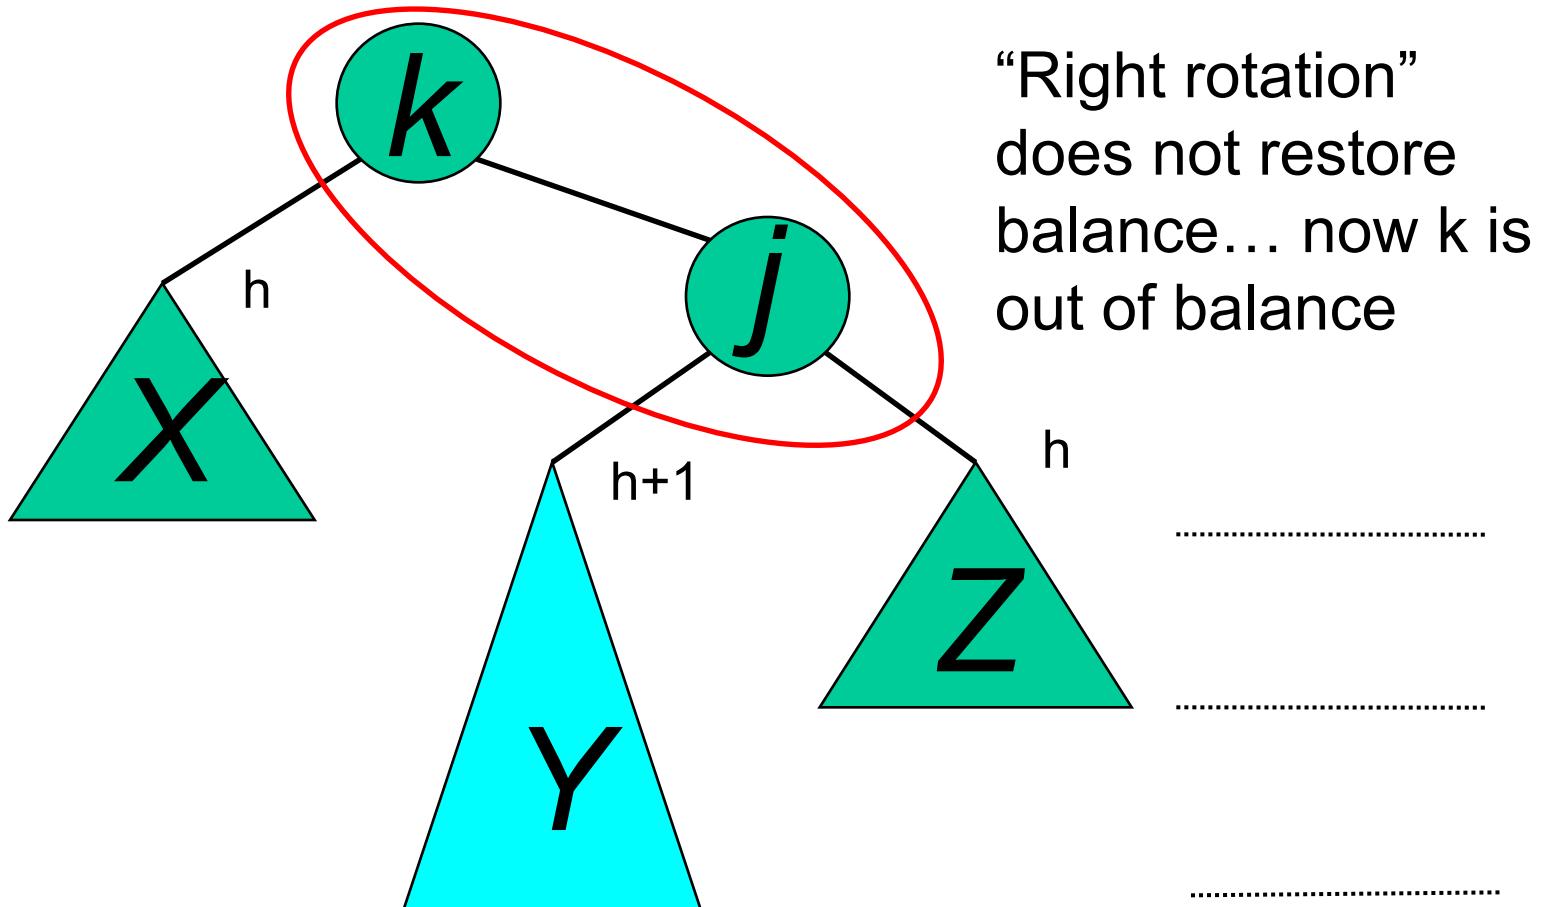
# AVL Insertion: Inside Case

Inserting into Y  
destroys the  
AVL property  
at node j

Does “right rotation”  
restore balance?

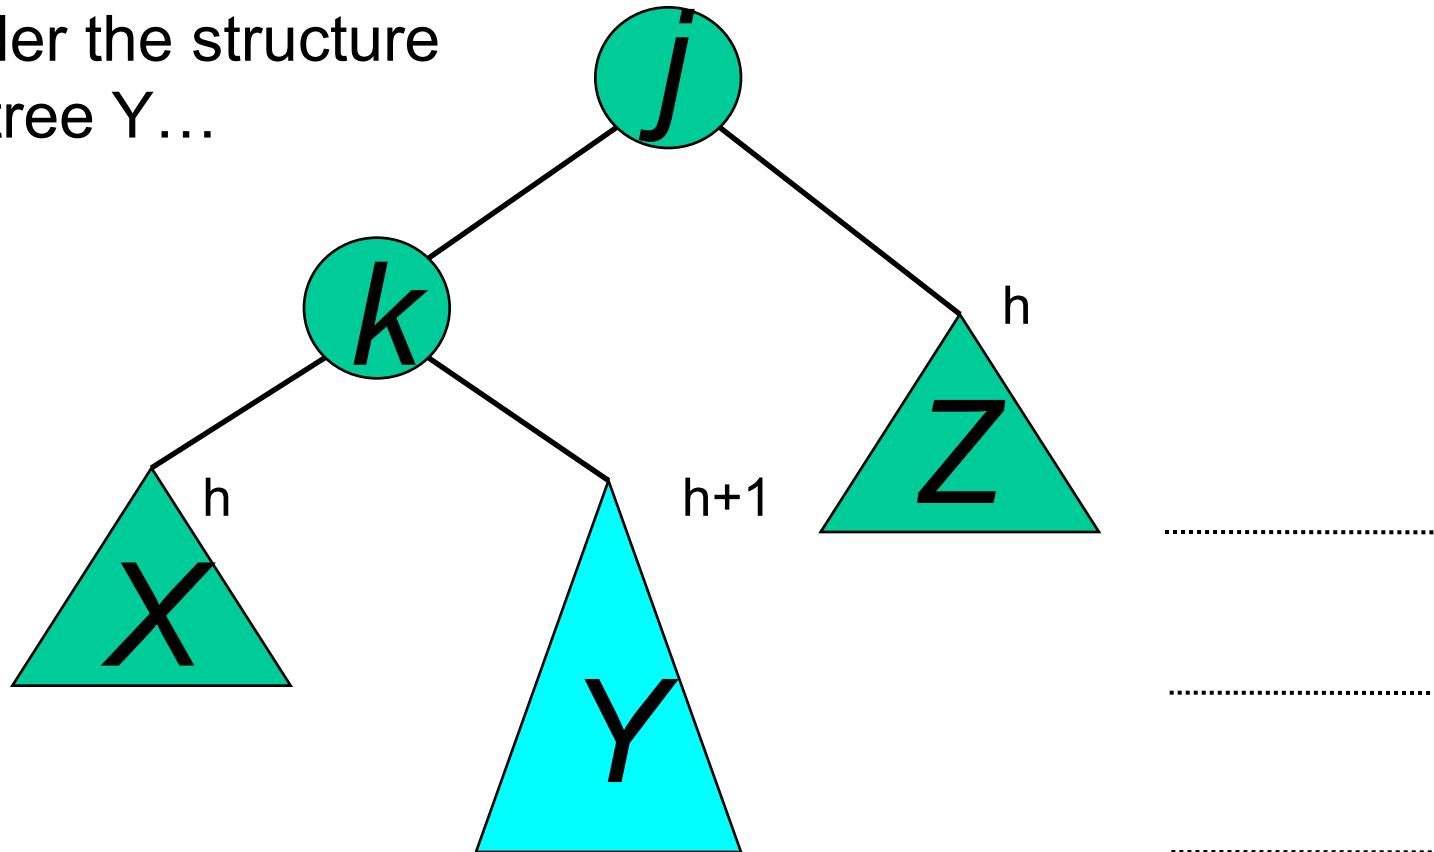


# AVL Insertion: Inside Case

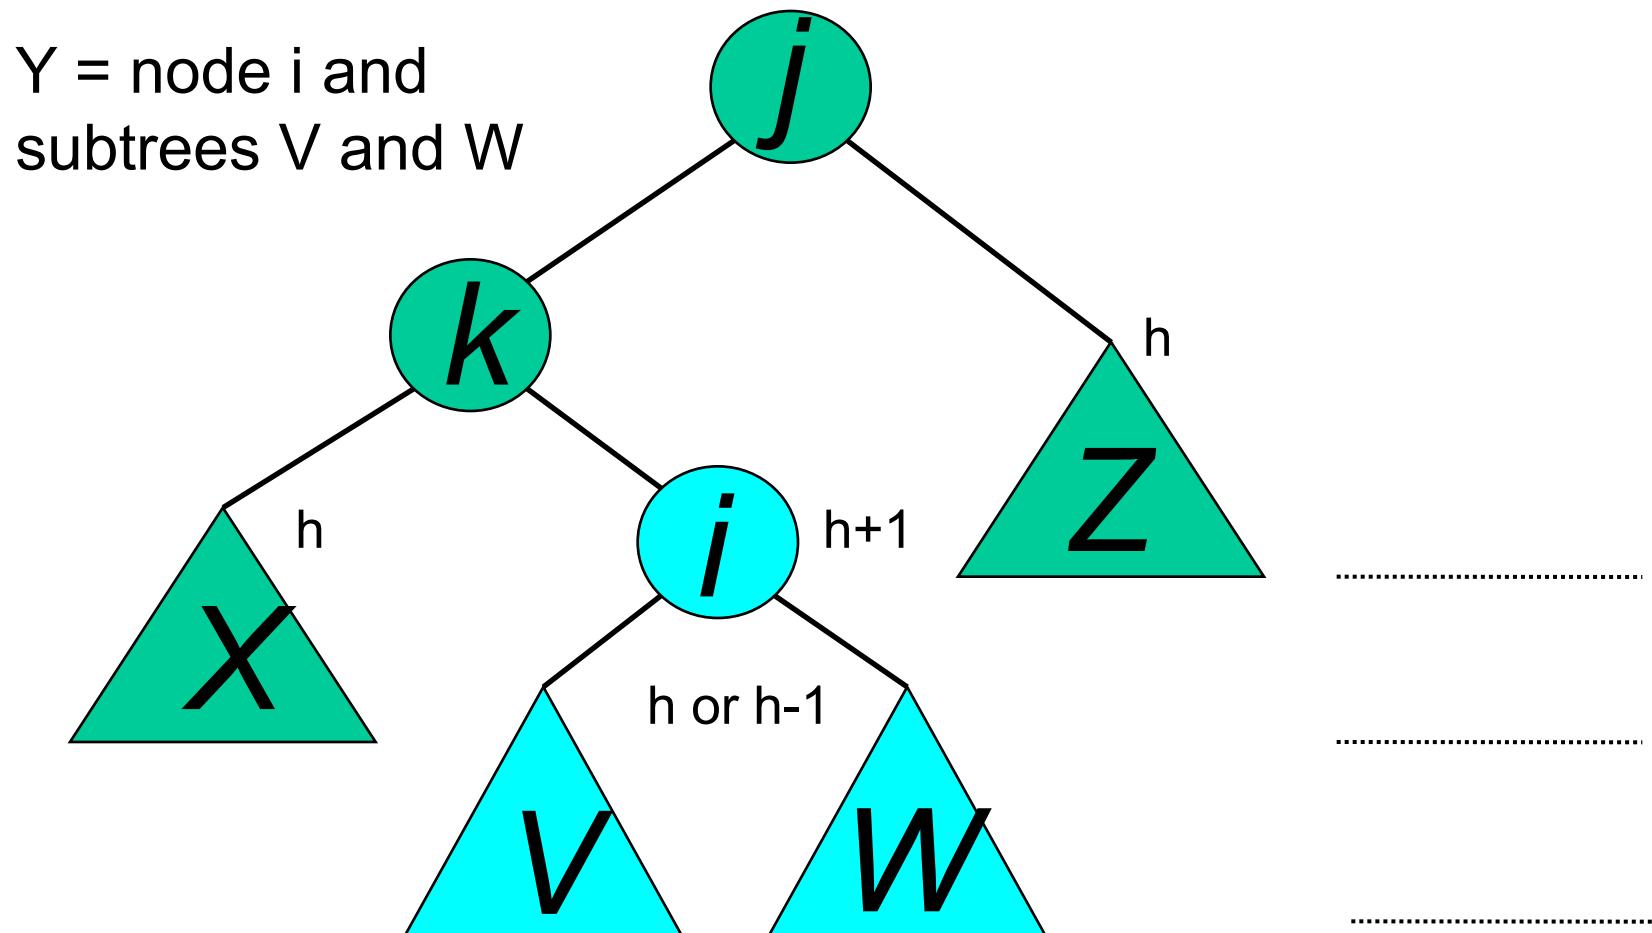


# AVL Insertion: Inside Case

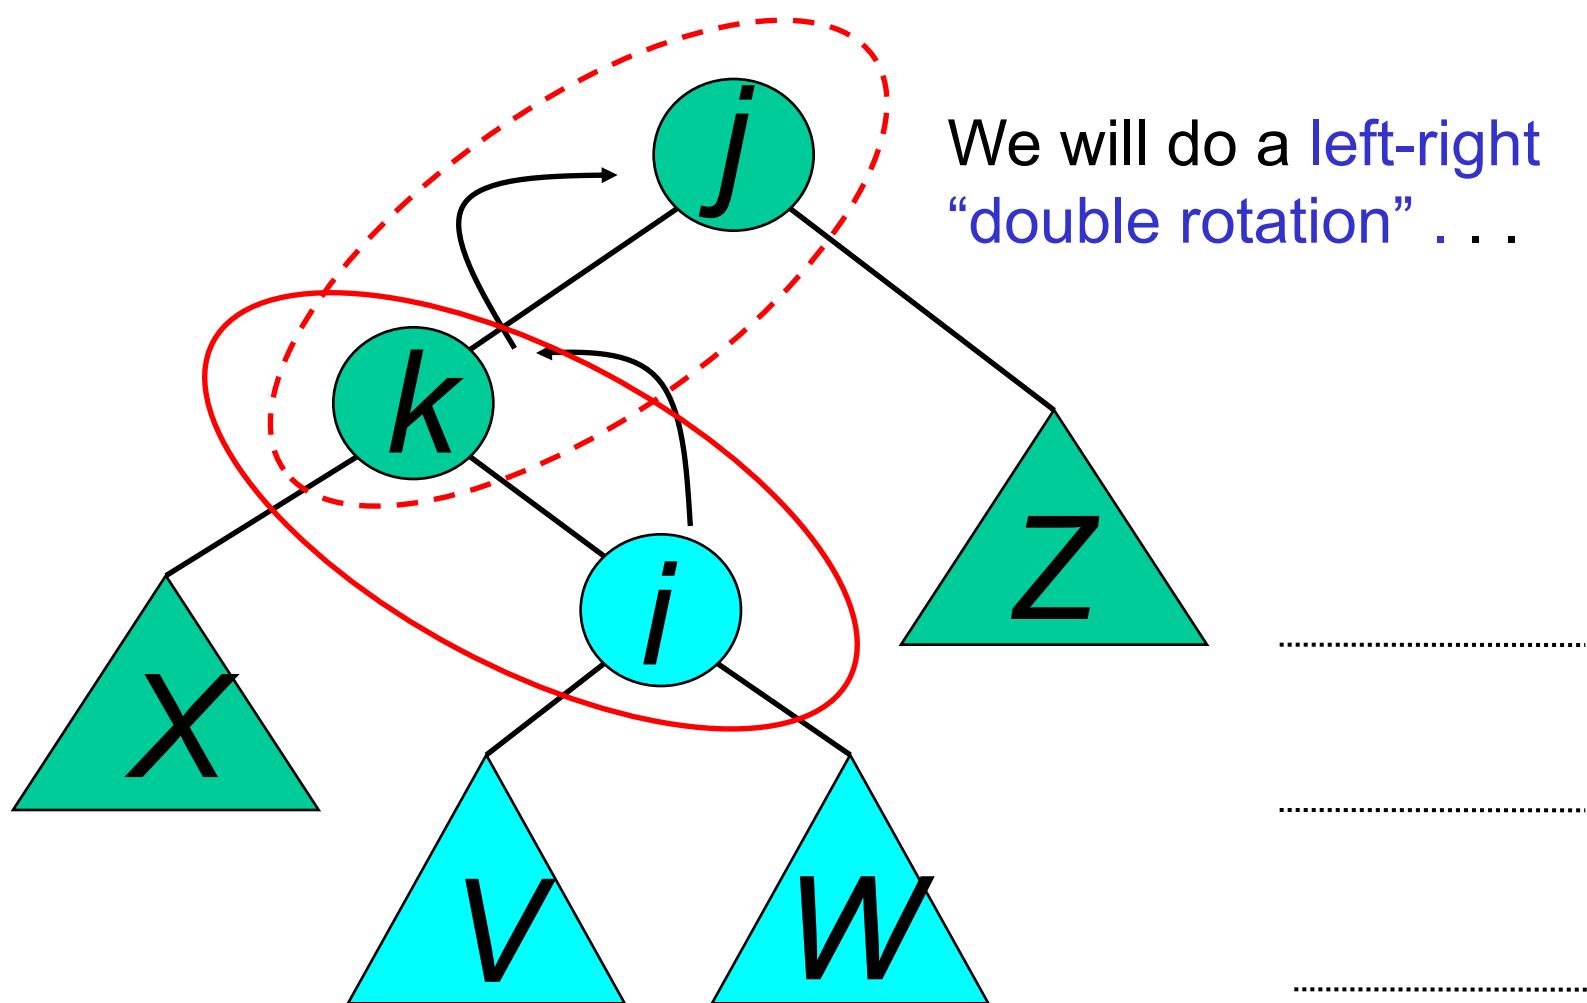
Consider the structure  
of subtree Y...



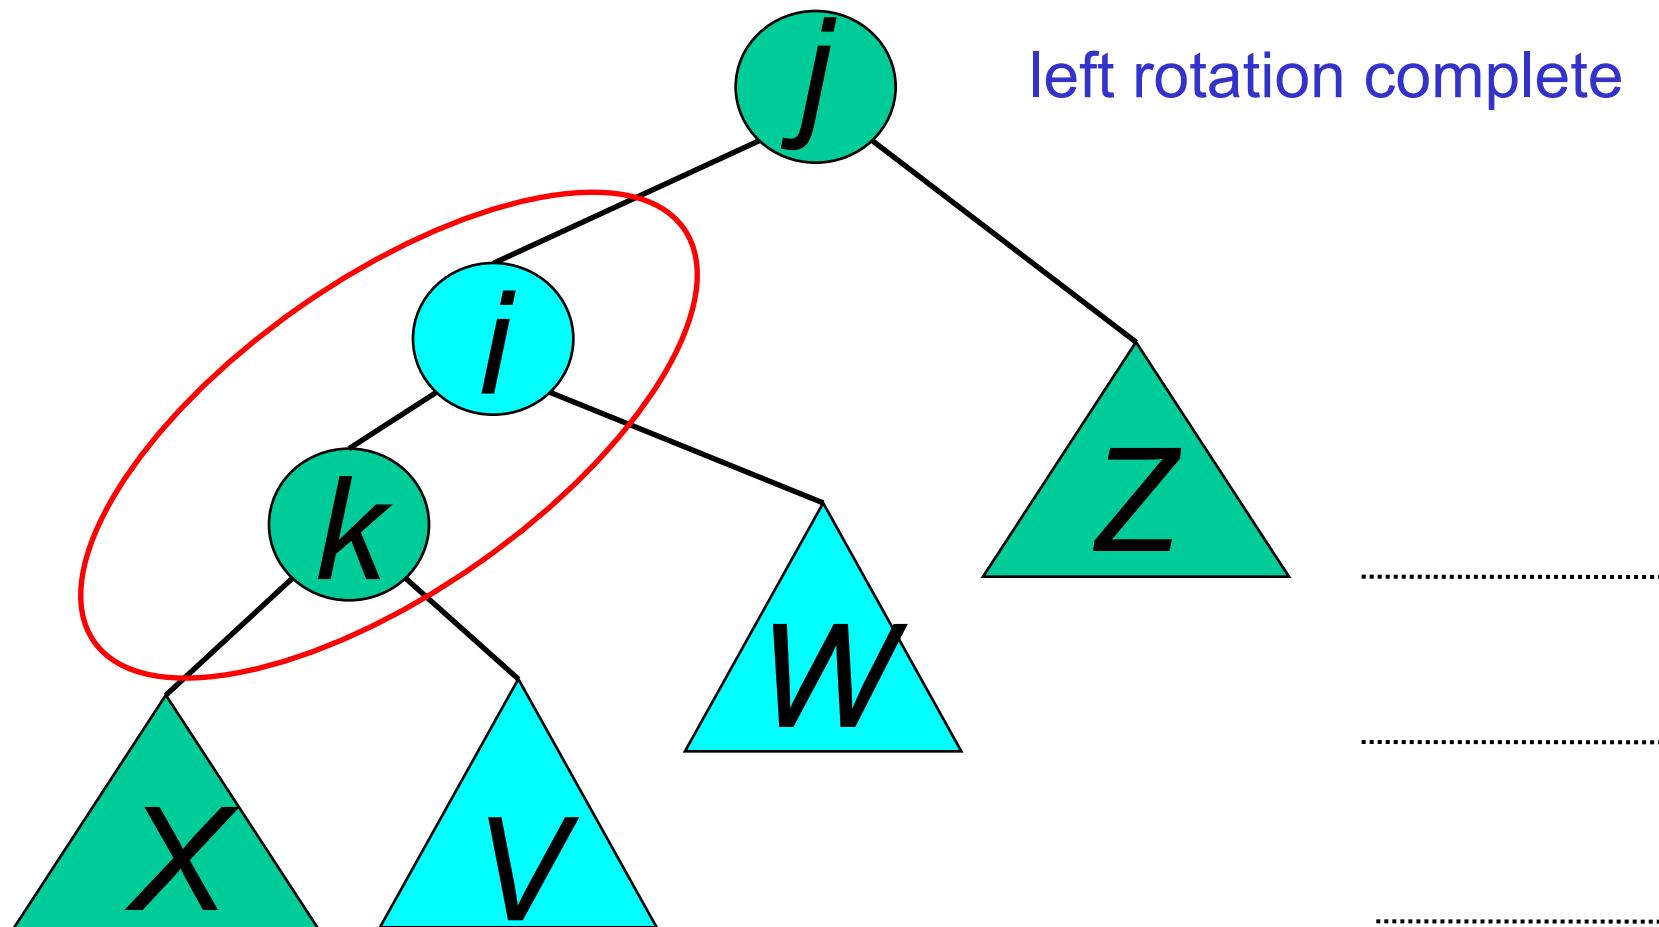
# AVL Insertion: Inside Case



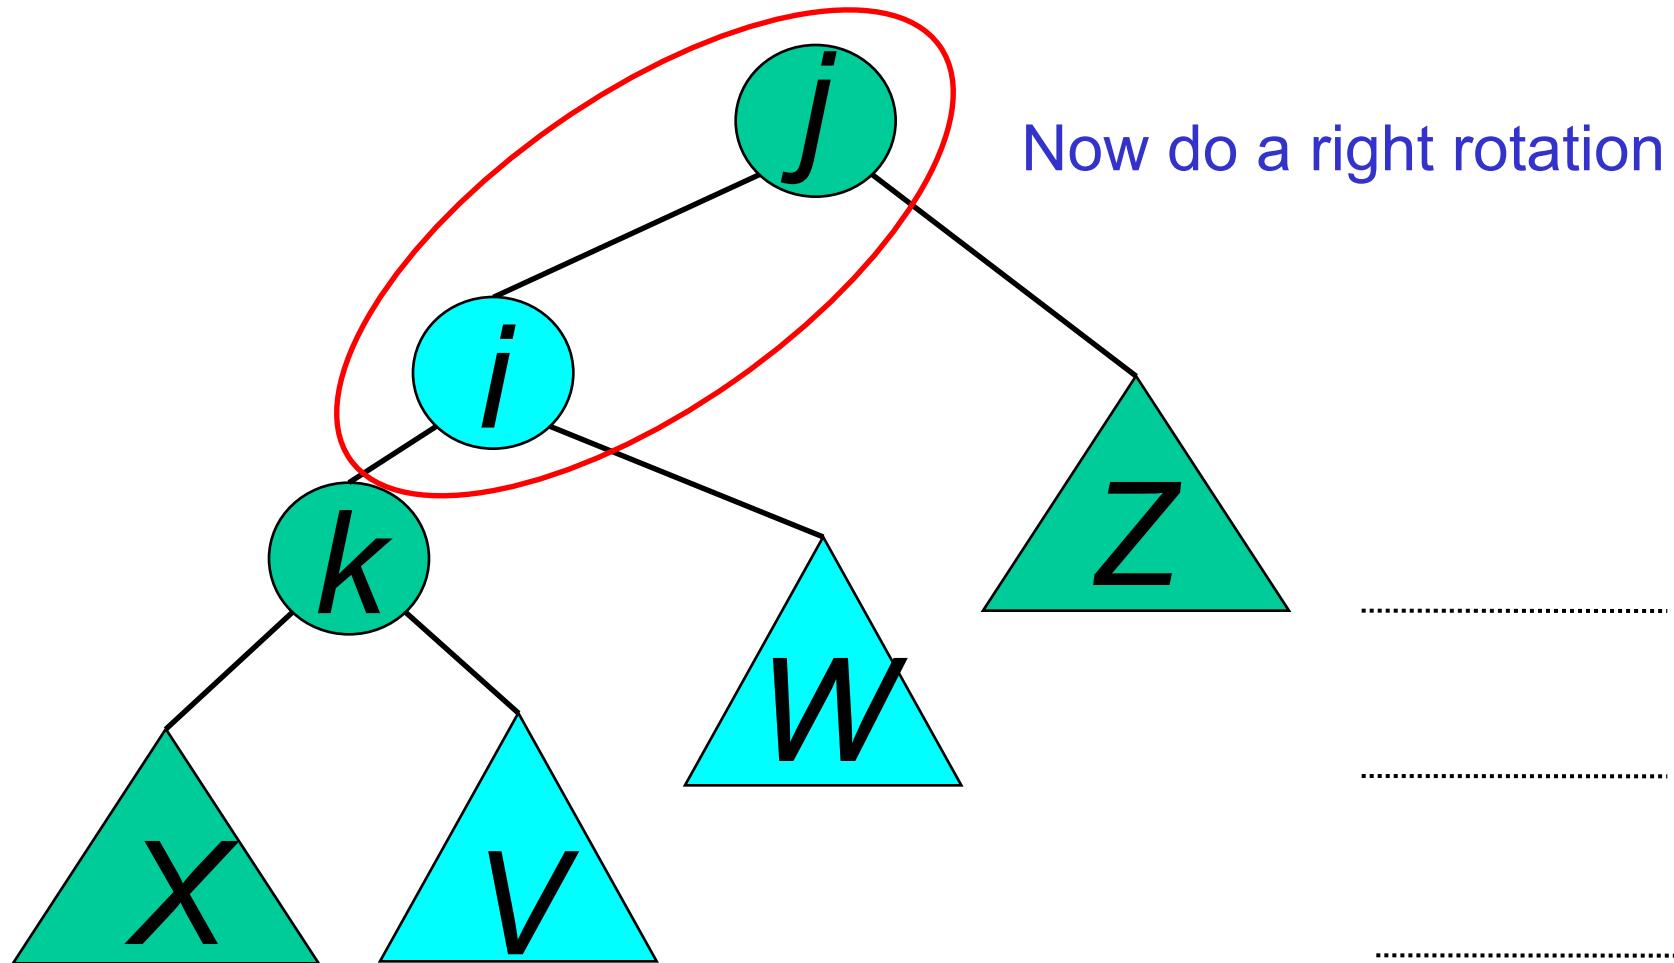
# AVL Insertion: Inside Case



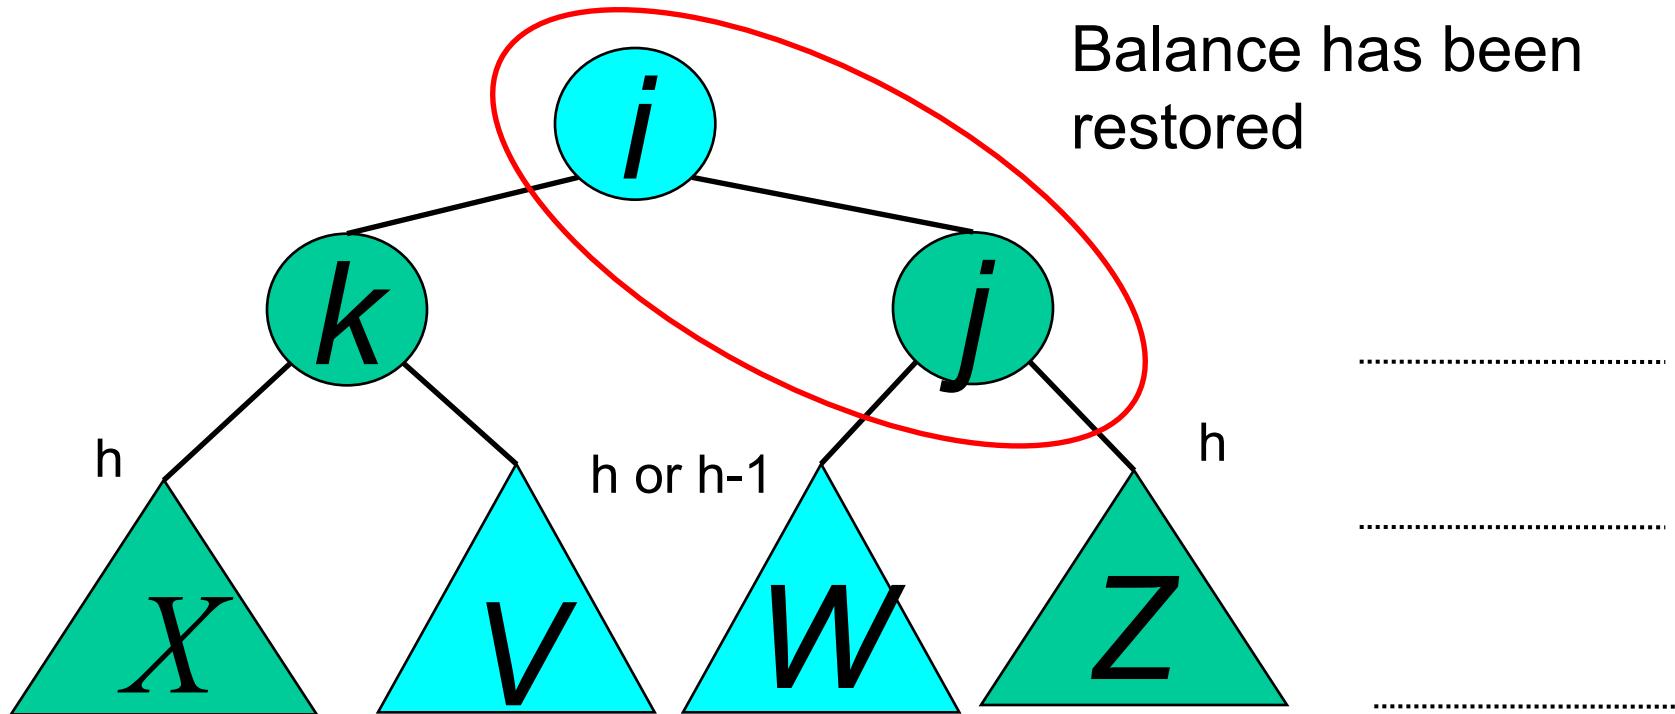
# Double rotation : first rotation



# Double rotation : second rotation

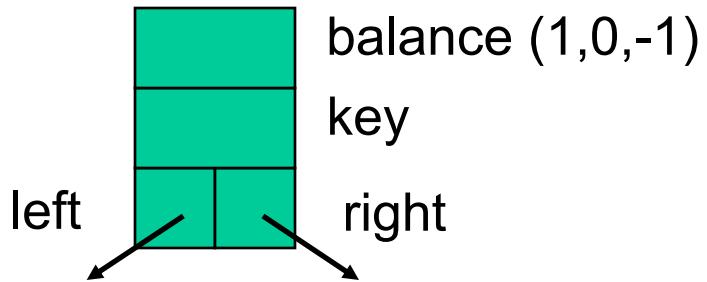


# Double rotation : second rotation



# Implementation

---



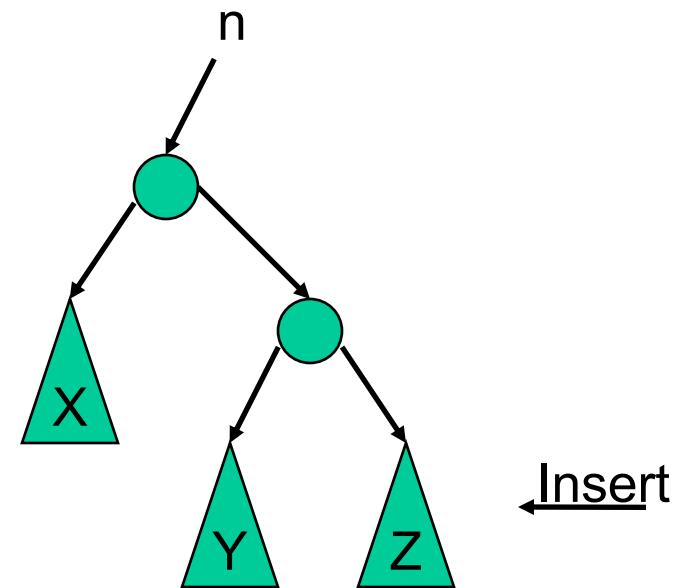
No need to keep the height; just the difference in height,  
i.e. the **balance** factor; this has to be modified on the path of  
insertion even if you don't perform rotations

Once you have performed a rotation (single or double) you  
won't need to go back up the tree

# Single Rotation

```
RotateFromRight(n : reference node pointer) {  
    p : node pointer;  
    p := n.right;  
    n.right := p.left;  
    p.left := n;  
    n := p  
}
```

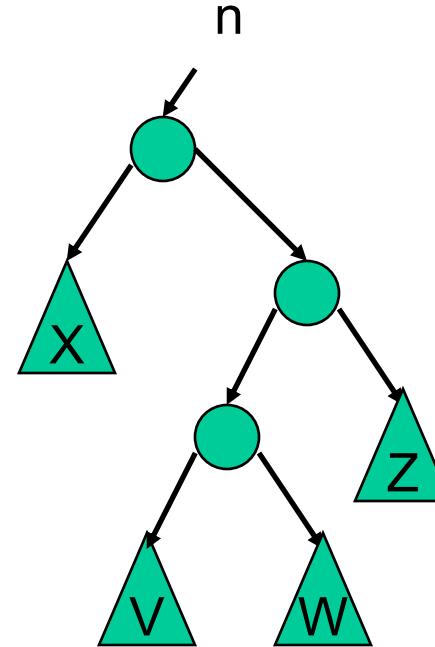
You also need to modify the heights or balance factors of n and p



# Double Rotation

- Implement Double Rotation in two lines.

```
DoubleRotateFromRight (n : reference node pointer) {  
    ???  
}
```



# Insertion in AVL Trees

---

- Insert at the leaf (as for all BST)
  - › only nodes on the path from insertion point to root node have possibly changed in height
  - › So after the Insert, go back up to the root node by node, updating heights
  - › If a new balance factor (the difference  $h_{\text{left}} - h_{\text{right}}$ ) is 2 or  $-2$ , adjust tree by *rotation* around the node

# Insert in BST

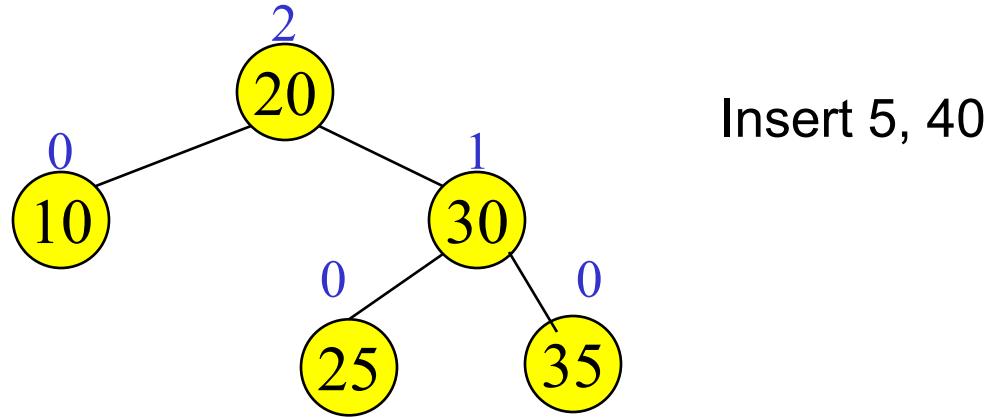
---

```
Insert(T : reference tree pointer, x : element) : integer {
    if T = null then
        T := new tree; T.data := x; return 1;
        //the links to children are null
    case
        T.data = x : return 0; //Duplicate do nothing
        T.data > x : return Insert(T.left, x);
        T.data < x : return Insert(T.right, x);
    endcase
}
```

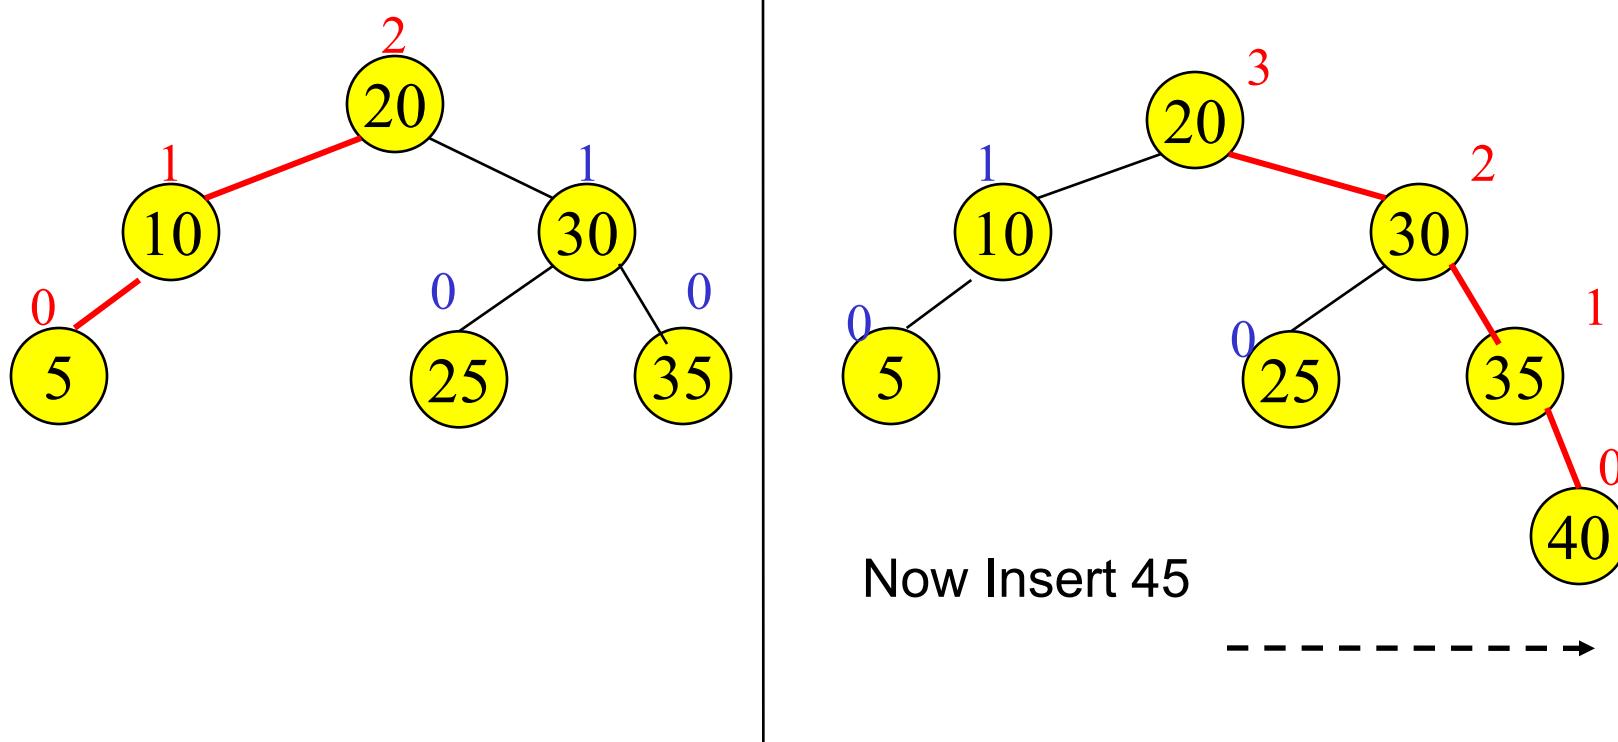
# Insert in AVL trees

```
Insert(T : reference tree pointer, x : element) : {
    if T = null then
        {T := new tree; T.data := x; height := 0; return;}
    case
        T.data = x : return ; //Duplicate do nothing
        T.data > x : Insert(T.left, x);
            if ((height(T.left)- height(T.right)) = 2) {
                if (T.left.data > x) then //outside case
                    T = RotatefromLeft(T);
                else                                //inside case
                    T = DoubleRotatefromLeft(T);}
        T.data < x : Insert(T.right, x);
            code similar to the left case
    endcase
    T.height := max(height(T.left),height(T.right)) +1;
    return;
}
```

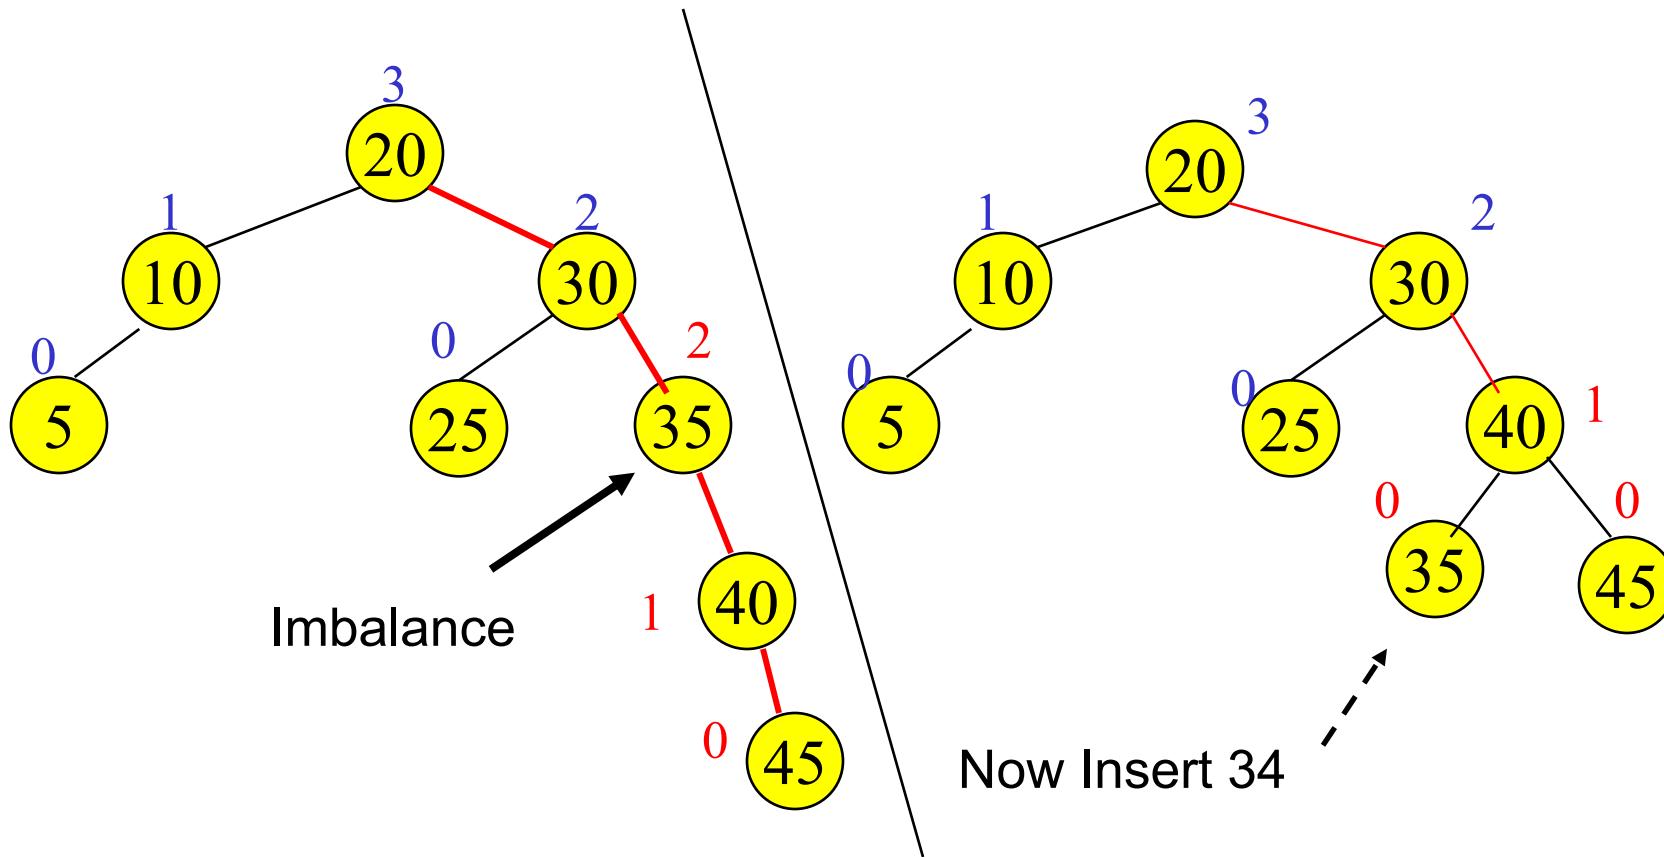
# Example of Insertions in an AVL Tree



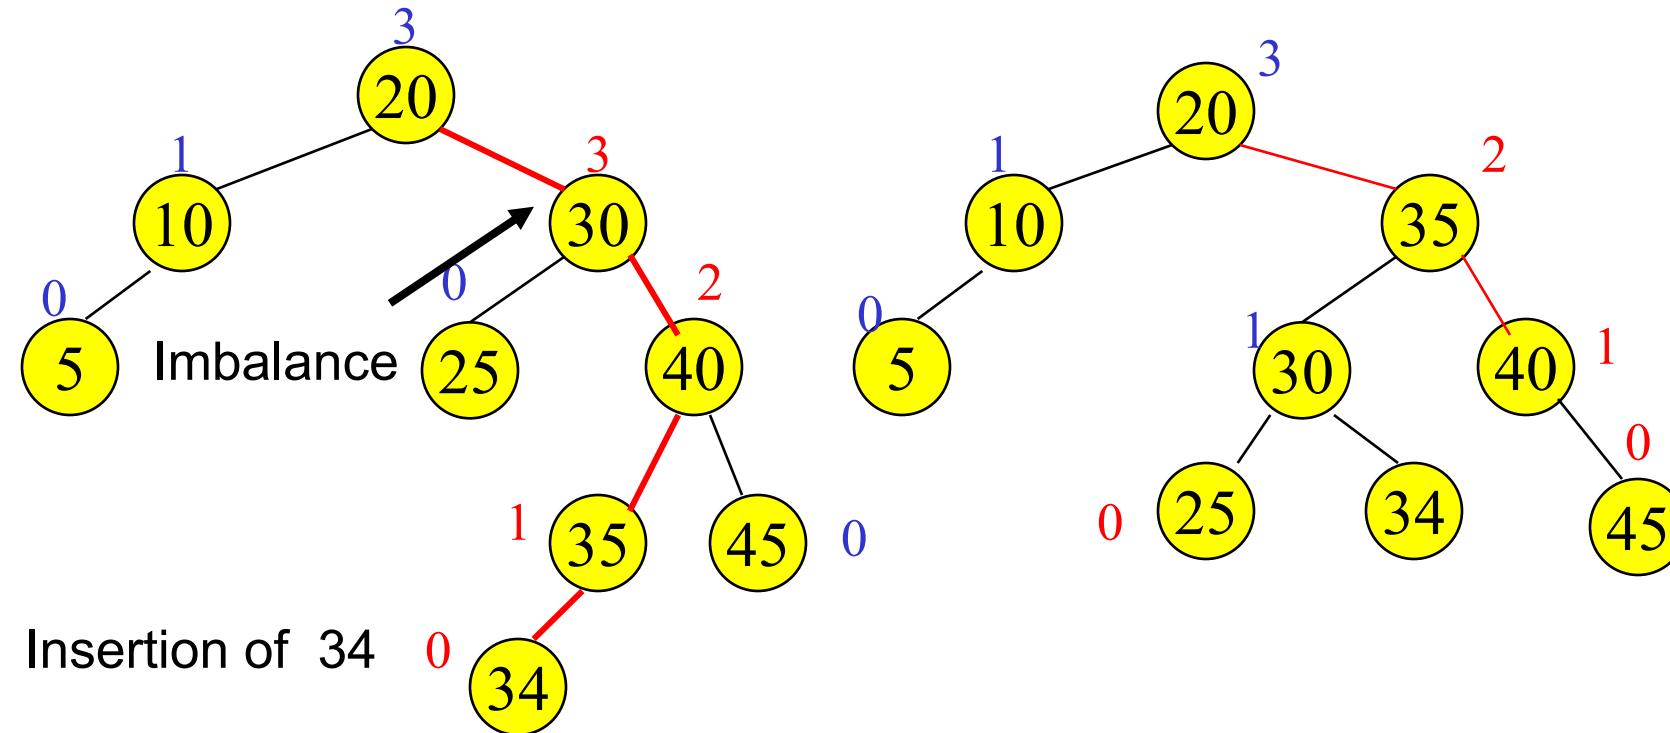
# Example of Insertions in an AVL Tree



# Single rotation (outside case)



# Double rotation (inside case)



# AVL Tree Deletion

---

- Similar but more complex than insertion
  - › Rotations and double rotations needed to rebalance
  - › Imbalance may propagate upward so that many rotations may be needed.

# Pros and Cons of AVL Trees

---

## Arguments for AVL trees:

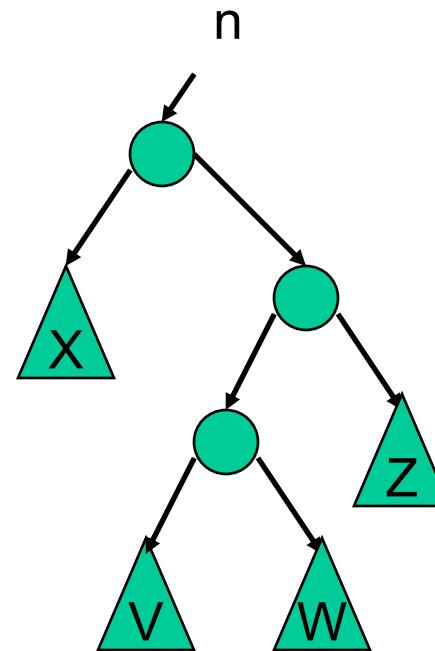
1. Search is  $O(\log N)$  since AVL trees are **always balanced**.
2. Insertion and deletions are also  $O(\log n)$
3. The height balancing adds no more than a constant factor to the speed of insertion.

## Arguments against using AVL trees:

1. Difficult to program & debug; more space for balance factor.
2. Asymptotically faster but rebalancing costs time.
3. Most large searches are done in database systems on disk and use other structures (e.g. B-trees).
4. May be OK to have  $O(N)$  for a single operation if total run time for many consecutive operations is fast (e.g. Splay trees).

# Double Rotation Solution

```
DoubleRotateFromRight (n : reference node pointer) {  
    RotateFromLeft (n.right) ;  
    RotateFromRight (n) ;  
}
```



# Advantages / Disadvantages of AVL Trees

AVL Trees	
+	-
Guaranteed $\mathcal{O}(\log_\phi n)$ performance of <b>all operations</b> .	Spend a lot of time ( <i>almost after every operation!</i> ) on rotations, which are themselves expensive operations
<i>Biggest disadvantages of AVL Trees</i>	Need to store at least 4 bits per node to preserve balance information (plus the unit cost of updating it)
• • •	Complex implementation (deletions in particular)
• • •	

# Maintain the pros, kill the cons

---

- Is there any way we can maintain mission-critical logarithmic worst-case complexity for search, while simultaneously shedding the storage cost and rotation overhead of AVL Trees?
- Yup! Two solutions:
  - 2-3 trees implemented via Red-Black Binary Search Trees (RBBSTs)
  - B-Trees (BTs)
- Today, we examine the simplest B-tree, known as a 2-3 tree

# 2-3 Trees

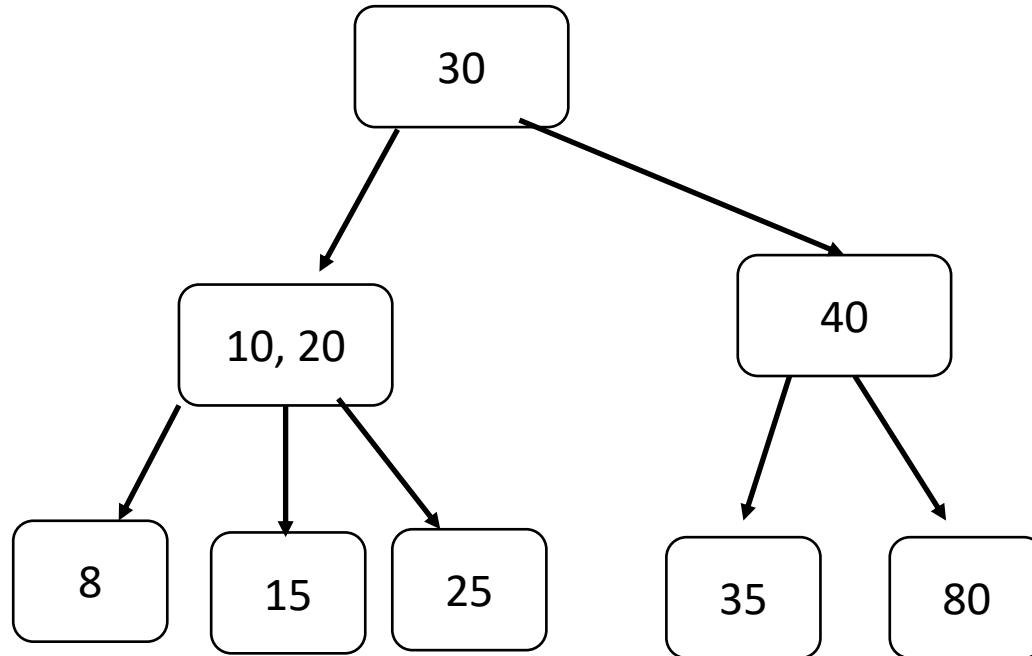
---

# 2-3 Trees

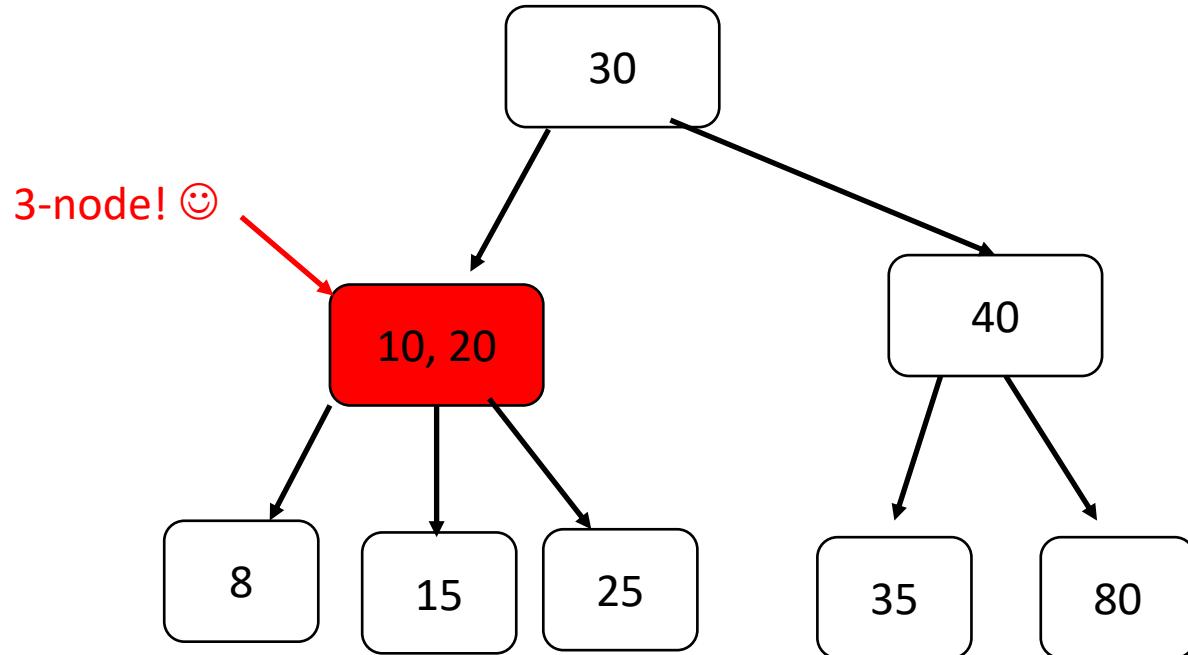
---

- In addition to their other properties, 2-3 trees (and B-trees) will have **perfect balance** all the time!
- Core idea: 2 children per node are ok, **more are better (for height)**.
- 2-3 trees will have two different nodes inside them!
  - a) A **BST-like node** with just one key, and two pointers to subtrees (**a 2-node**)
  - b) An expanded node with three keys and three pointers to subtrees (**a 3-node**)
    - **Leftmost** subtree contains keys  $< Key_1$
    - **Center** subtree contains keys in  $[Key_1, Key_2)$
    - **Rightmost** subtree contains keys  $> Key_2$

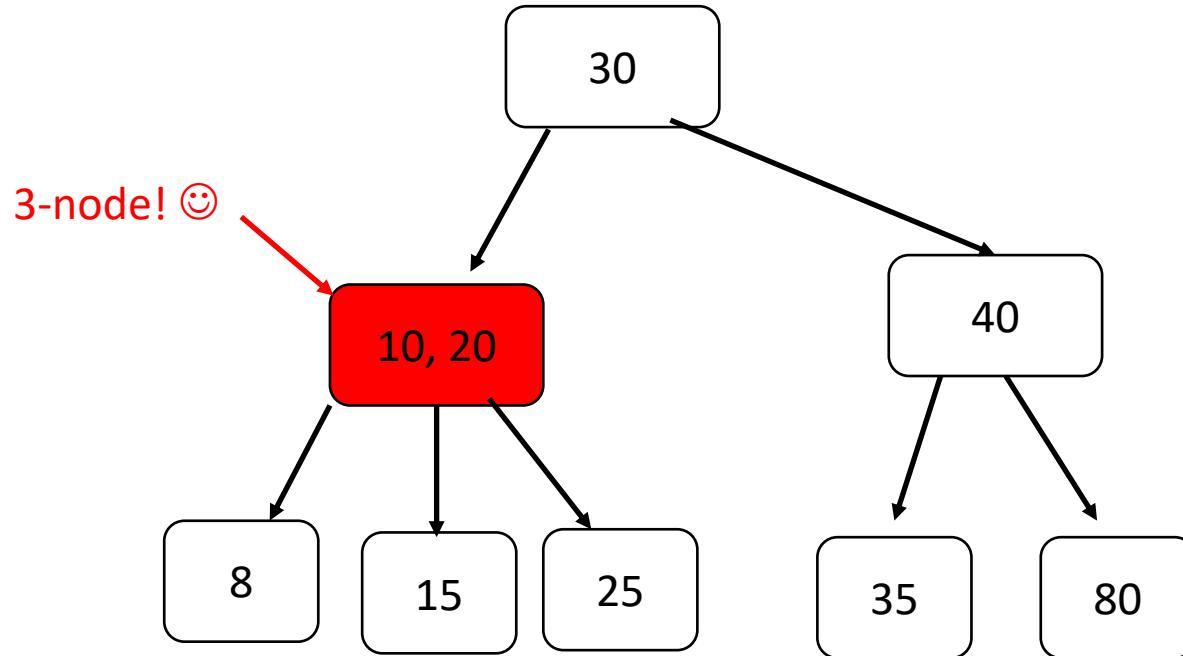
# 2-3 Tree example



# 2-3 Tree example



# 2-3 Tree example



The **minimum** height of a 2-3 tree with  $n$  keys is...

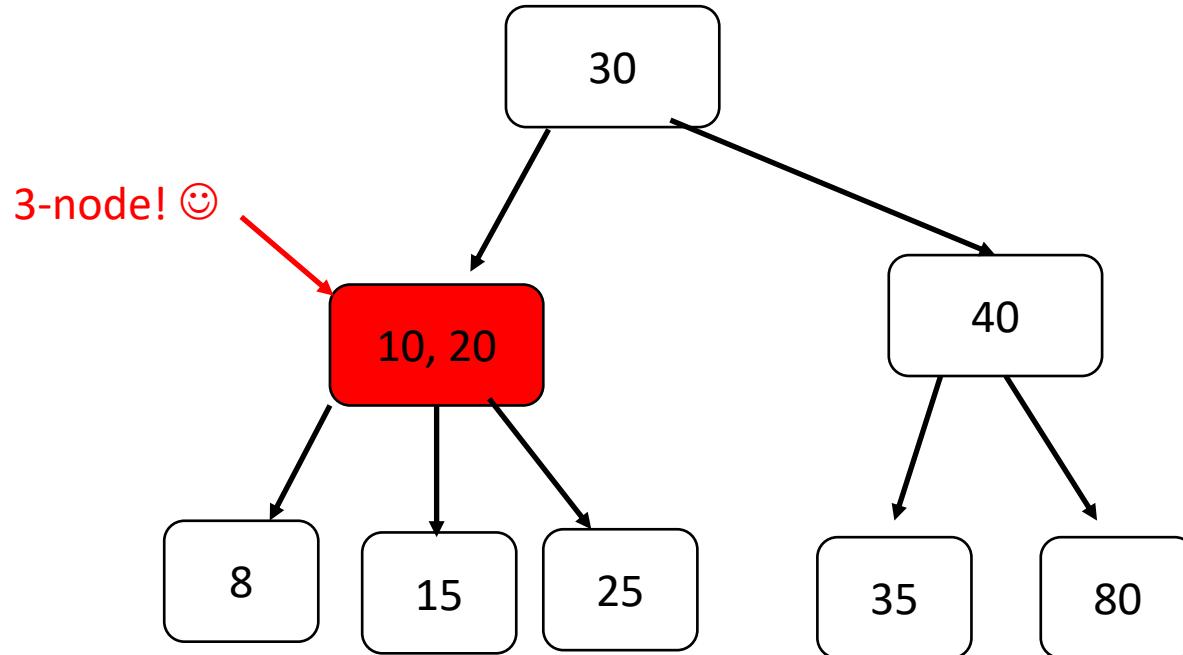
$\lfloor \log_2 n \rfloor$

$\lfloor \log_2 n \rfloor + 1$

$\lceil \log_2 n \rceil$

$\log_3 n$

# 2-3 Tree example



The **minimum** height of a 2-3 tree with  $n$  keys is...

$\lfloor \log_2 n \rfloor$

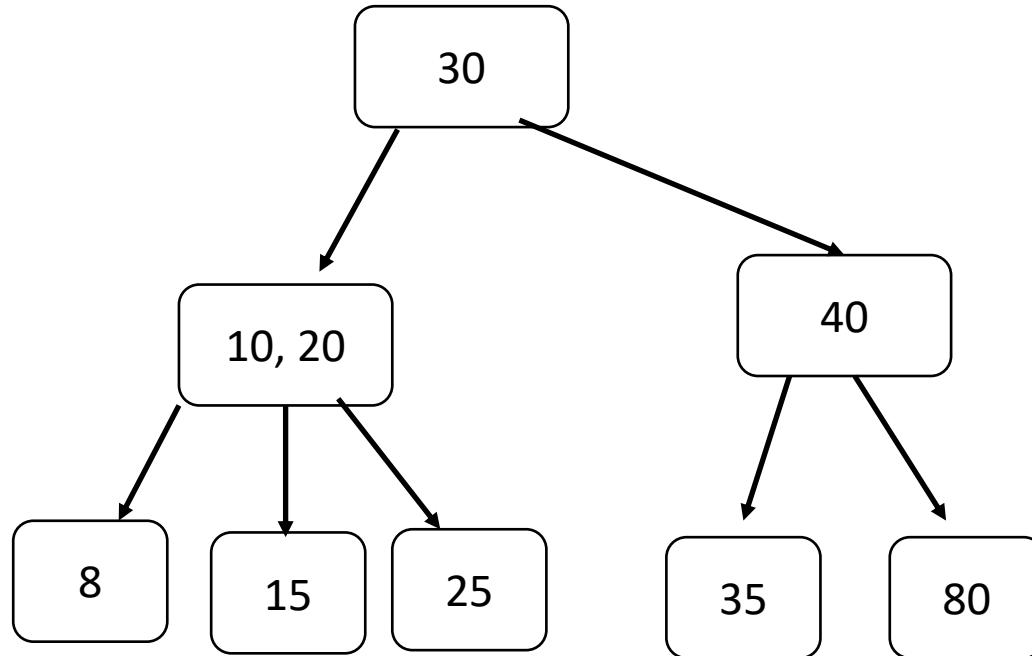
$\lfloor \log_2 n \rfloor + 1$

$\lceil \log_2 n \rceil$

$\lfloor \log_3 n \rfloor$

All 3-nodes!

# Search



- Searching a 2-3-Tree is done in **exactly the same way as a BST**.
- **Stronger efficiency guarantees** in place since the height is between  $\log_3 n$  and  $\log_2 n$ ! ☺

# Insertion

---

- Let's recall insertion in AVL ...
  - **AVL: Might trigger rotations** to re-balance the tree according to the AVL condition.
- This **will leave the tree unbalanced** according to the strict most criterion of... **same height everywhere!**
  - Only exception: AVL insertion places node in the only empty position at the leaves of an AVL Tree.

# Insertion

---

- Insertion on both of these tree has two phases...
  1. A **search** phase, where we **go down** to search for an empty position to put the key in (if it's not already in there), and
  2. A **rotation** phase where we go up to either rebalance subtrees (AVL)...

# Insertion

---

- Insertion on both of these trees has two phases...
  1. A **search** phase, where we **go down** to search for an empty position to put the key in (if it's not already in there), and
  2. A **rotation** phase where we go up to either rebalance subtrees (AVL) ...
- B-Trees have the same search phase, **BUT!** in order to make sure the entire tree is perfectly balanced, the second phase will rotate **exactly zero nodes!**

# Insertion

---

- Insertion on both of these trees has two phases...
  1. A **search** phase, where we **go down** to search for an empty position to put the key in (if it's not already in there), and
  2. A **rotation** phase where we go up to either rebalance subtrees (AVL)...
- B-Trees have the same search phase, **BUT!** in order to make sure the entire tree is perfectly balanced, the second phase will rotate **exactly zero nodes!**
- Instead, it *might* either make nodes **fatter** or blow them up into **two nodes**.

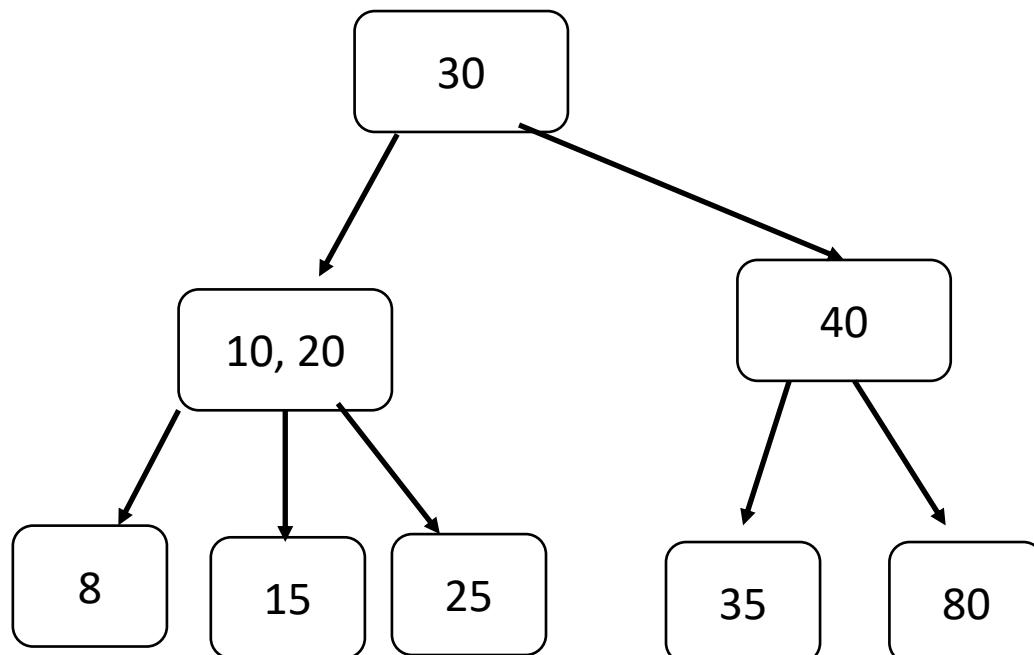
# Insertion

---

- Insertion on both of these trees has two phases...
  1. A **search** phase, where we **go down** to search for an empty position to put the key in (if it's not already in there), and
  2. A **rotation** phase where we go up to either rebalance subtrees (AVL)...
- B-Trees have the same search phase, **BUT!** in order to make sure the entire tree is perfectly balanced, the second phase will rotate **exactly zero nodes!**
- Instead, it *might* either make nodes **fatter** or blow them up into **two nodes**.
  - We'll see how this happens **right now**.

# Insertion

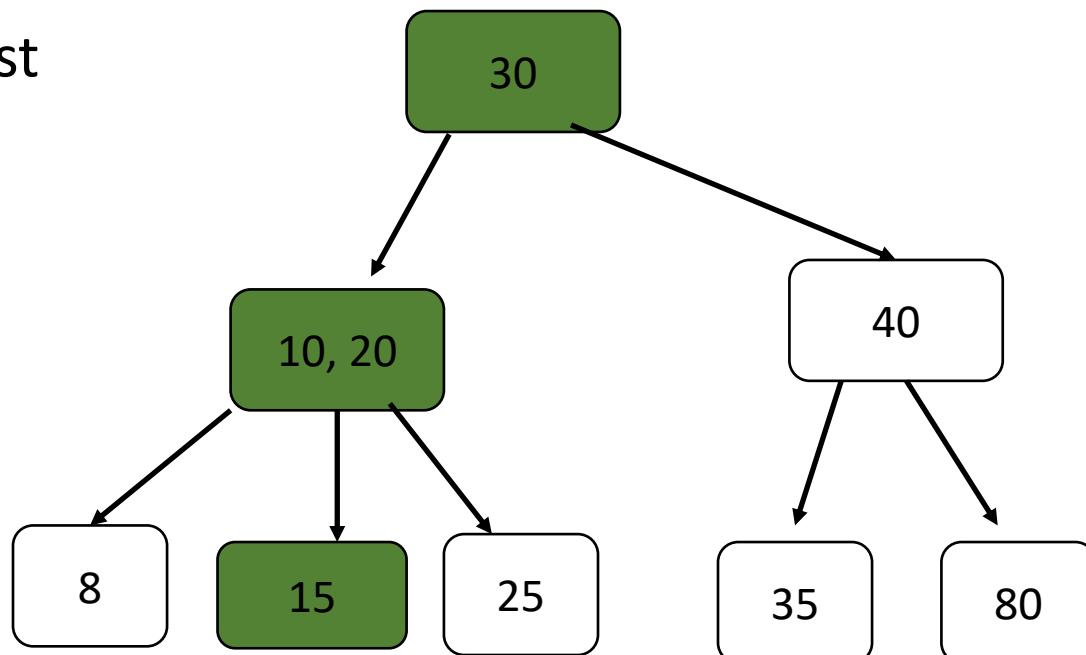
- Insert 18



# Insertion

- Insert 18

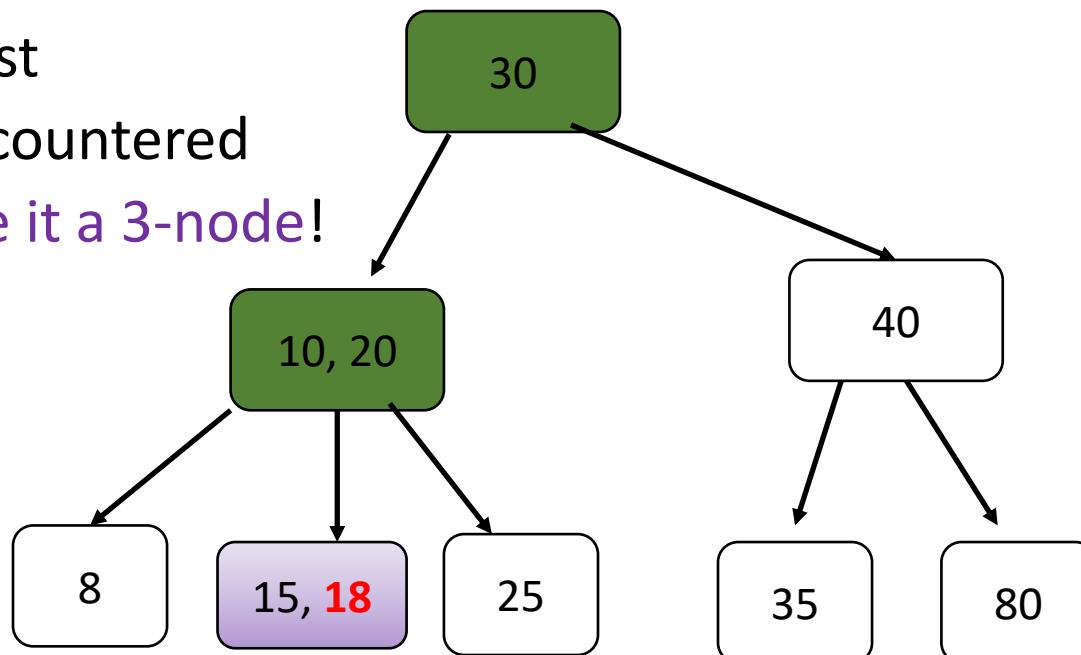
1. We search first



# Insertion

- Insert 18

1. We search first
2. Last node encountered  
is 2-node: make it a 3-node!

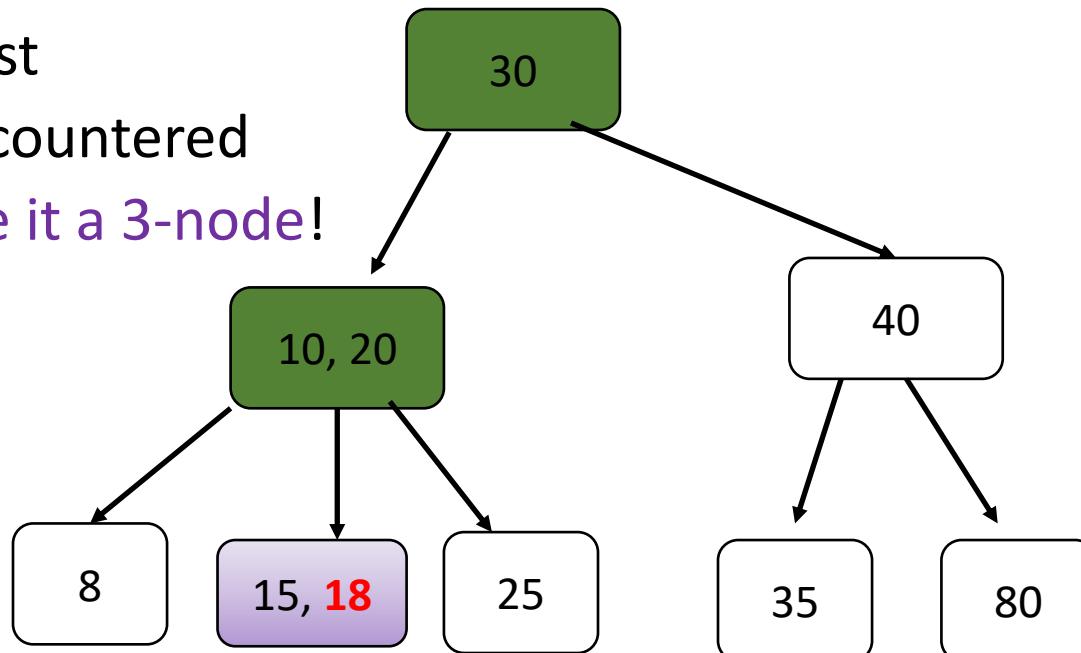


# Insertion

- Insert 18

1. We search first
2. Last node encountered  
is 2-node: make it a 3-node!

- Done! ☺

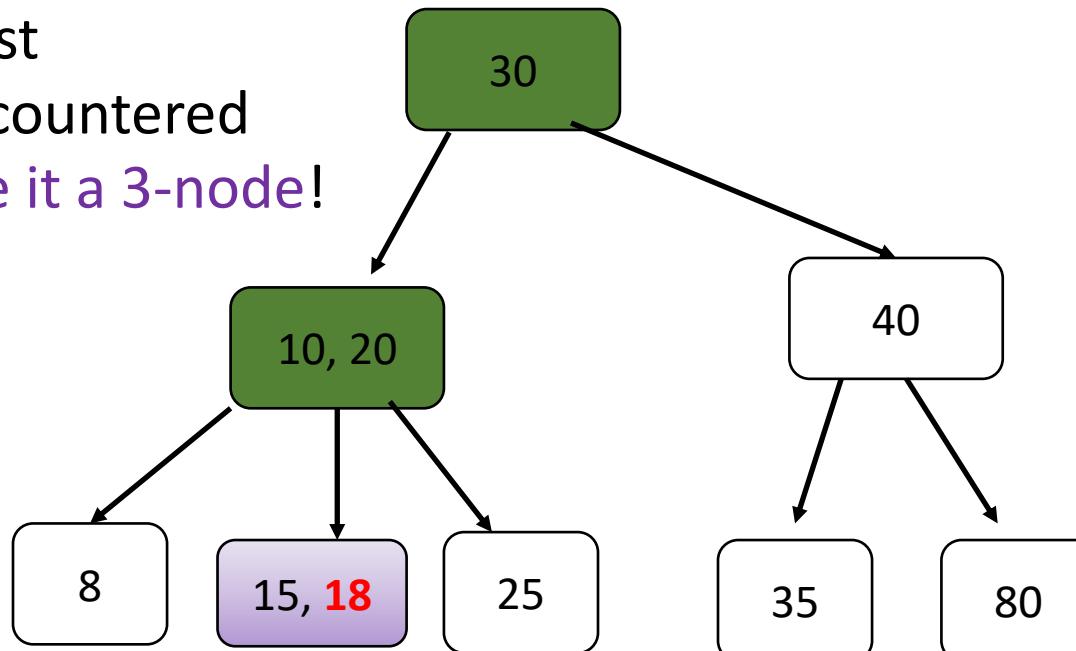


# Insertion

- Insert 18

1. We search first
2. Last node encountered  
is 2-node: make it a 3-node!

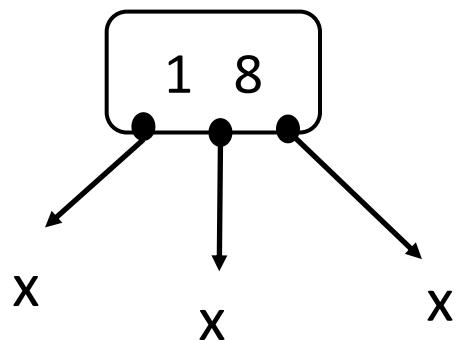
- Done! 😊



- But what if I now wanted to insert 17? I don't have space for it!

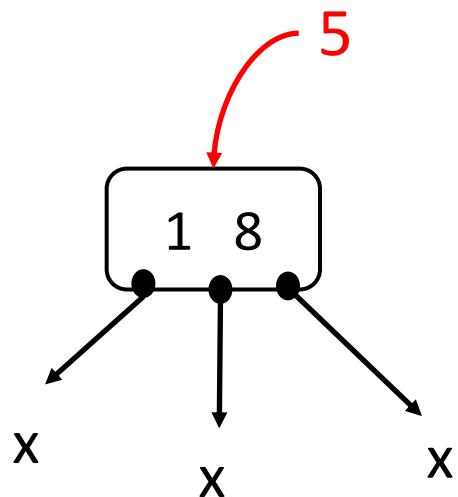
# Insertion

- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.



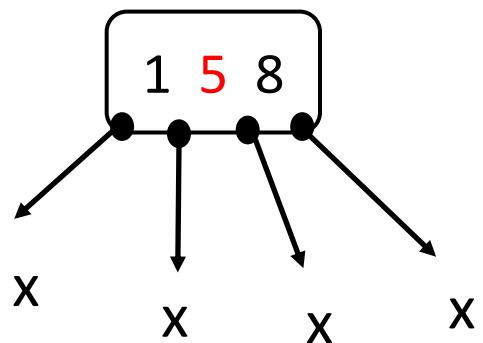
# Insertion

- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.



# Insertion

- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.

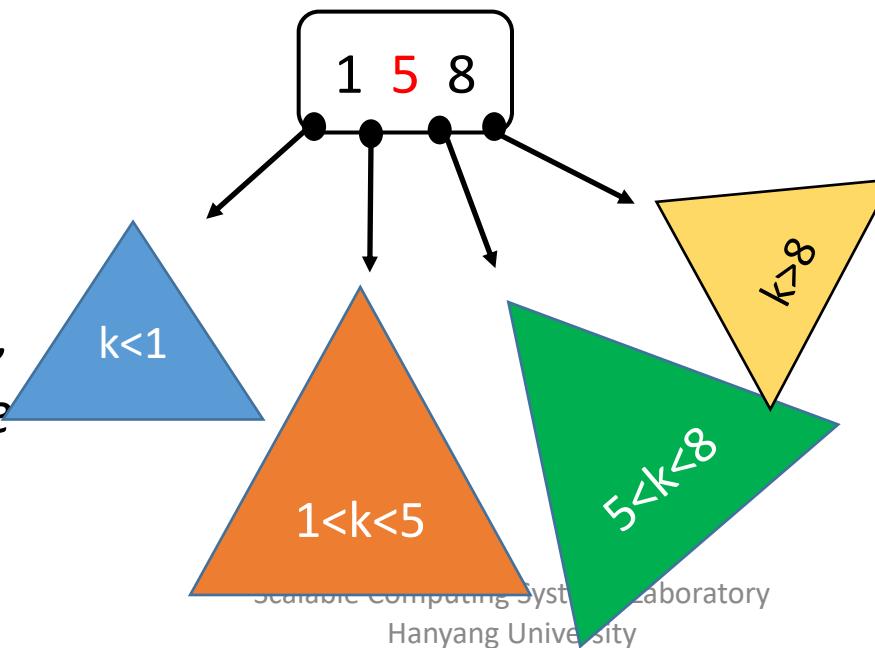


Temporarily assume we can expand the 3-node to a 4-node!

# Insertion

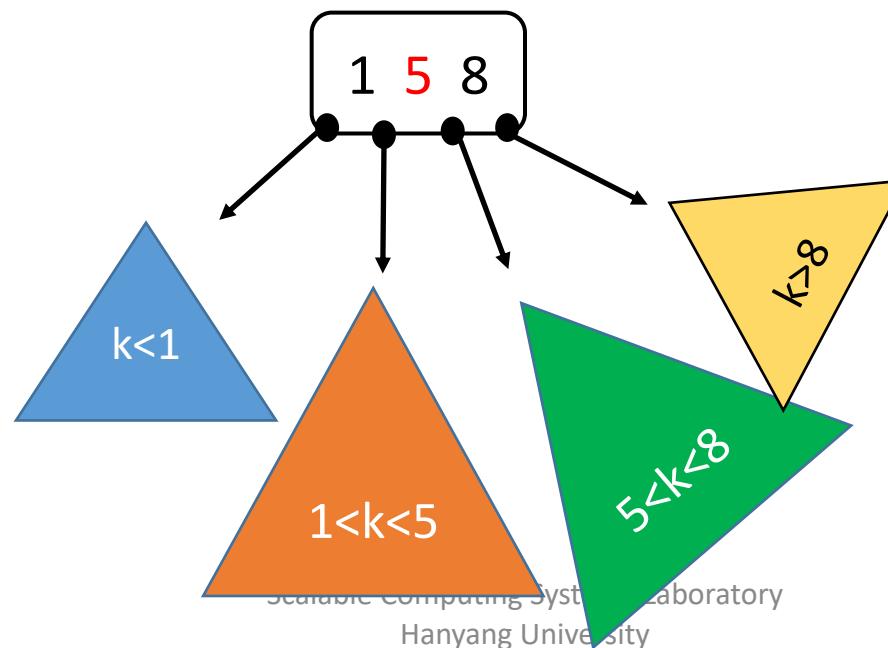
- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.

*(if the links were non-null, those would be the subtrees / key ranges they would point to)*



# Insertion

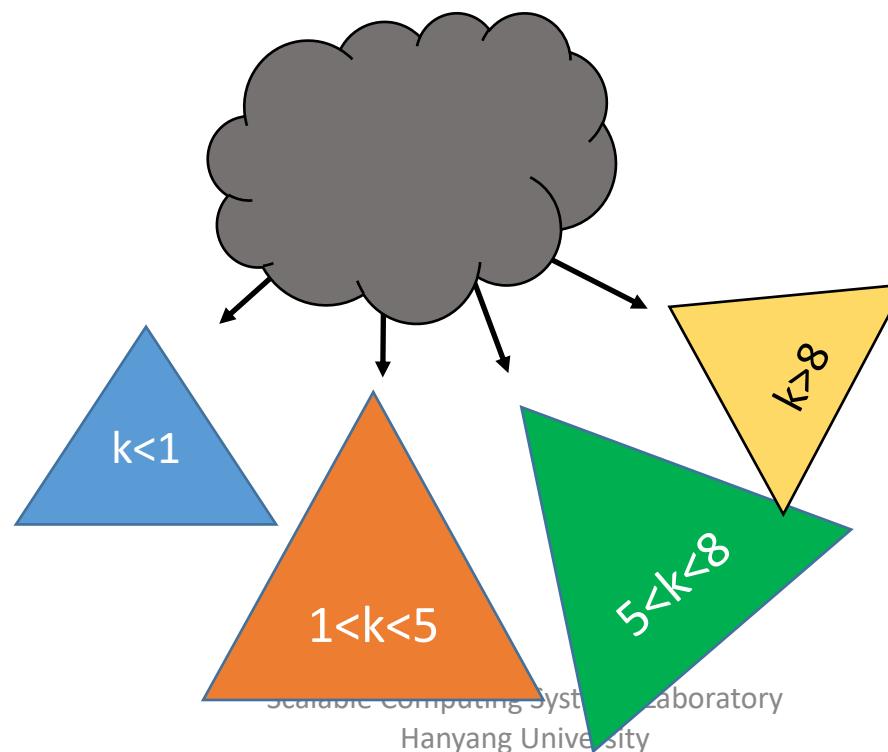
- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.



Of course, this is **unacceptable**. Our solution:

# Insertion

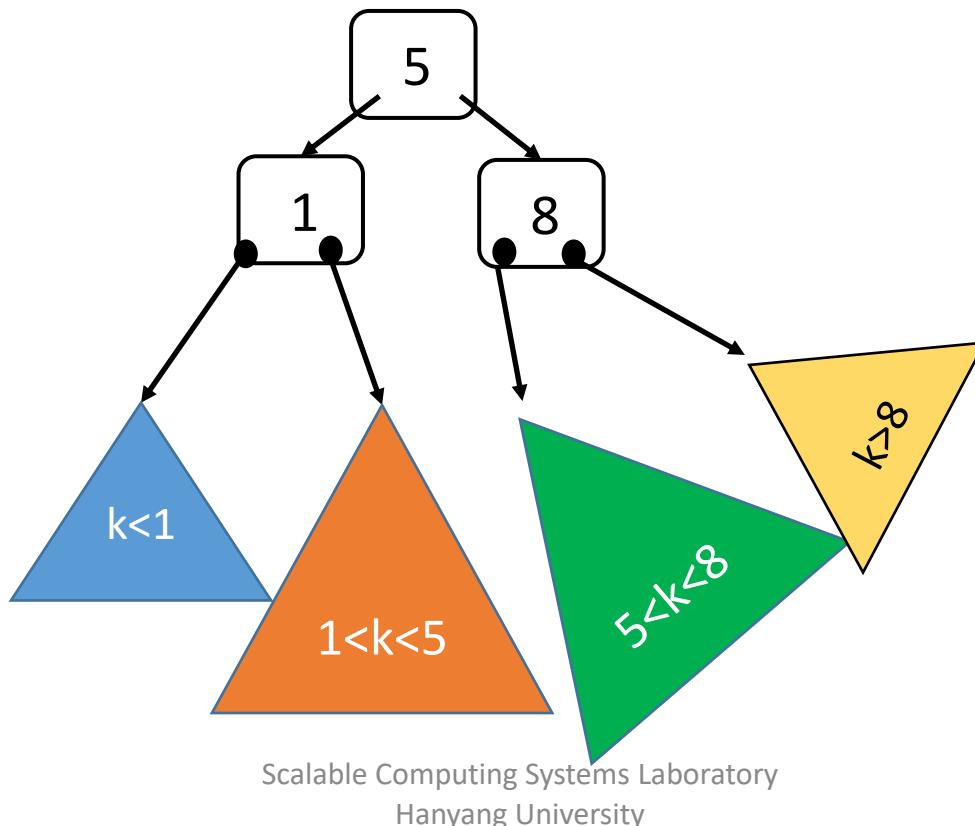
- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.



Of course, this is **unacceptable**. Our solution: **Blow the node apart!**

# Insertion

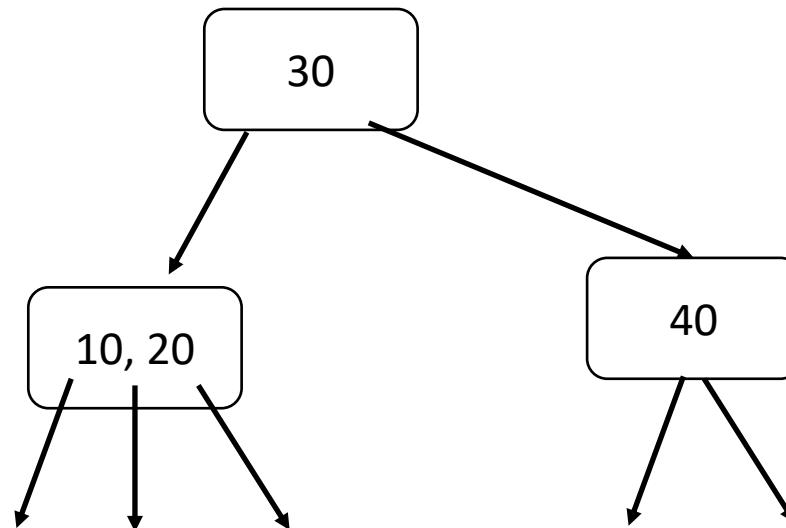
- To understand what to do in this case, consider inserting **5** in this “stub” 2-3 tree.



We are left with **three 2-nodes** that point to the correct subtrees!

# For you!

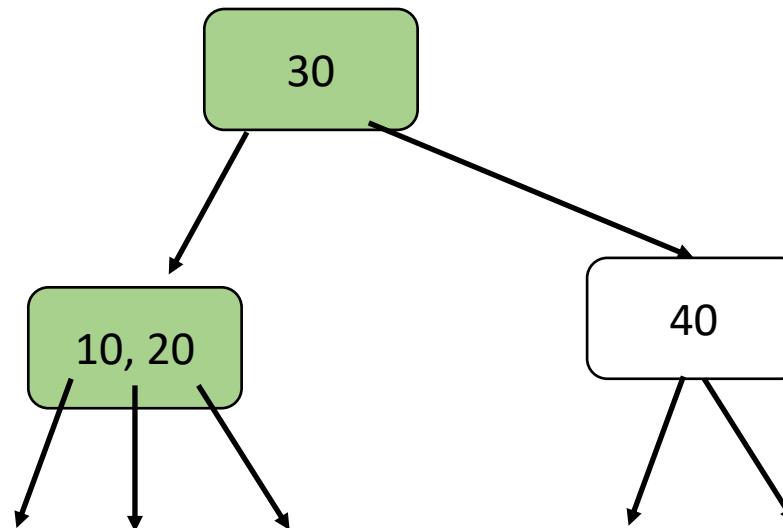
- Insert 15 in this 2-3 tree please! 😊



# Walkthrough...

- Insert 15 in this 2-3 tree please! 😊

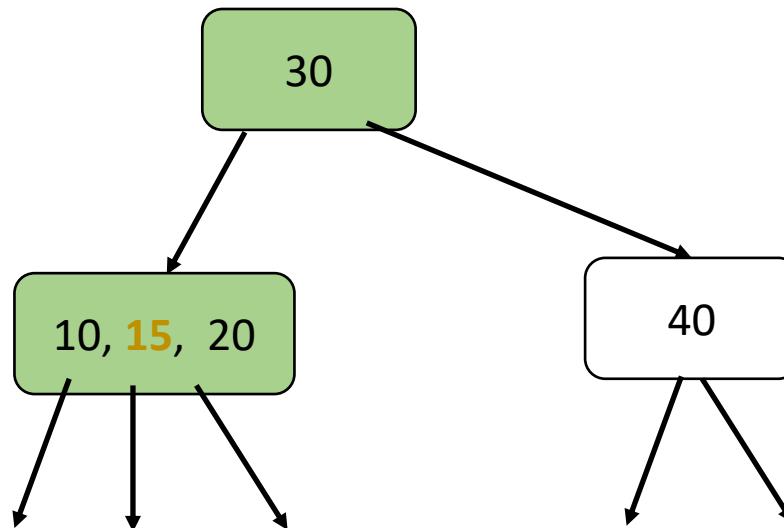
1. Search



# Walkthrough...

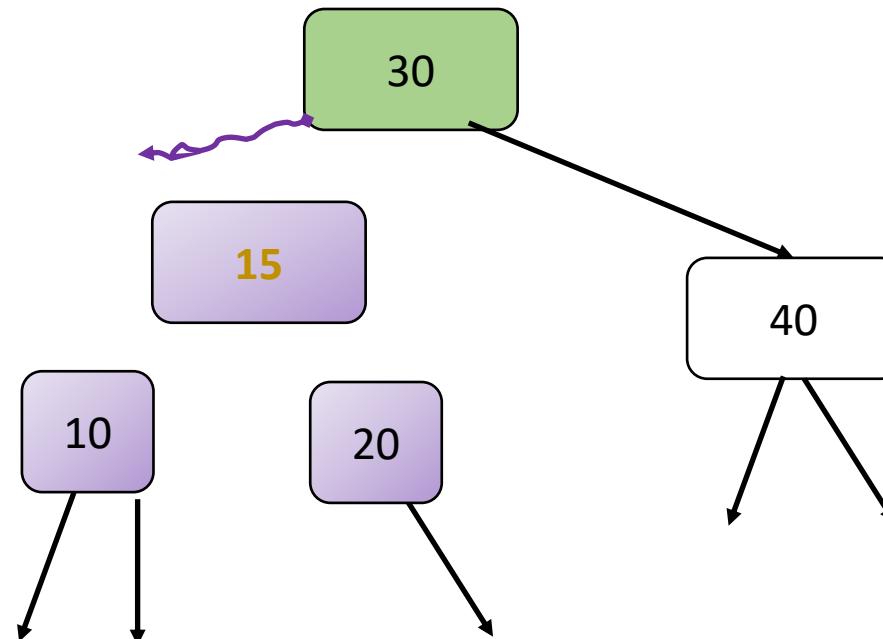
- Insert 15 in this 2-3 tree please! ☺

1. Search
2. Pretend



# Walkthrough...

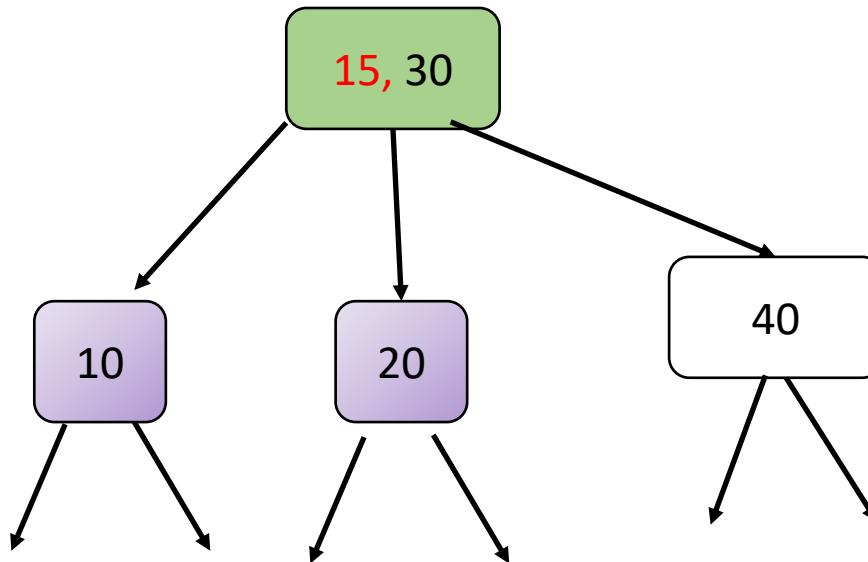
- Insert 15 in this 2-3 tree please! ☺



1. Search
2. Pretend
3. Kaboom

# Walkthrough...

- Insert 15 in this 2-3 tree please! ☺

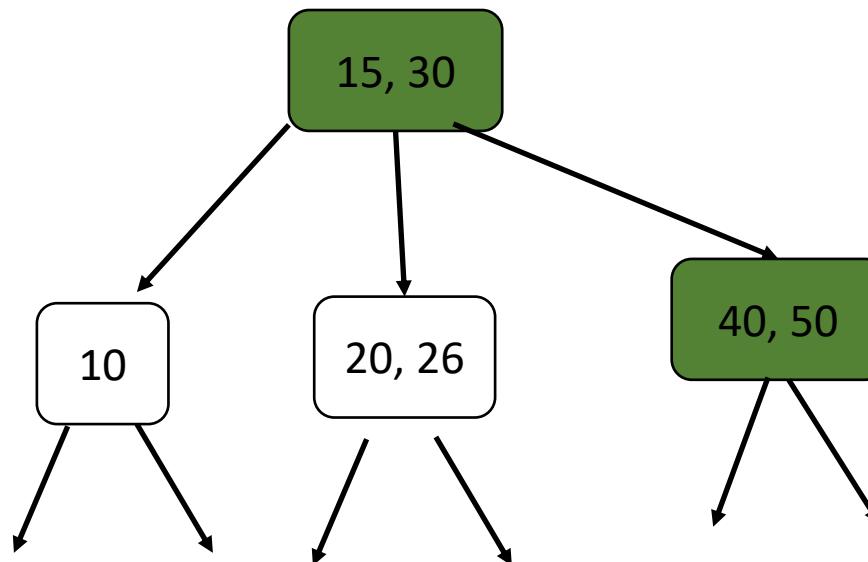


1. Search
2. Pretend
3. Kaboom
4. Merge with 2-node root, distributing children accordingly!

# For you!

- Let's now insert 43 in this tree...

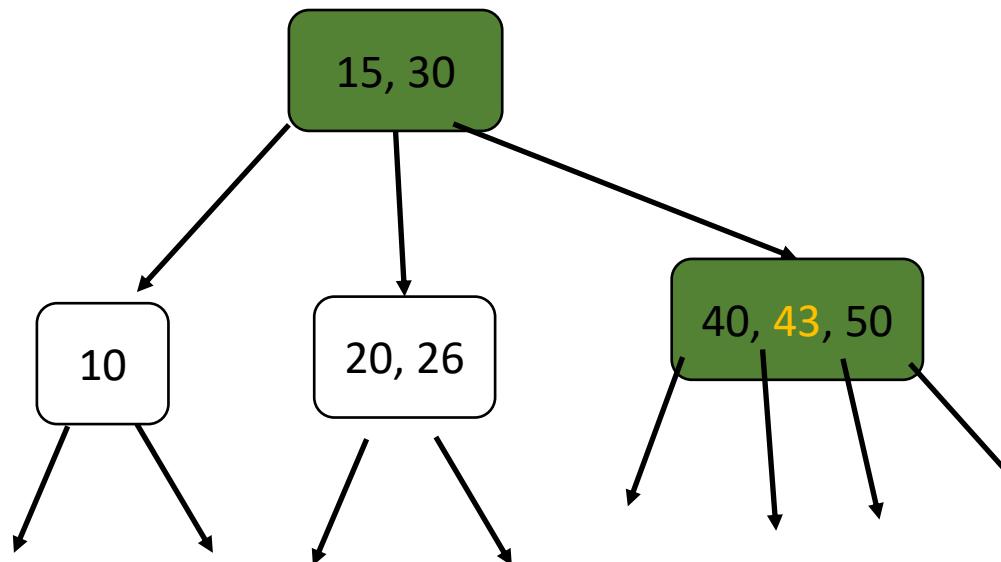
## 1. Search



# For you!

- Let's now insert 43 in this tree...

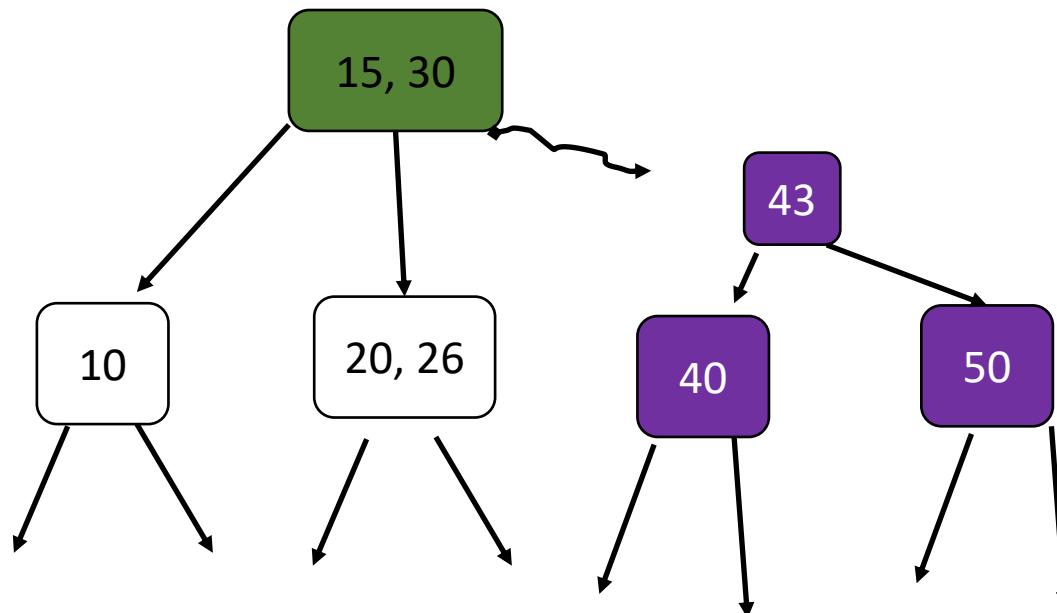
- Search
- Pretend



# For you!

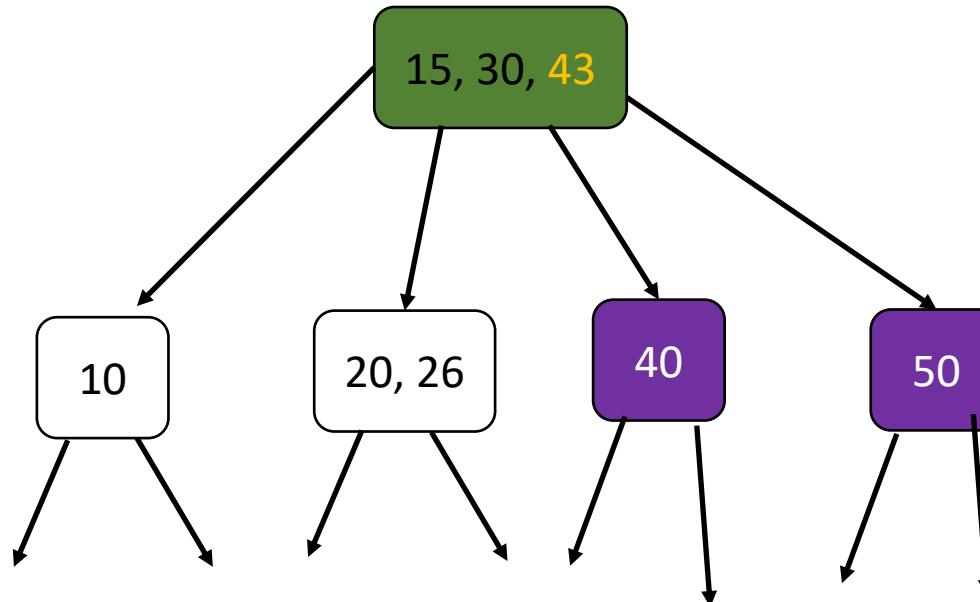
- Let's now insert 43 in this tree...

- Search
- Pretend
- Split



# For you!

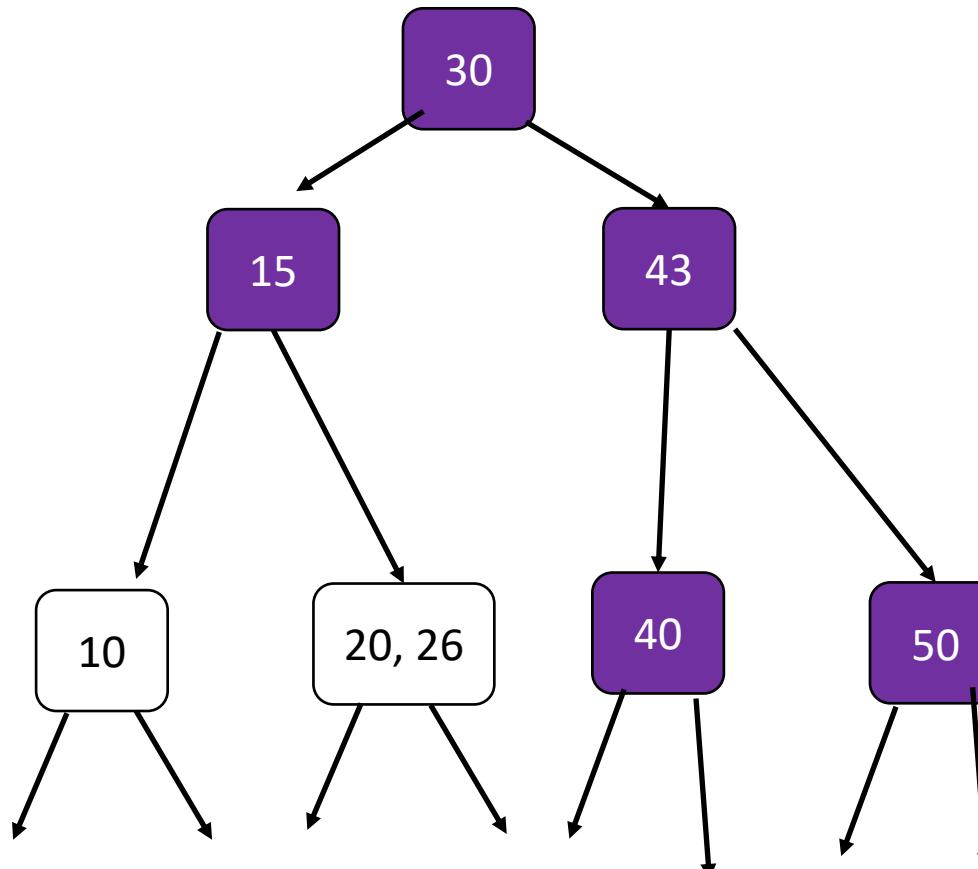
- Let's now insert 43 in this tree...



- Search
- Pretend
- Split
- Pretend

# For you!

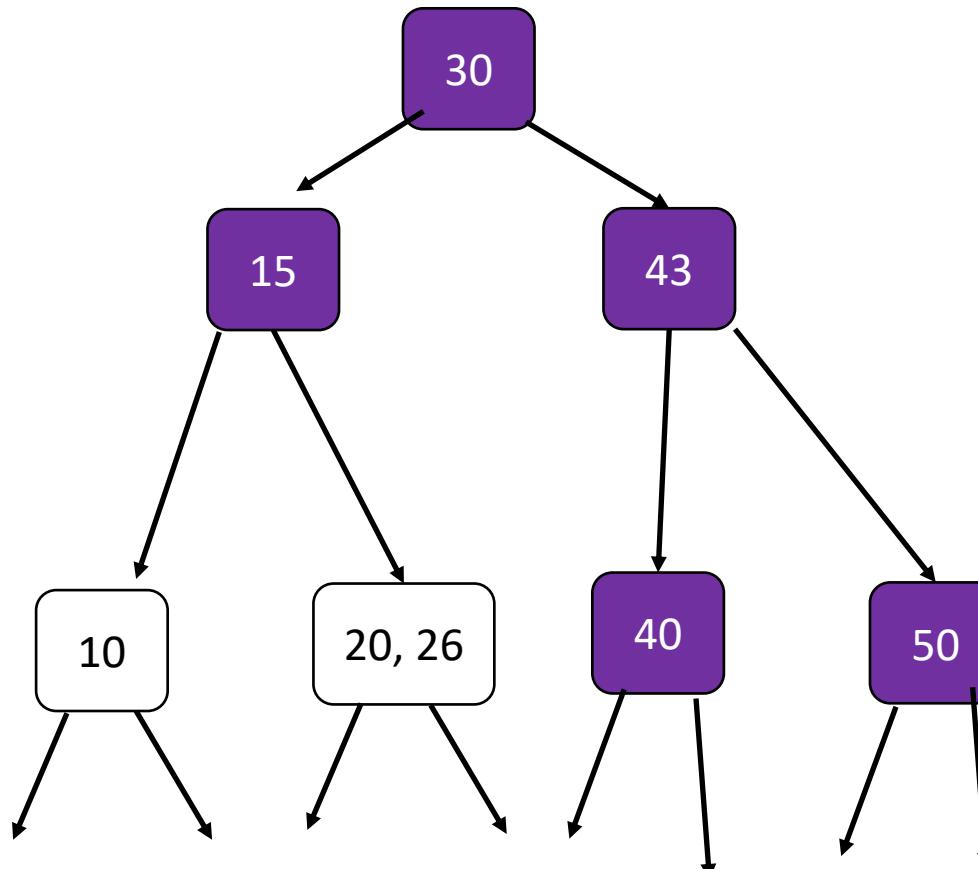
- Let's now insert 43 in this tree...



- Search
- Pretend
- Split
- Pretend
- Split

# For you!

- Let's now insert 43 in this tree...



- Search
- Pretend
- Split
- Pretend
- Split

Done!

# Deletion

---

- Insertions in 2-3 trees might transform a 2-node into a 3-node, or blow up a temporary 4-node into 3 2-nodes.
- Deletions, on the other hand, have the potential of not just making 3-nodes into 2-nodes, but making 2-nodes into ...

# Deletion

---

- Insertions in 2-3 trees might transform a 2-node into a 3-node, or blow up a temporary 4-node into 3 2-nodes.
- Deletions, on the other hand, have the potential of not just making 3-nodes into 2-nodes, but making 2-nodes into ... ?

# Deletion

---

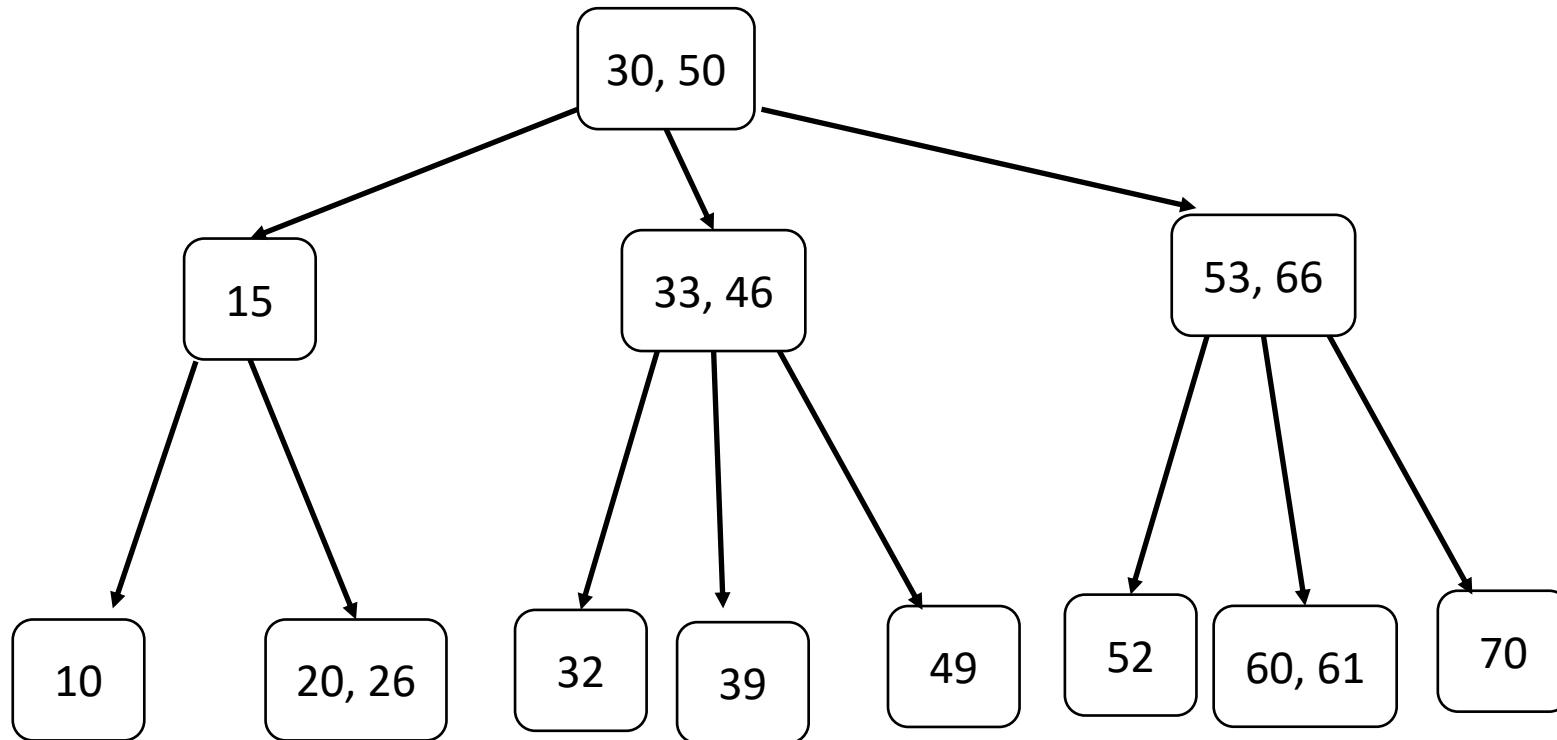
- Insertions in 2-3 trees might transform a 2-node into a 3-node, or blow up a temporary 4-node into 3 2-nodes.
- Deletions, on the other hand, have the potential of not just making 3-nodes into 2-nodes, but making 2-nodes into ... ?
- We'll see what the result of that will be soon.

# Deletion

---

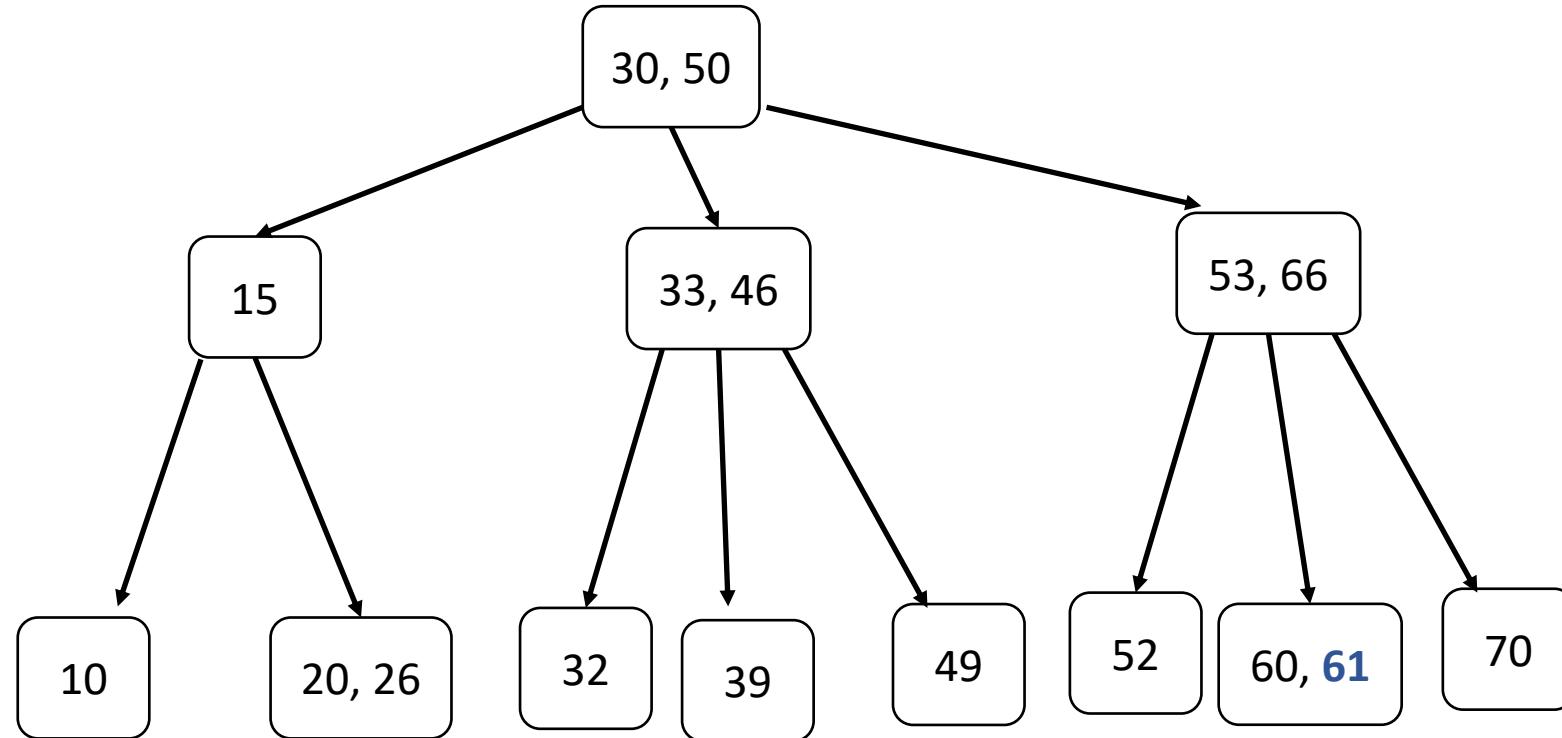
- Insertions in 2-3 trees might transform a 2-node into a 3-node, or blow up a temporary 4-node into 3 2-nodes.
- Deletions, on the other hand, have the potential of not just making 3-nodes into 2-nodes, but making 2-nodes into ... ?
- We'll see what the result of that will be soon.
- Key point: Since deleting an inner node always ends up being a deletion at the leaf level, we only care about leaf deletions.

# Deletion



# Deletion

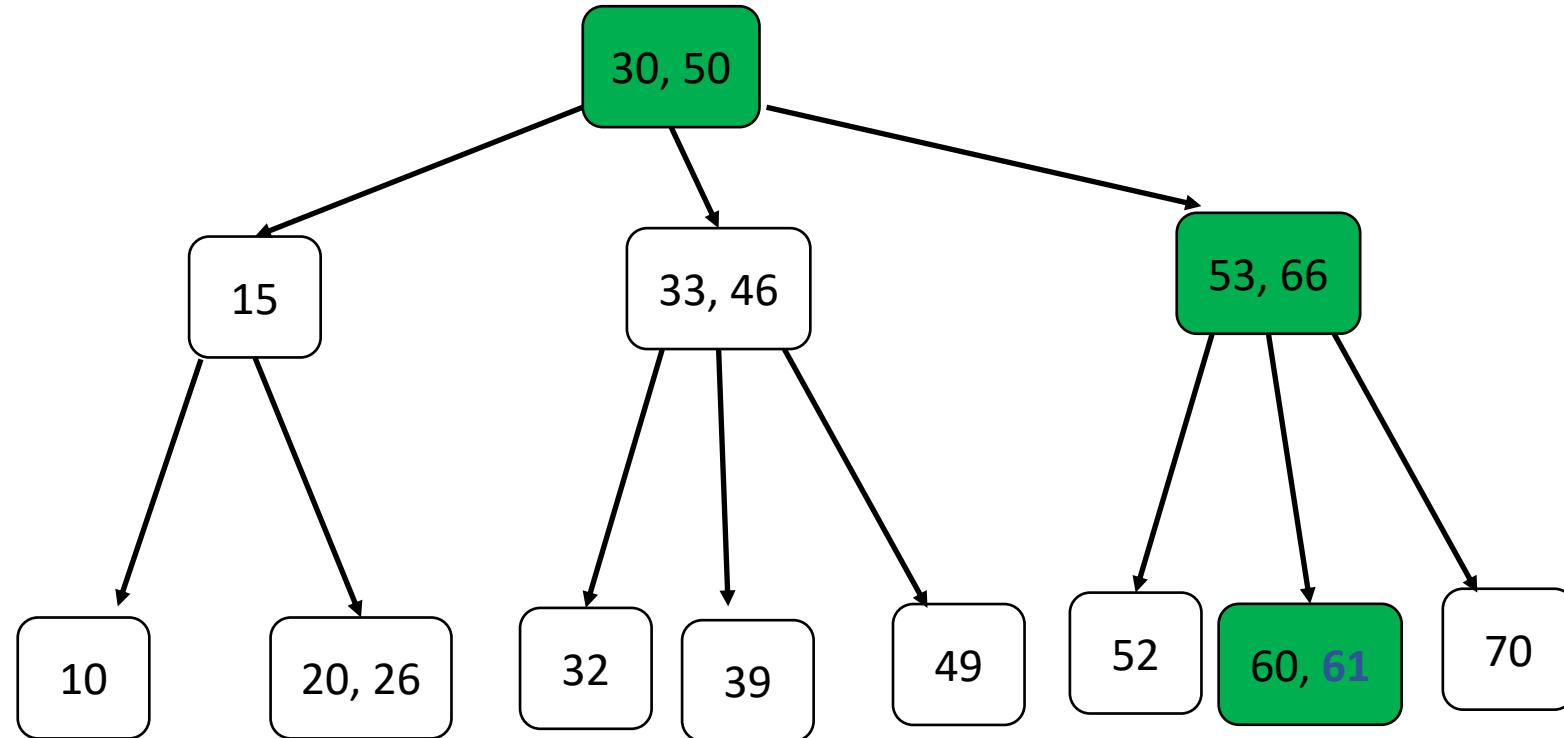
Delete 61



# Deletion

Delete 61

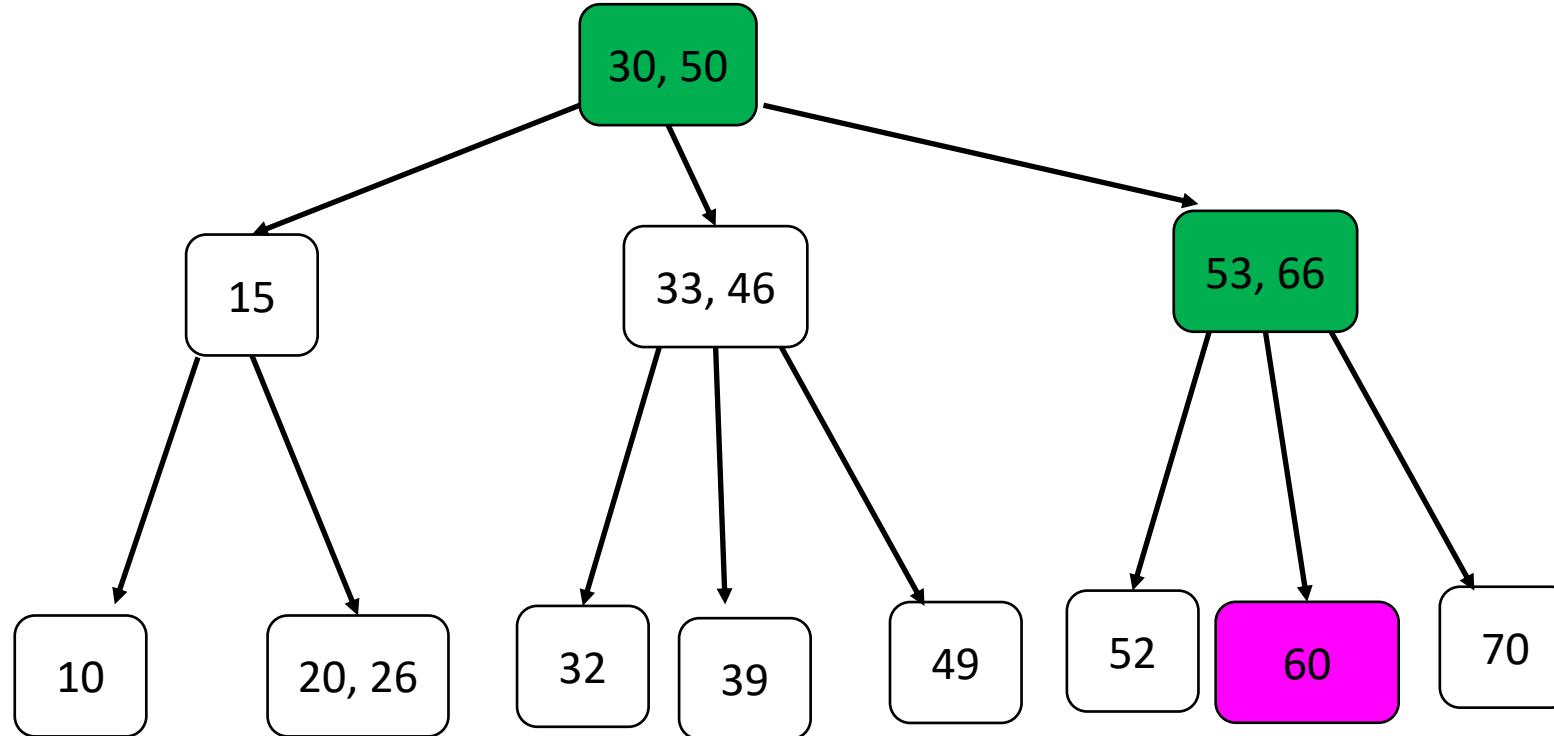
1. Search



# Deletion

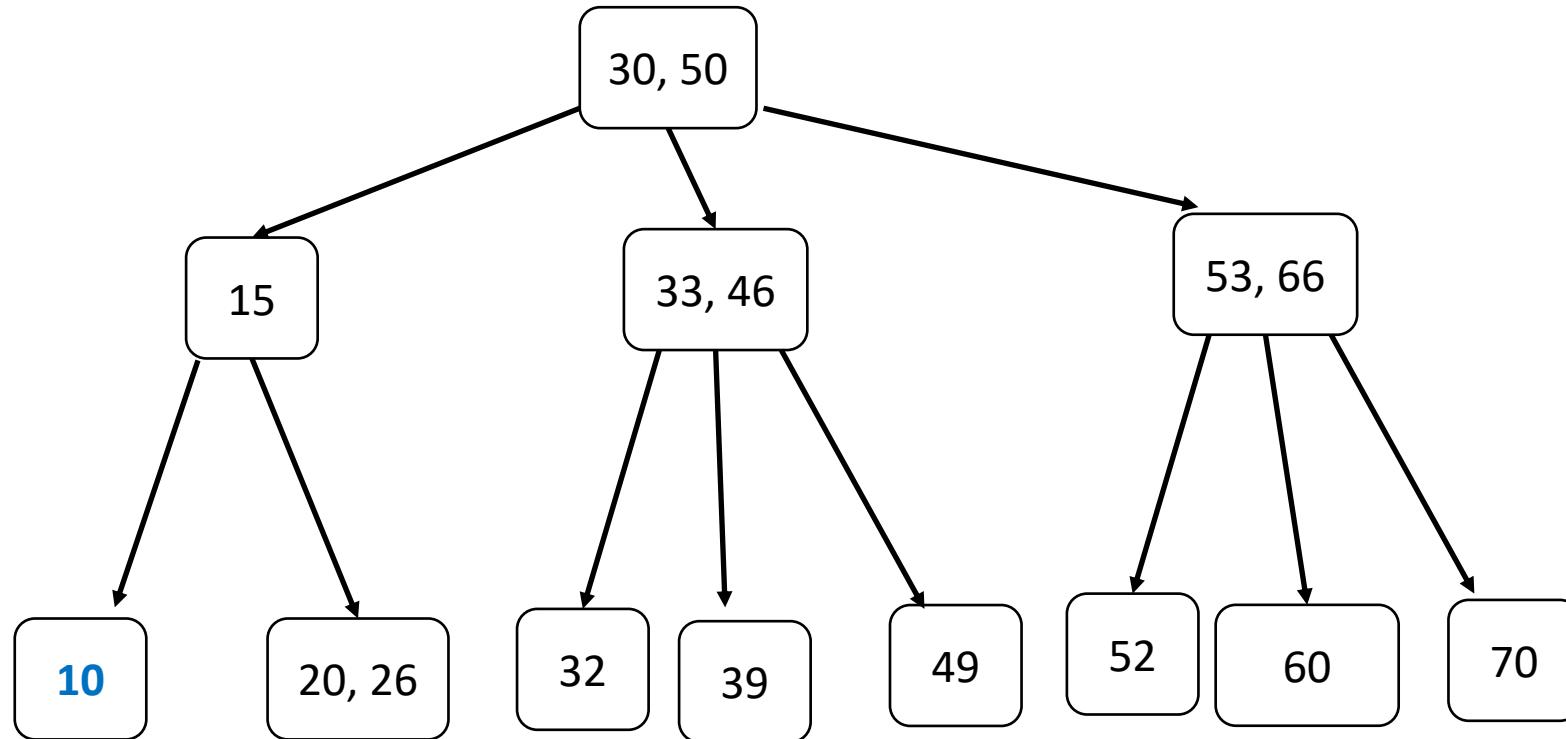
## Delete 61

1. Search
2. Since node that contains 61 is a 3-node, we just remove 61, make it a 2-node and we're done! 😊



# Deletion

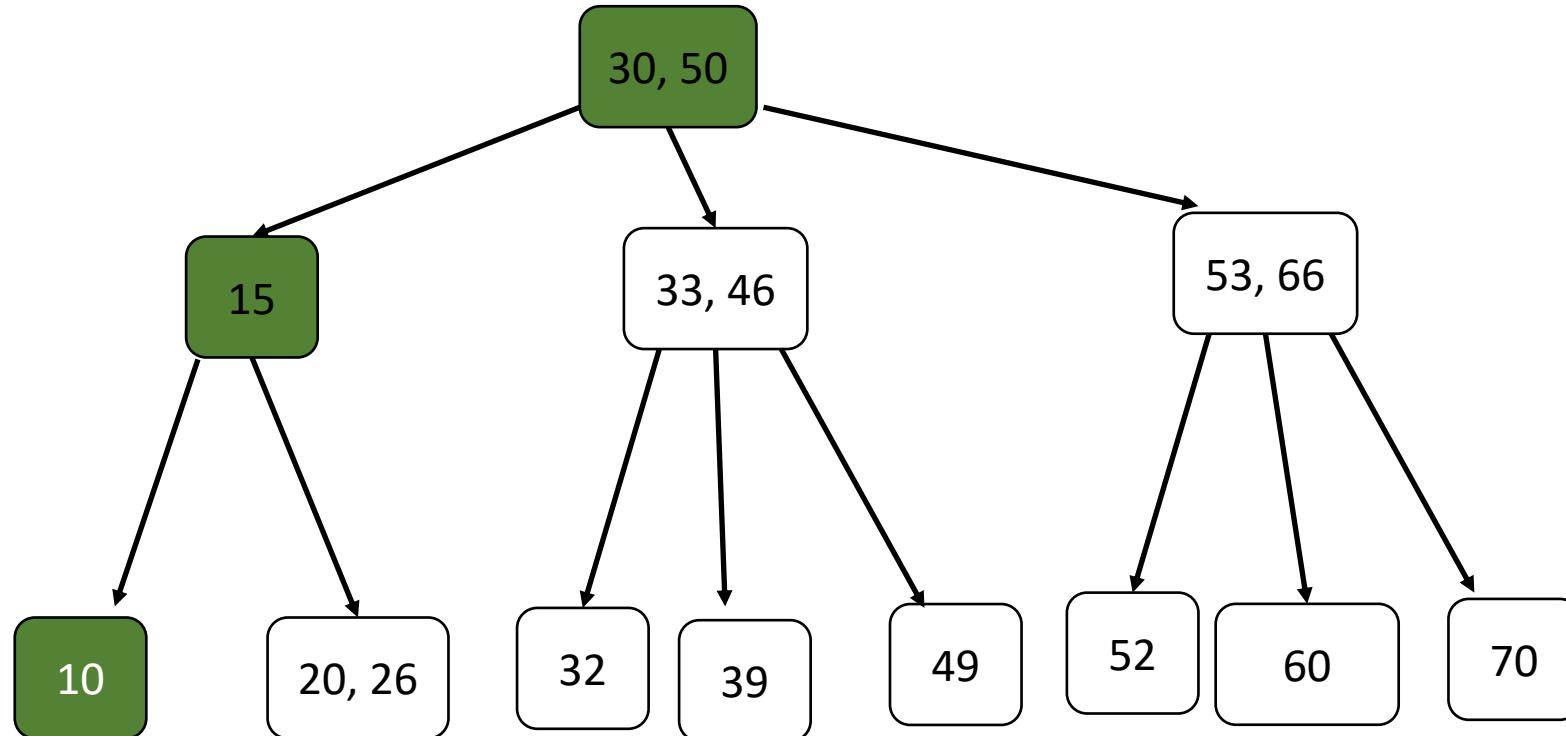
Delete 10



# Deletion

Delete 10

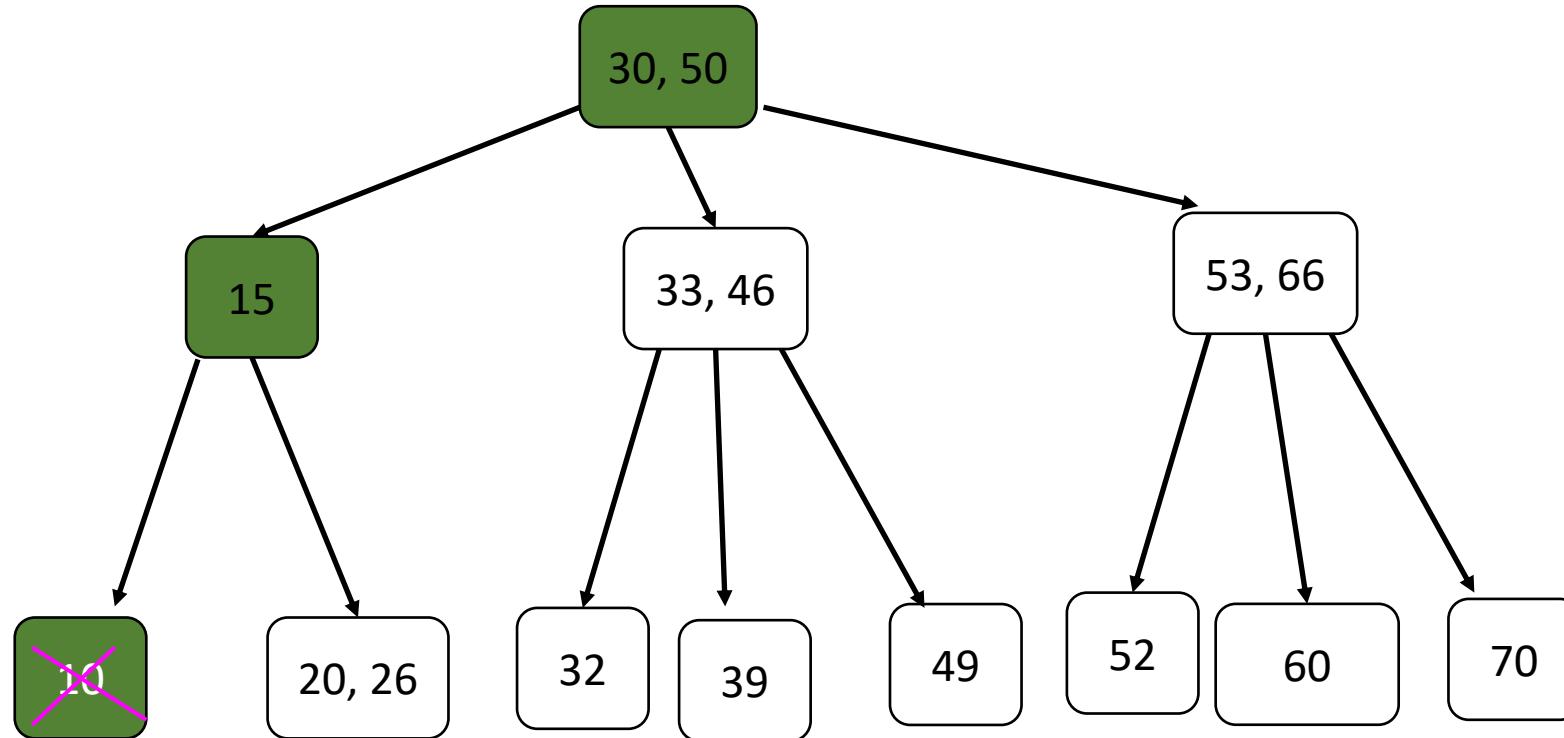
1. Search



# Deletion

## Delete 10

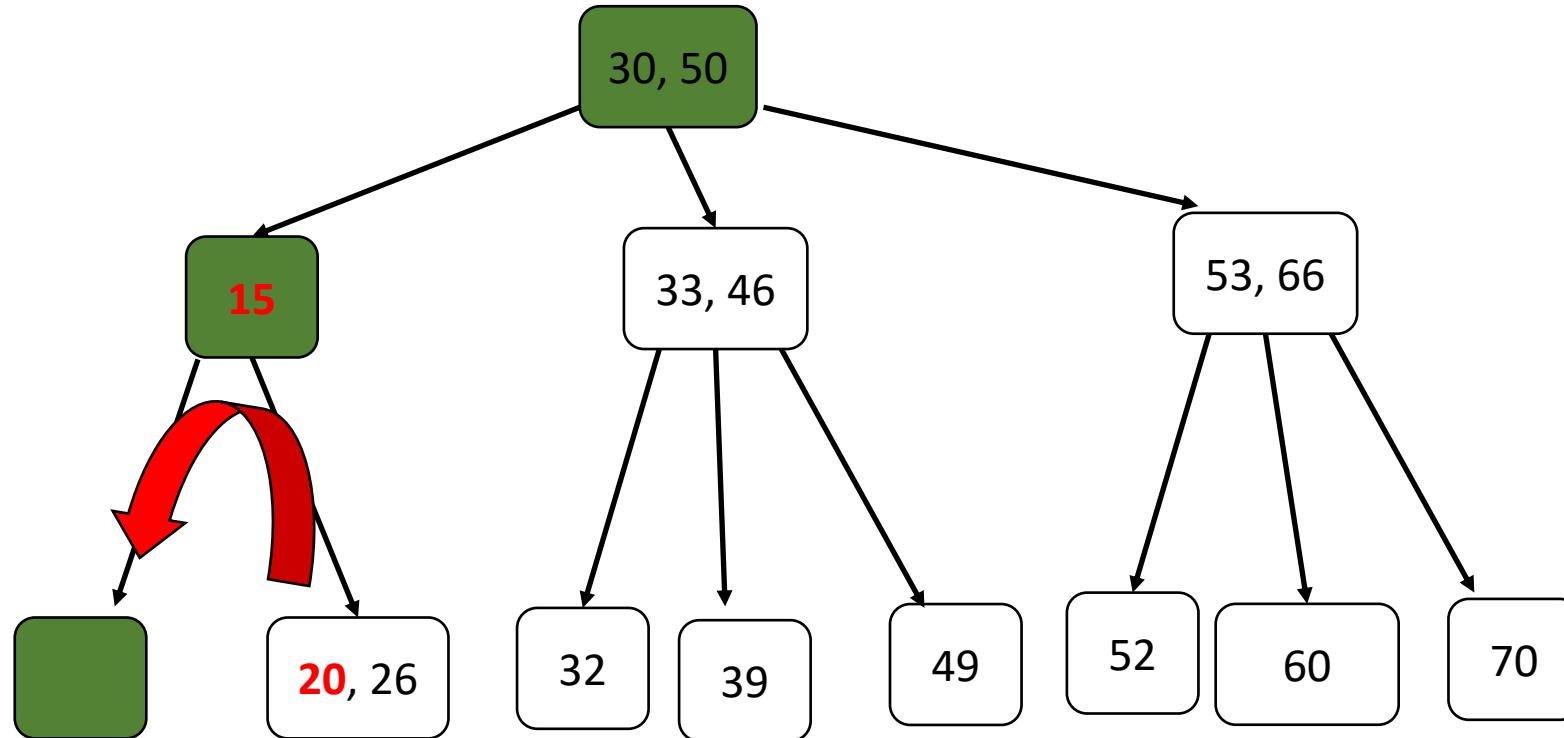
1. Search
2. Deleting 10 leads to a node underflow... 😞



# Deletion

## Delete 10

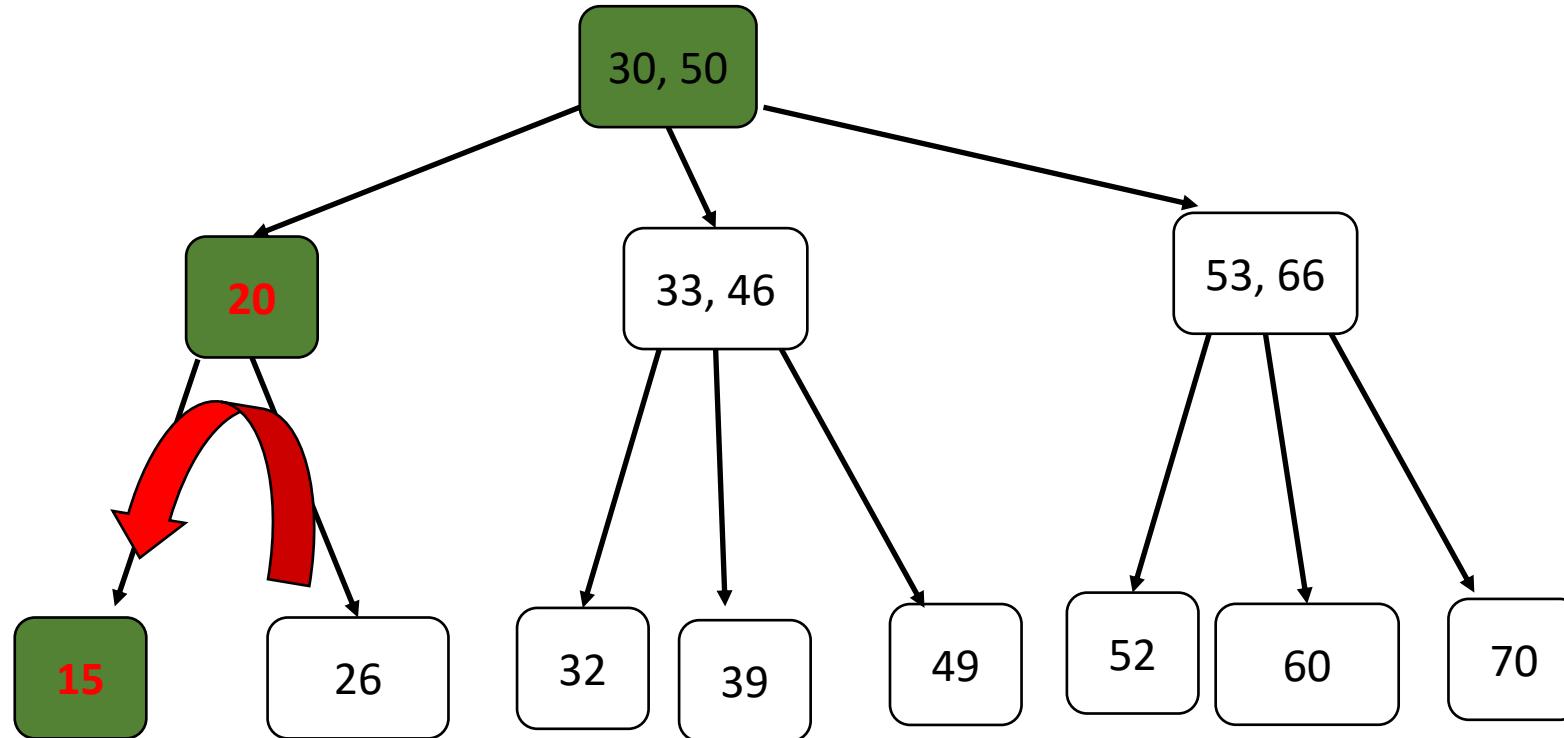
1. Search
2. Deleting 10 leads to a node underflow  
w... 😞
3. The solution: rotate a key from our sibling! 😊



# Deletion

## Delete 10

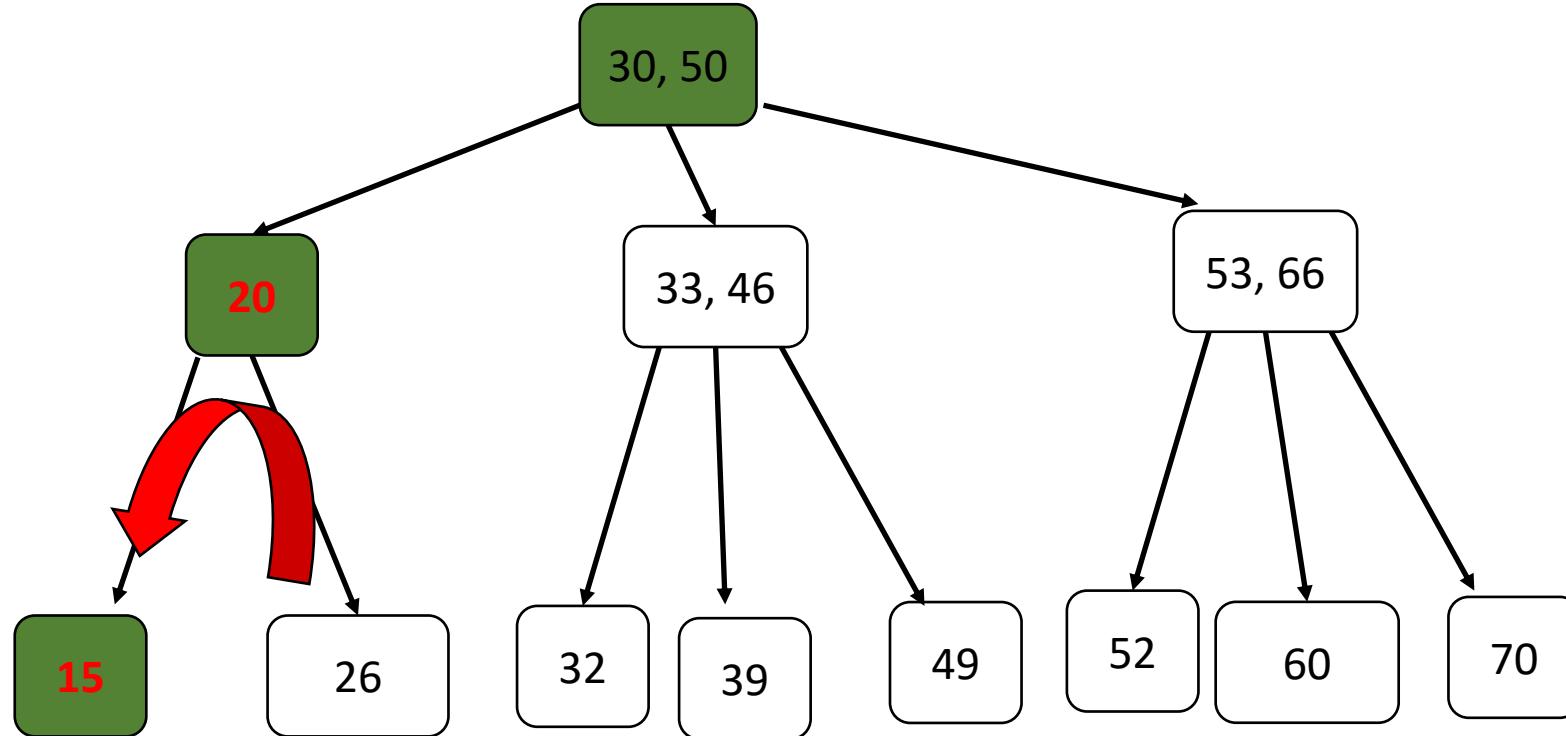
1. Search
2. Deleting 10 leads to a node underflow  
w... 😞
3. The solution: rotate a key from our sibling! 😊



# Deletion

## Delete 10

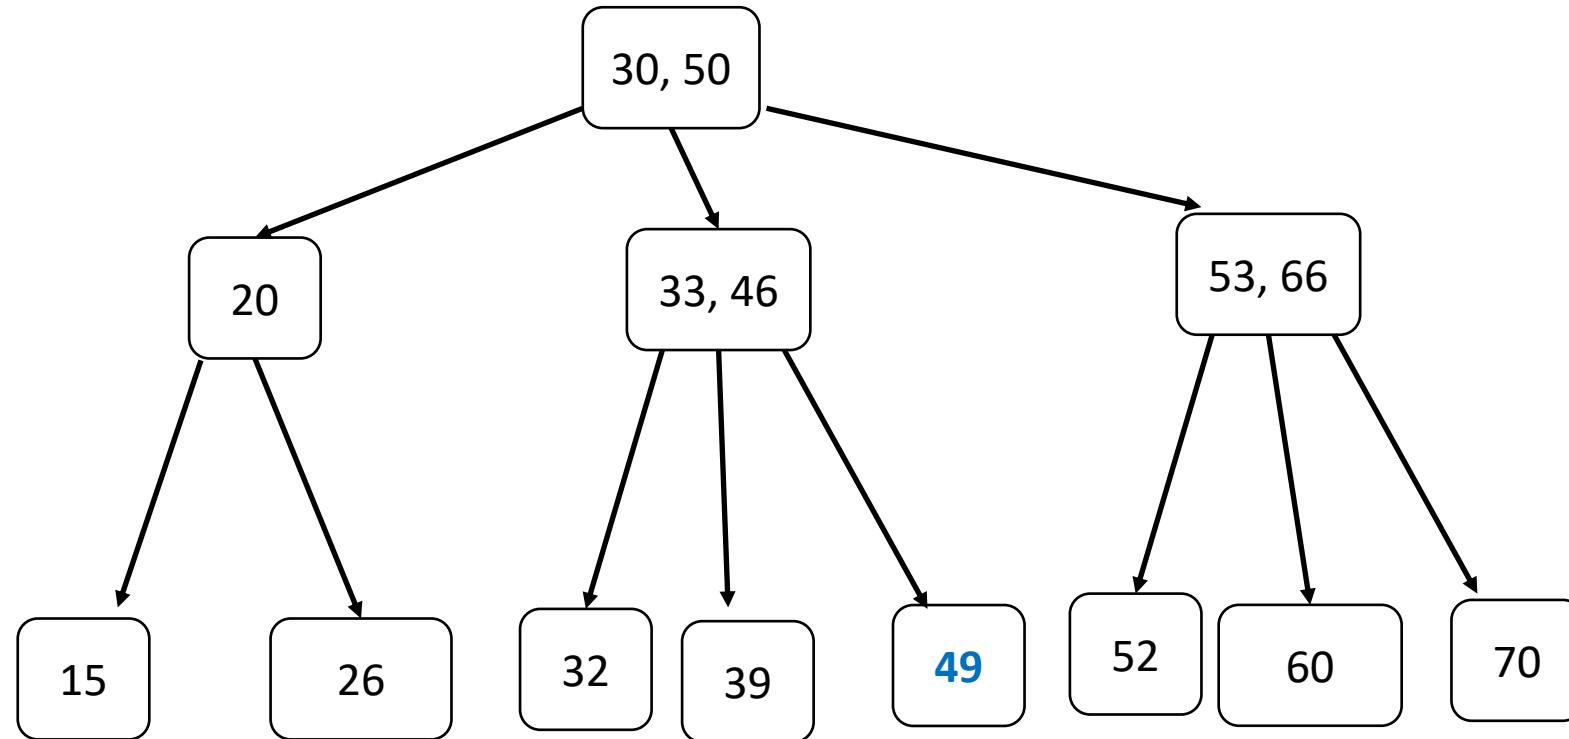
1. Search
2. Deleting 10 leads to a node underflow  
w... 😞
3. The solution: rotate a key from our sibling! 😊



These operations are called “key rotations”, to disambiguate them from the node rotations that occur in AVL and Splay trees.

# Deletion

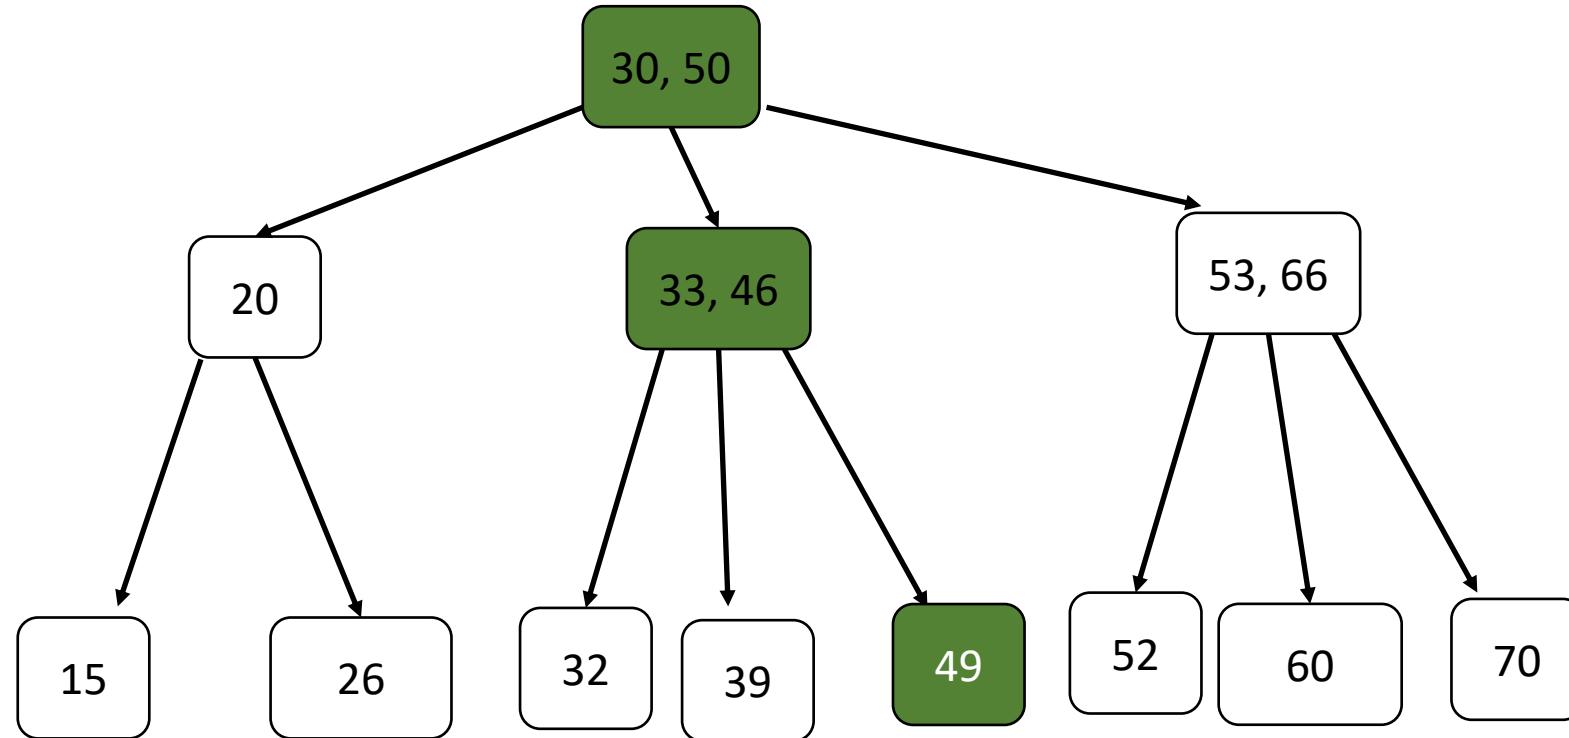
Delete 49



# Deletion

Delete 49

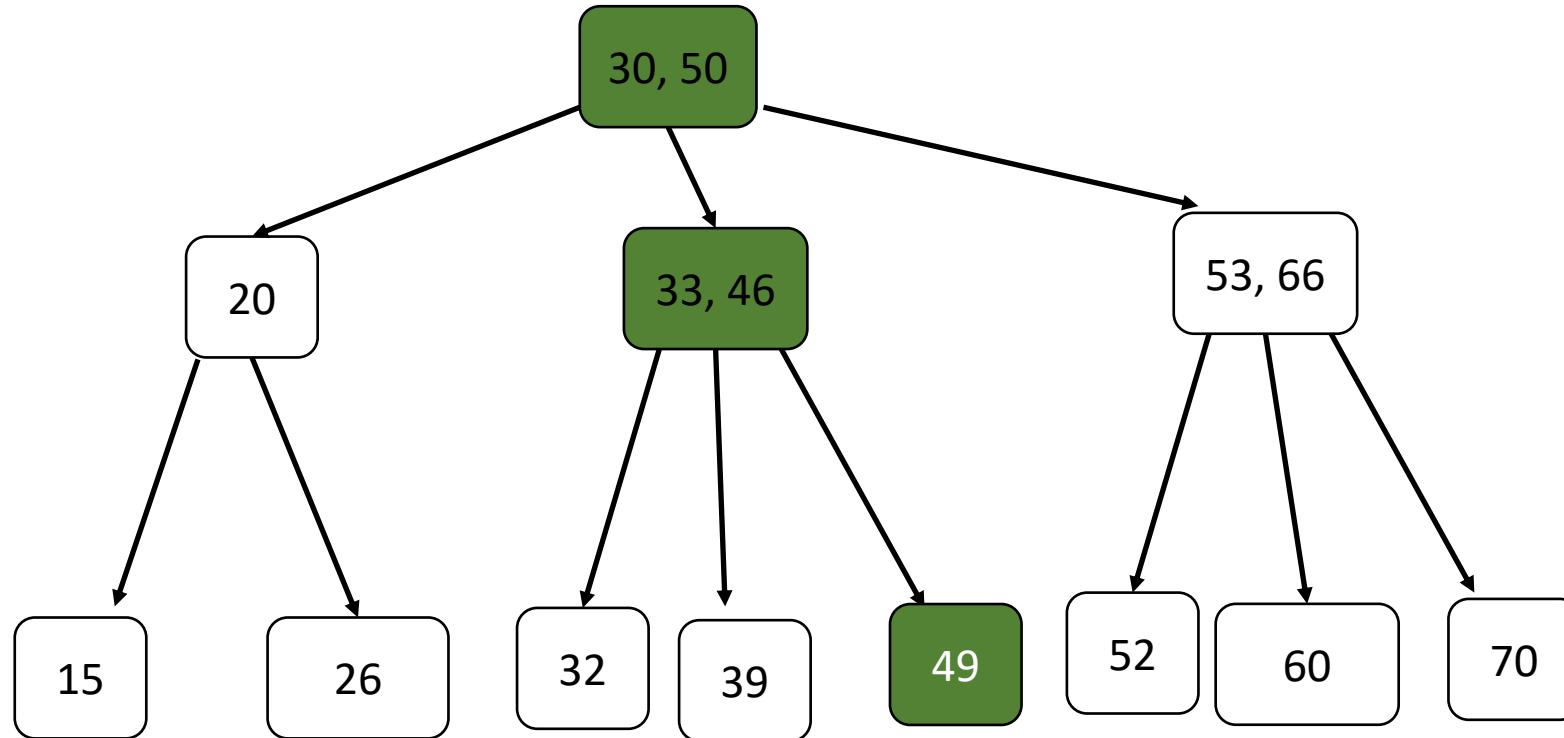
1. Search



# Deletion

Delete 49

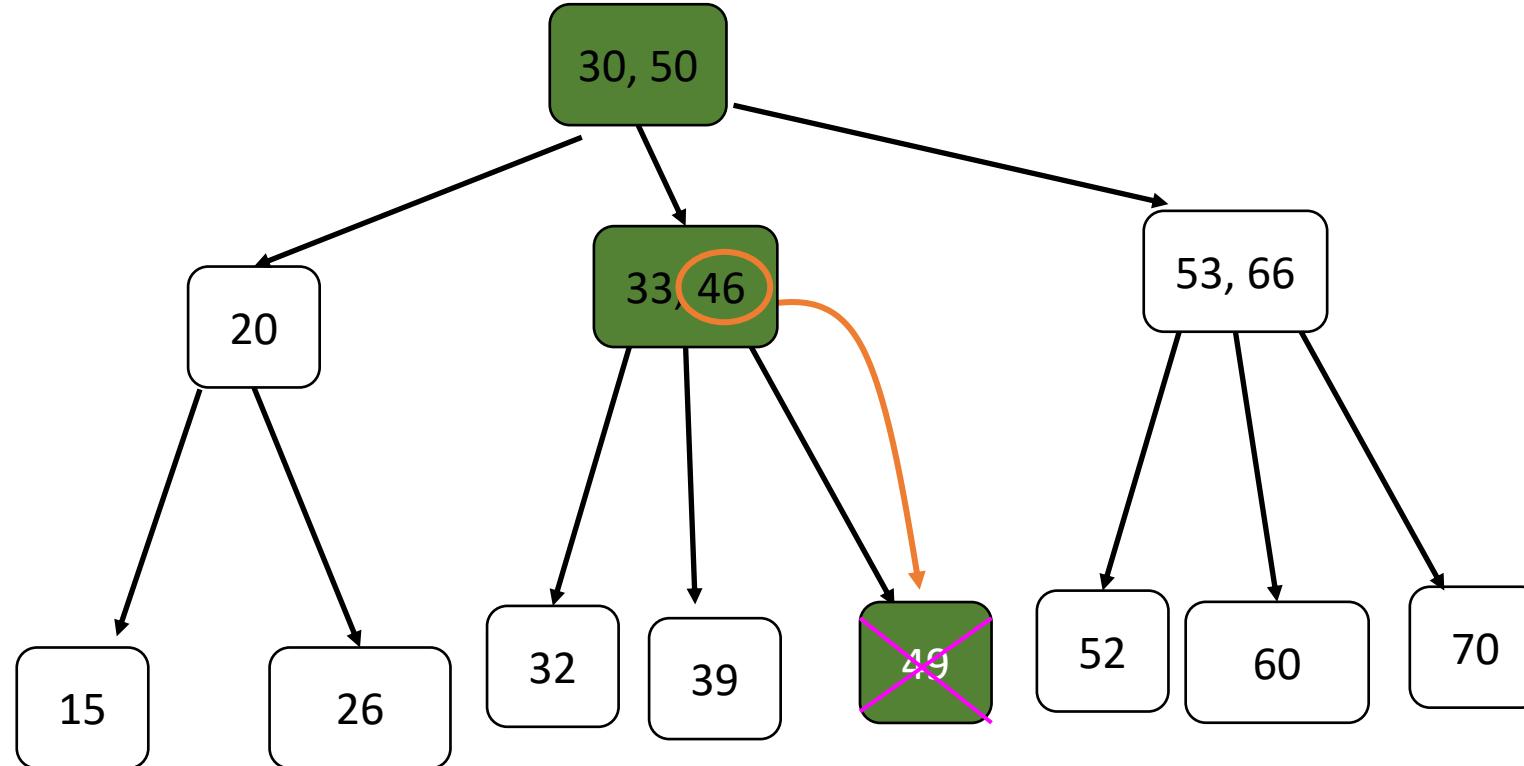
1. Search
2. 49 is contained in a 2-node, and both our siblings are 2-nodes as well...  
:(



# Deletion

## Delete 49

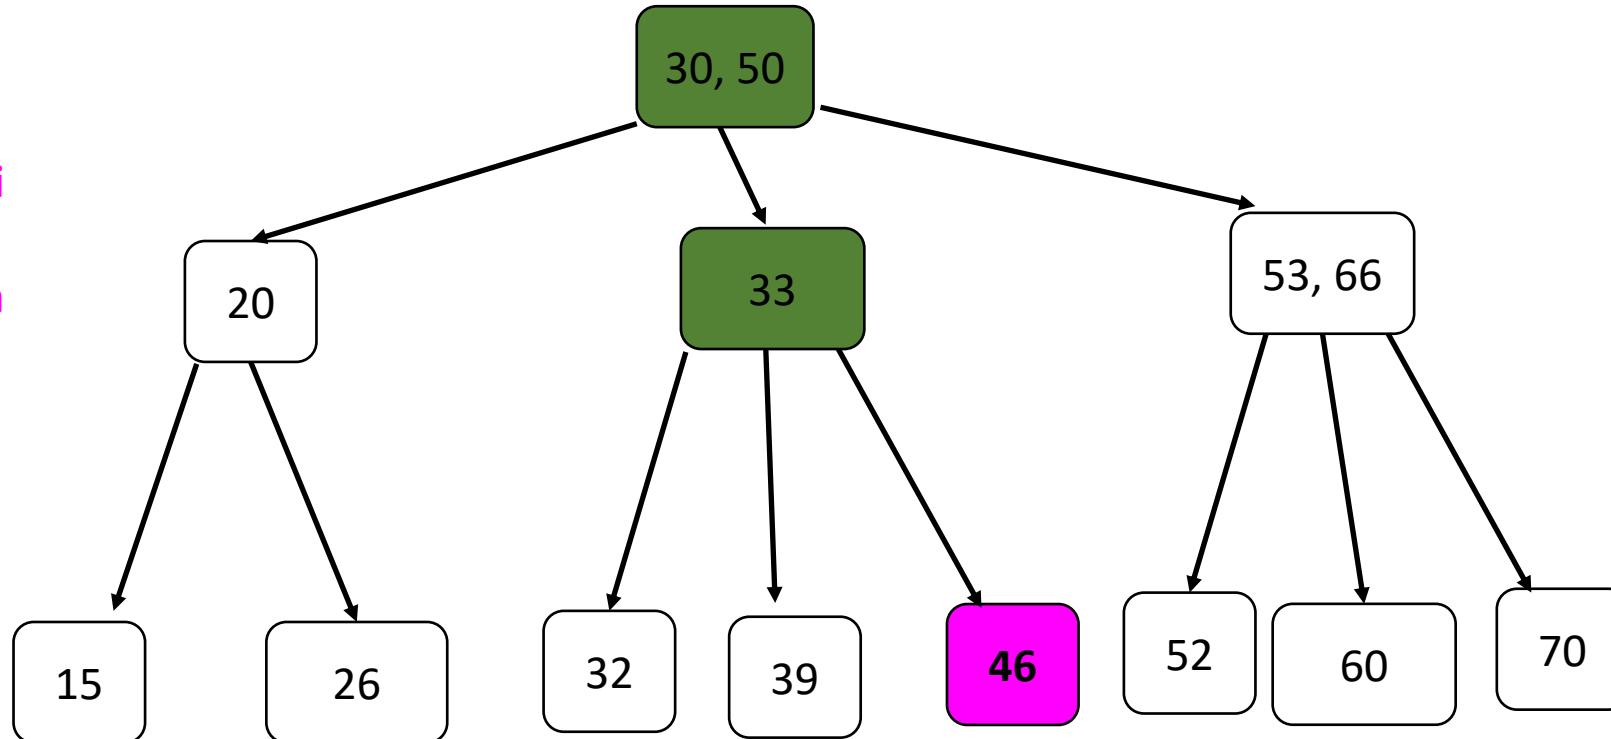
1. Search
2. 49 is contained in a 2-node, and both our siblings are 2-nodes as well...  
:(
3. Solution: Grab a key from the parent!



# Deletion

## Delete 49

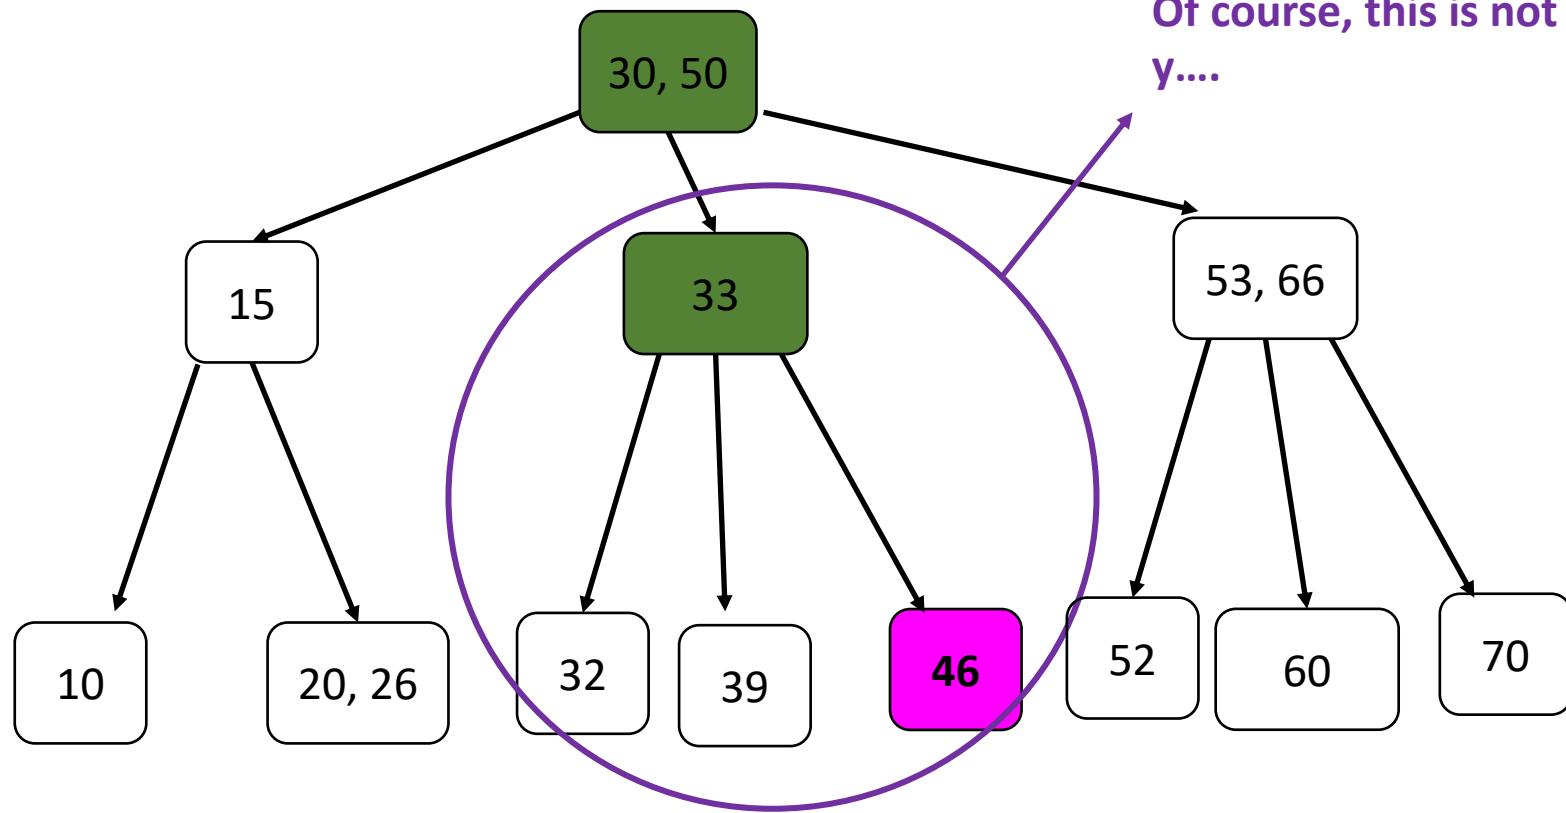
1. Search
2. 49 is contained in a 2-node, so deleting the node would create an imbalance... 😞
3. Solution: Grab a key from the parent!



# Deletion

## Delete 49

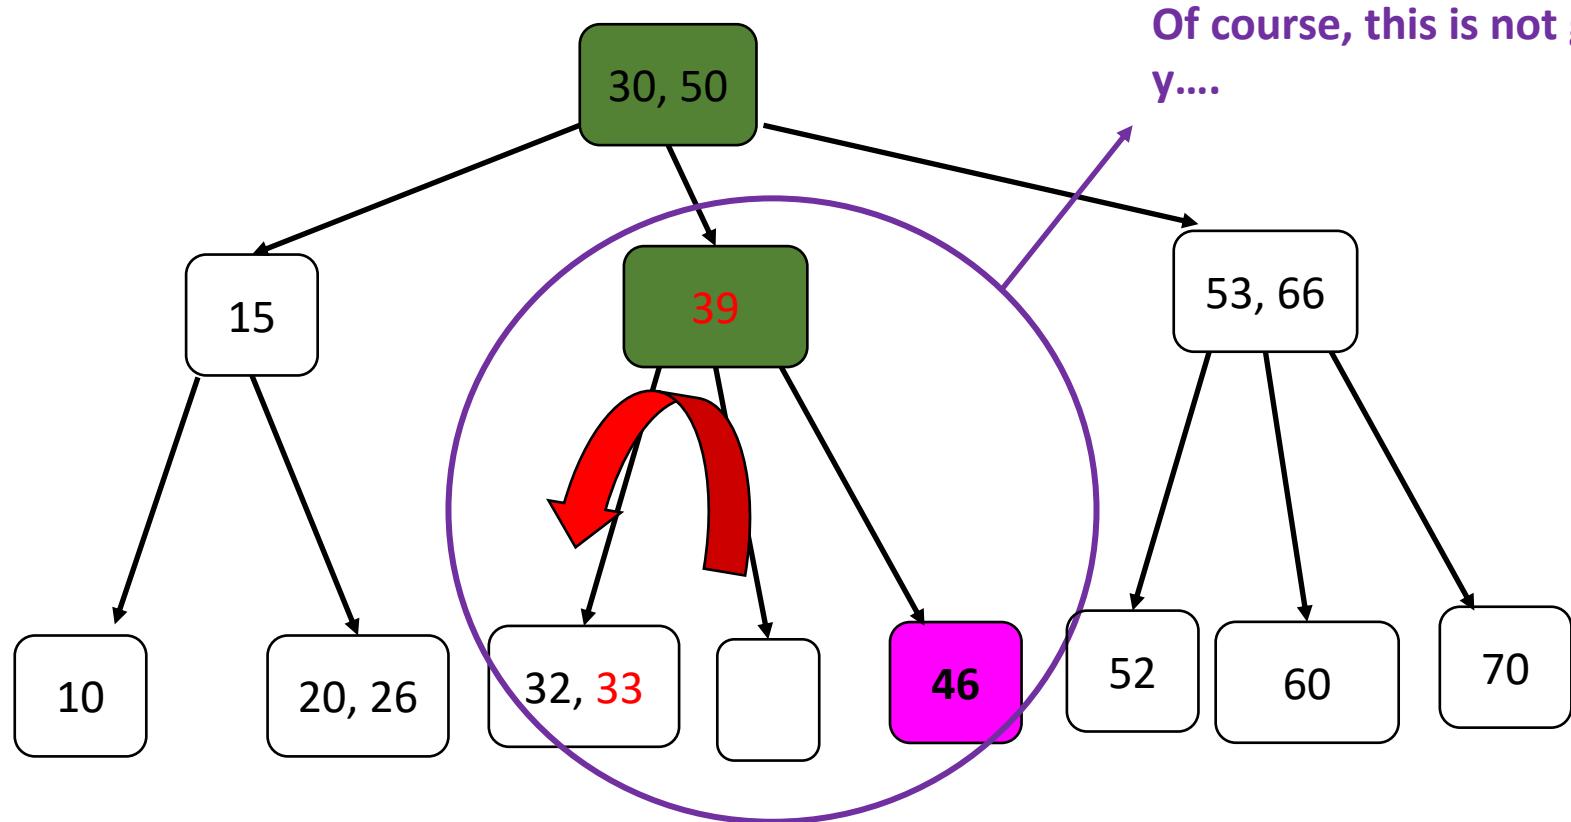
1. Search
2. 49 is contained in a 2-node, so deleting the node would create an imbalance... 😞
3. Solution: Grab a key from the parent!



# Deletion

## Delete 49

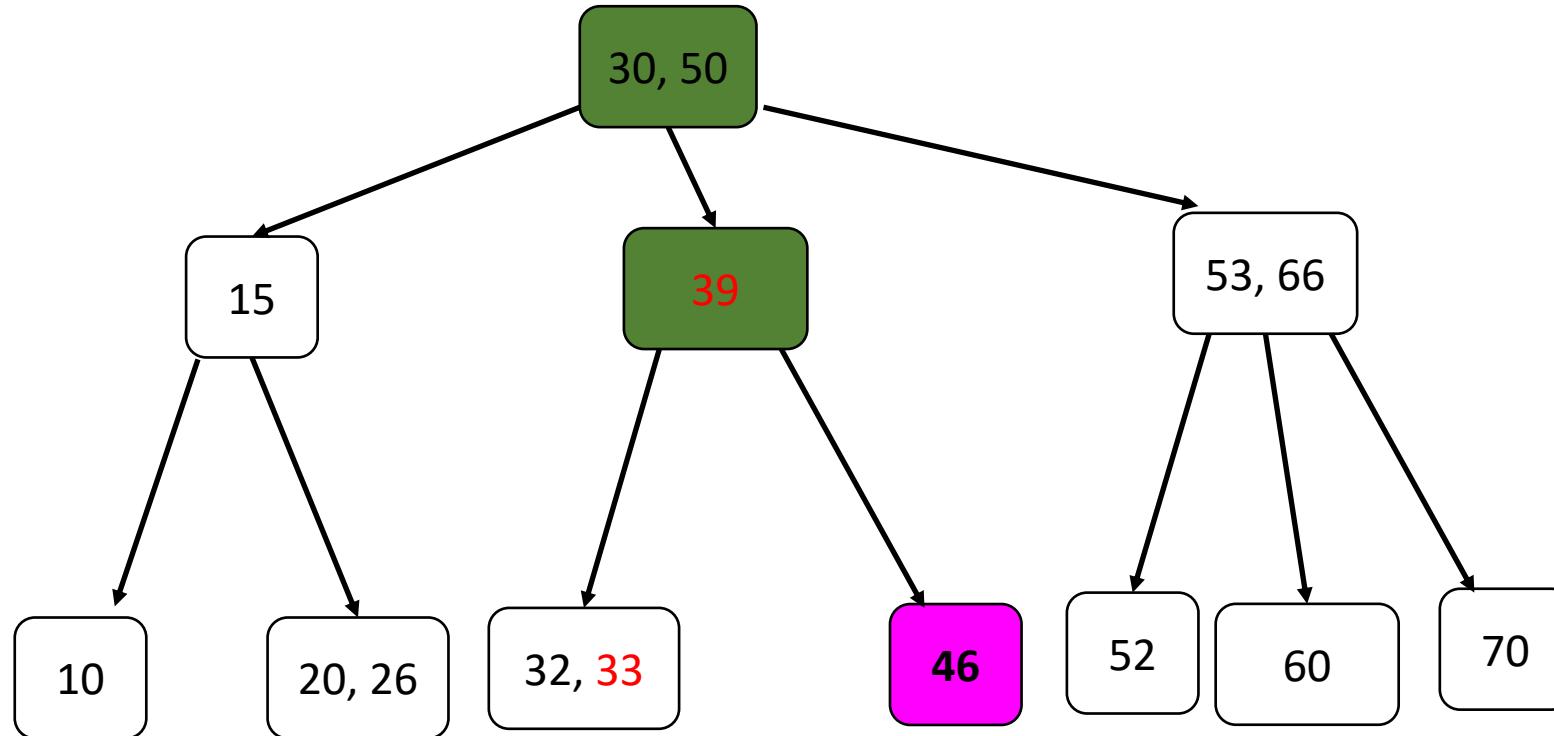
1. Search
2. 49 is contained in a 2-node, so deleting the node would create an imbalance... 😞
3. Solution: Grab a key from the parent!
4. To make the parent be a well-defined 2-node, rotate the middle key to the left!



# Deletion

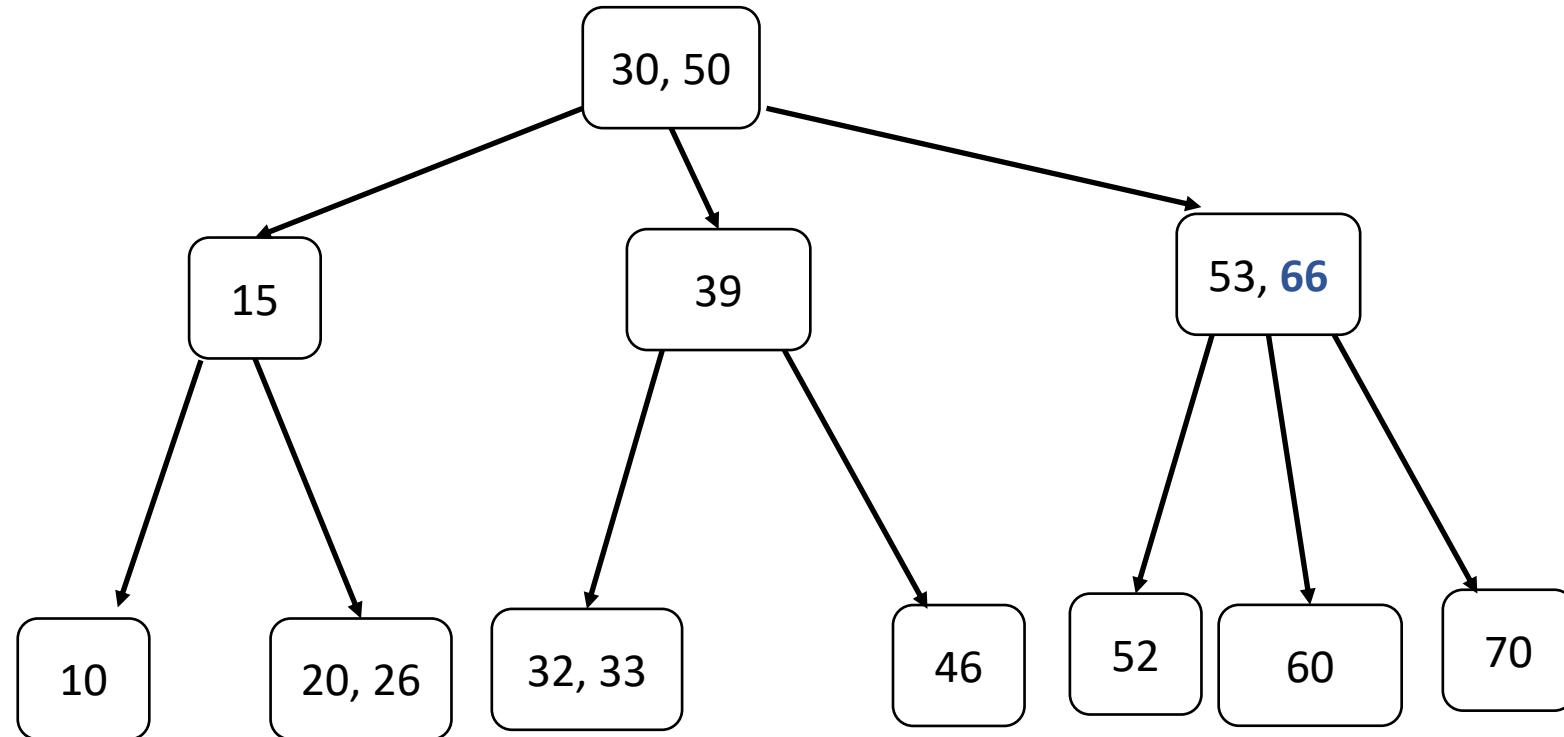
## Delete 49

1. Search
2. 49 is contained in a 2-node, so deleting the node would create an imbalance... 😞
3. Solution: Grab a key from the parent!
4. To make the parent be a well-defined 2-node, rotate the middle key to the left!



# Deletion

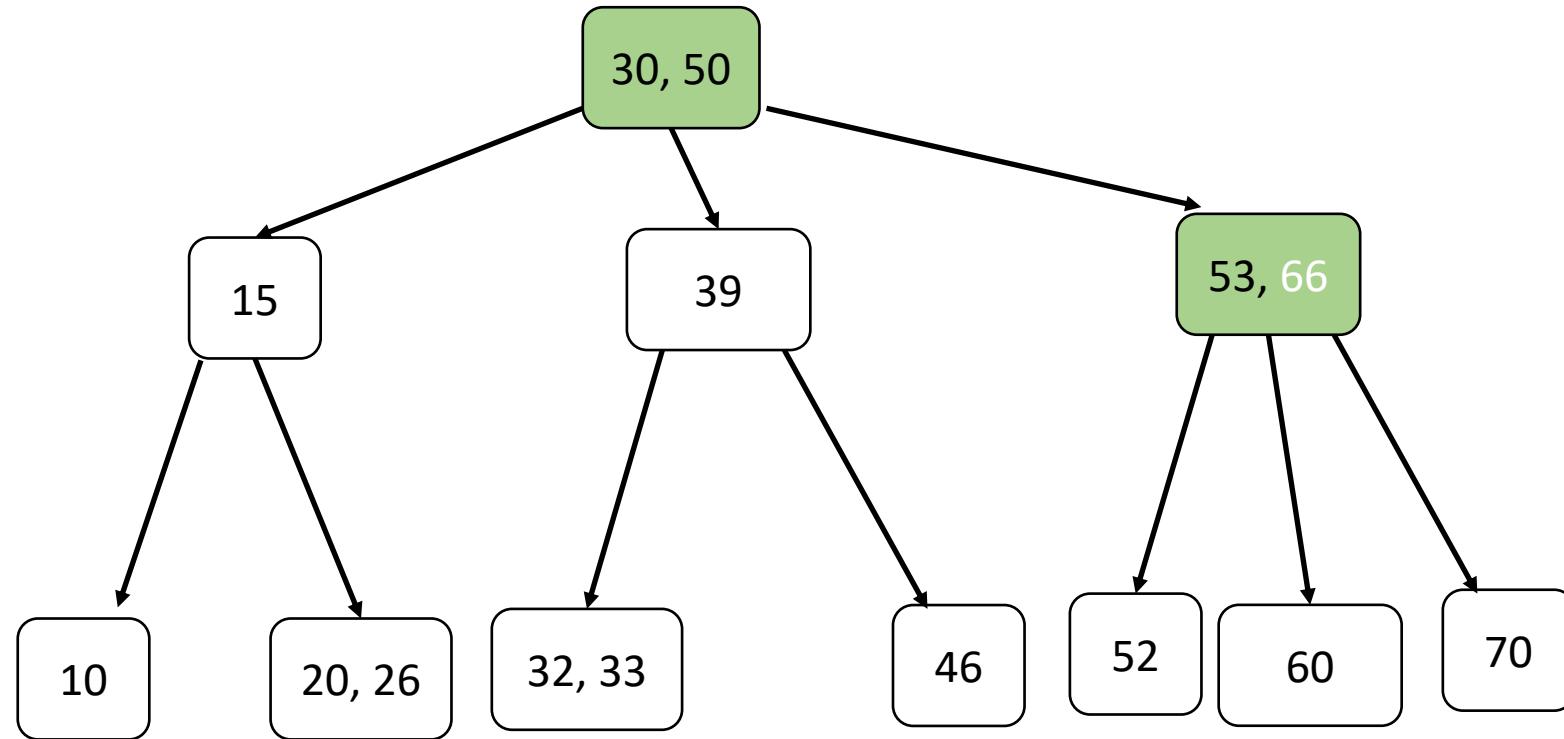
Delete 66



# Deletion

Delete 66

1. Search

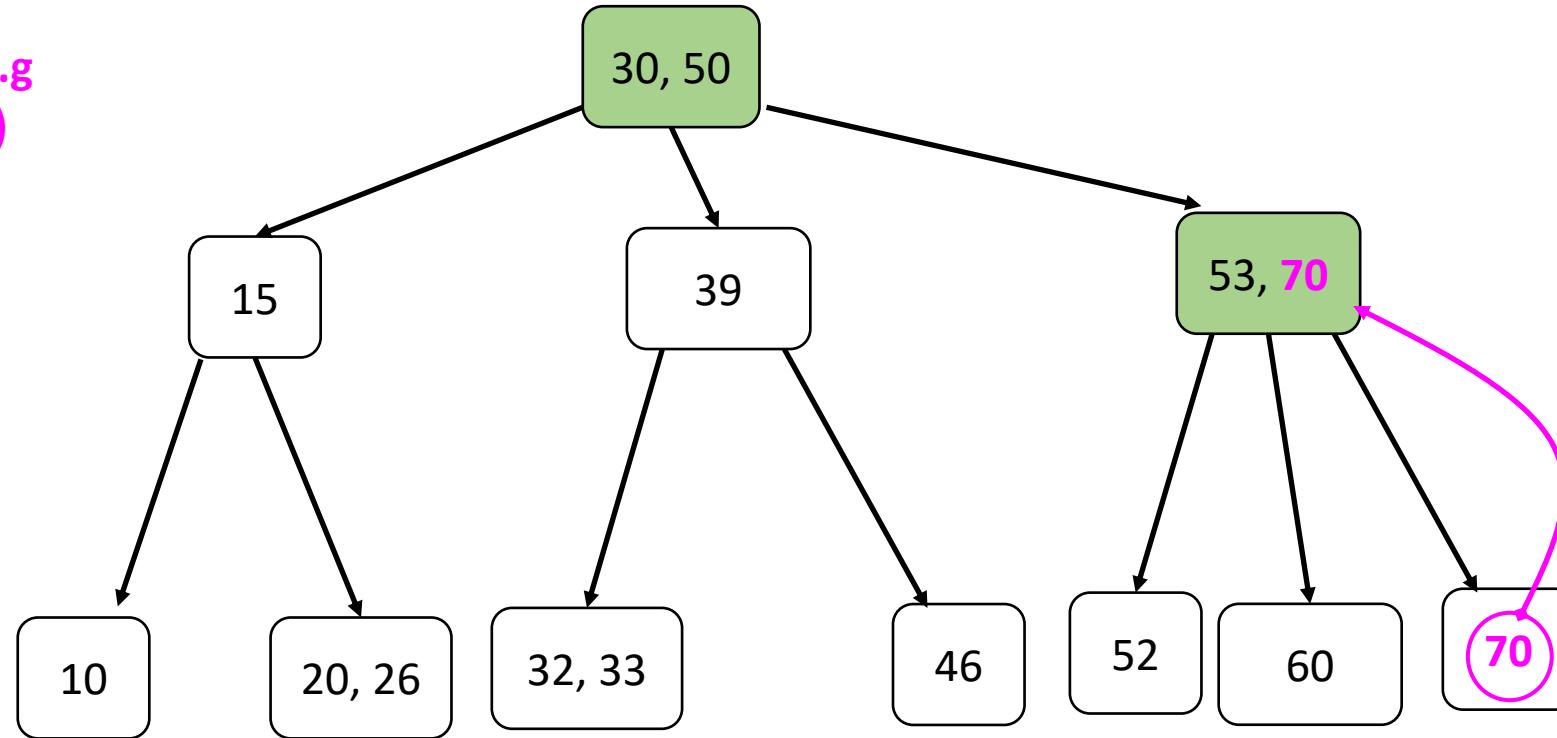


# Deletion

Delete 66

1. Search

2. Have to replace key (e.g  
with inorder successor...)



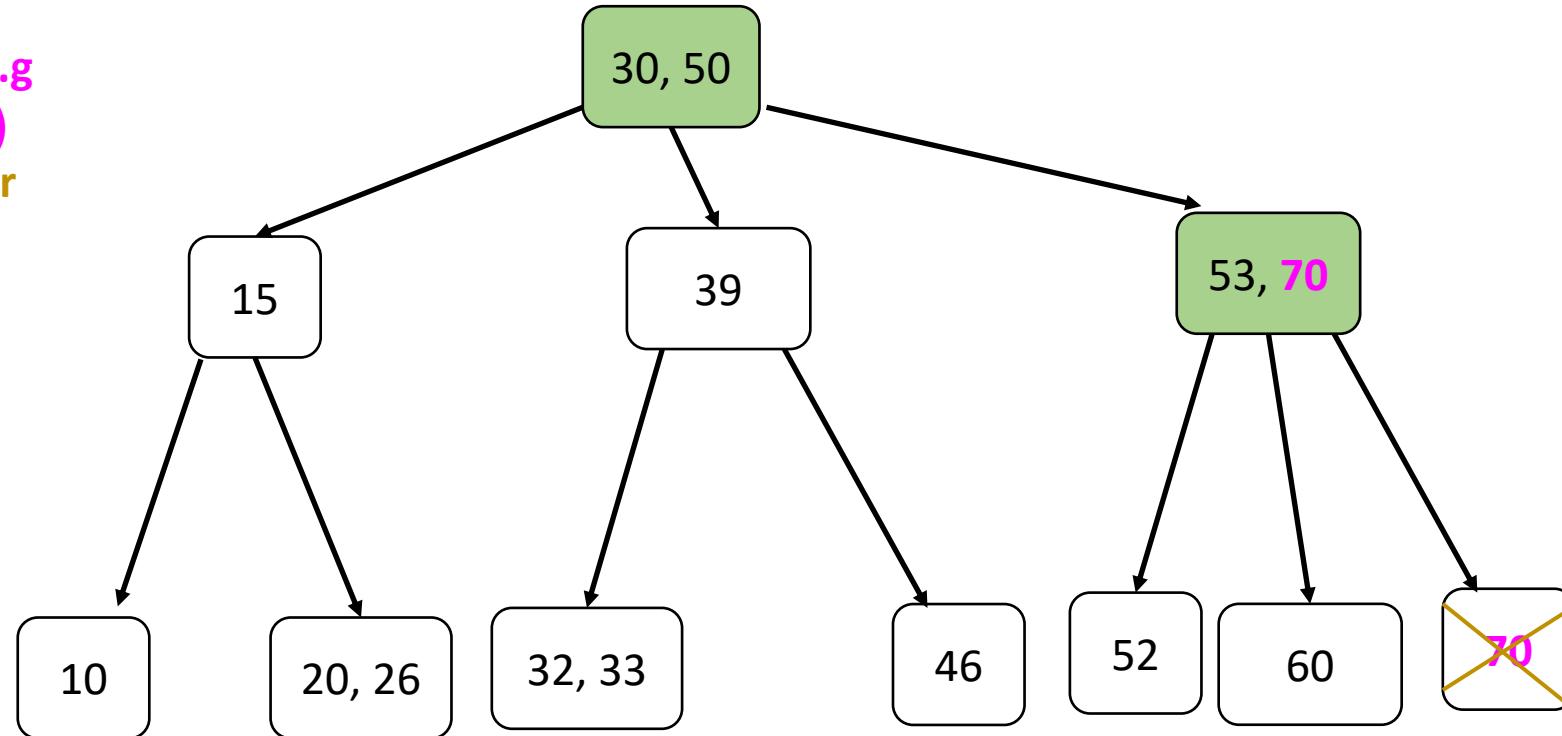
# Deletion

Delete 66

1. Search

2. Have to replace key (e.g  
with inorder successor...)

3. Have to delete the inor  
der succesor



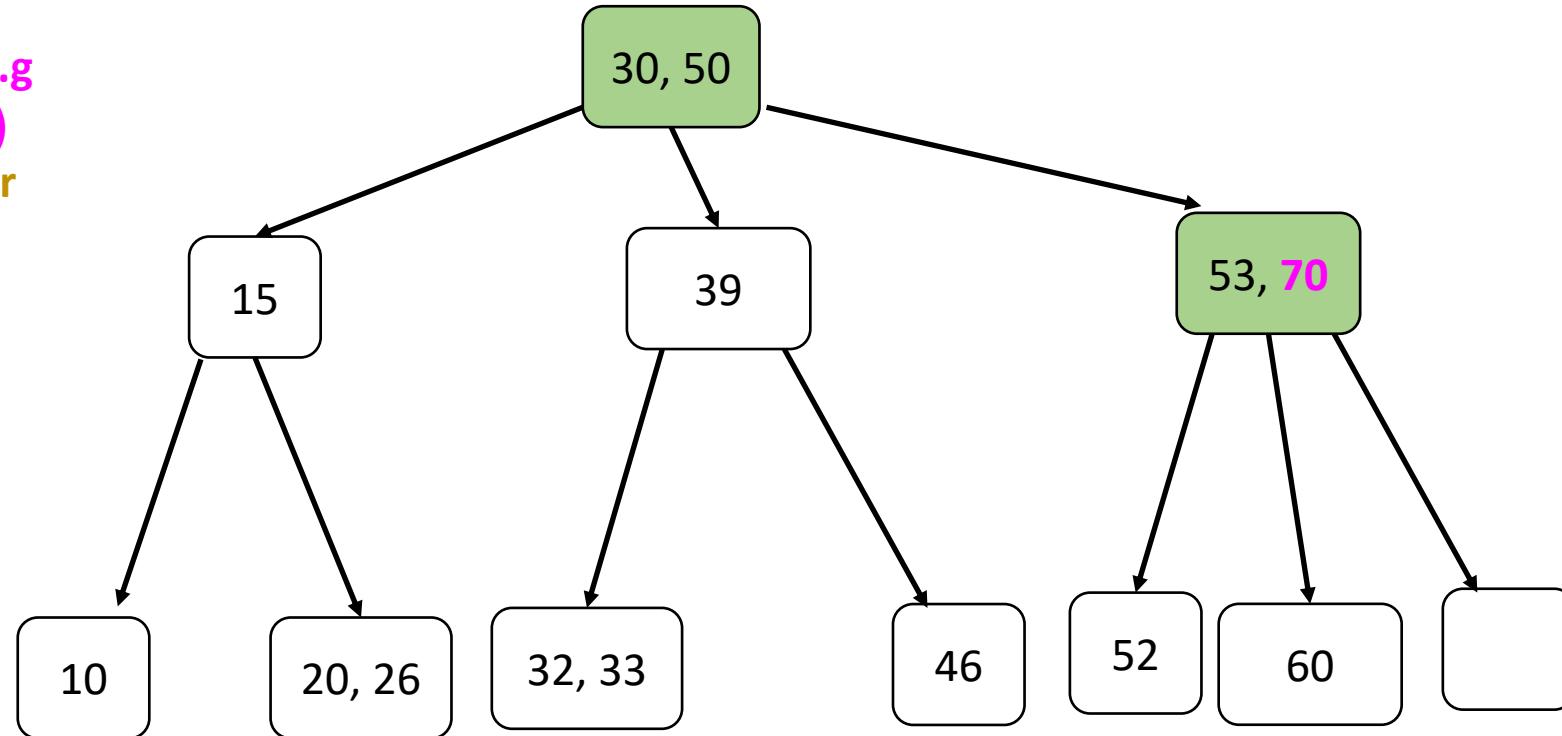
# Deletion

Delete 66

1. Search

2. Have to replace key (e.g  
with inorder successor...)

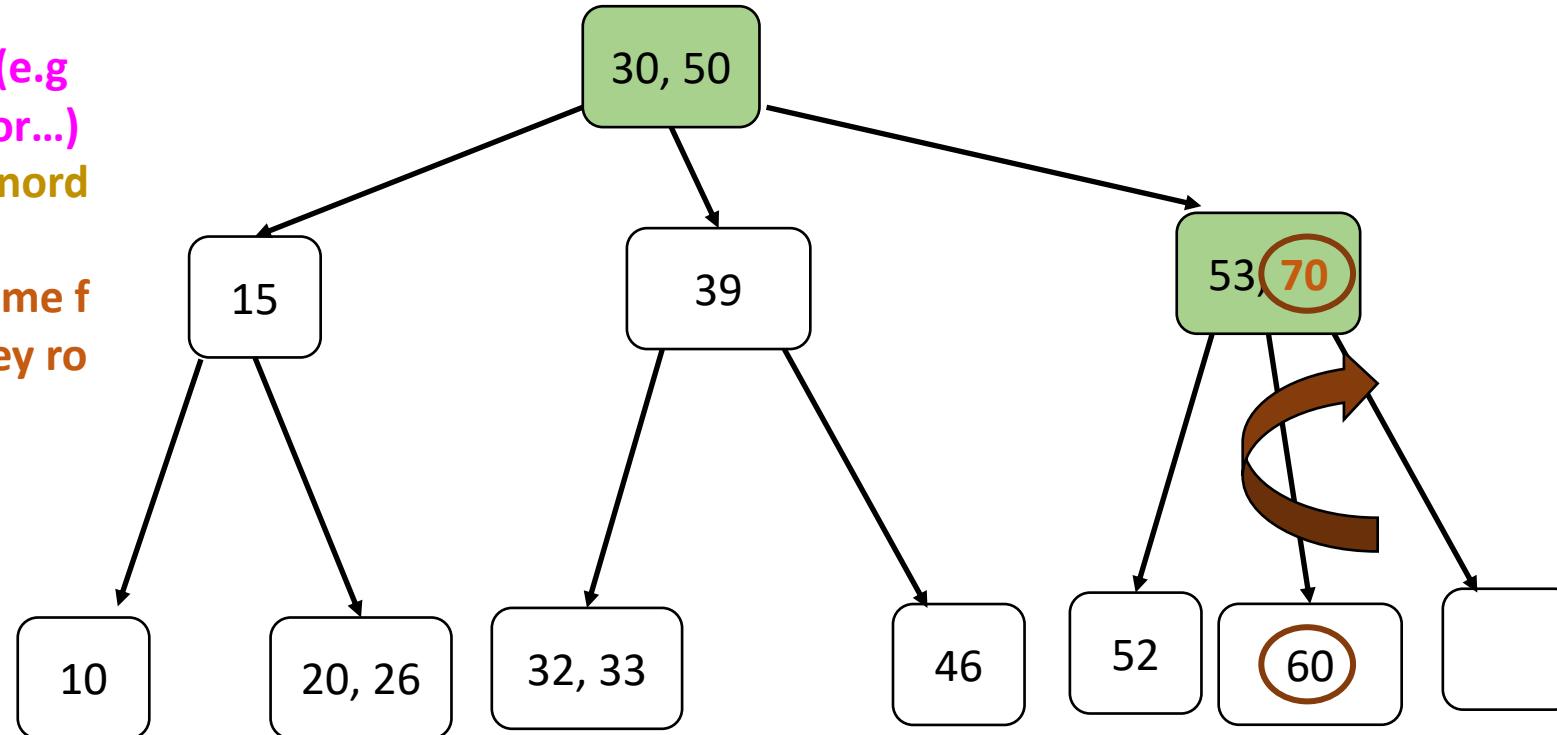
3. Have to delete the inor  
der succesor



# Deletion

Delete 66

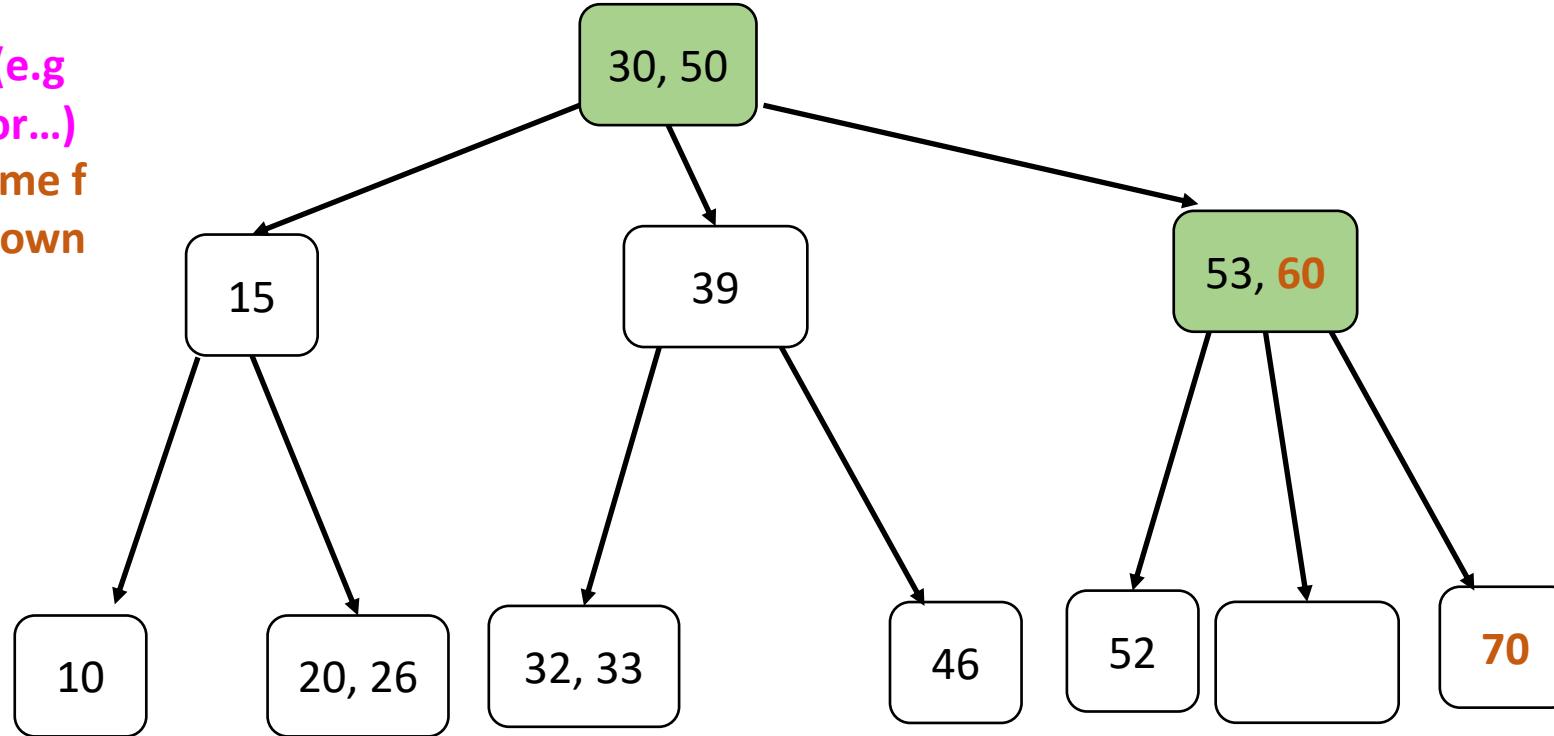
1. Search
2. Have to replace key (e.g with inorder successor...)
3. . Have to delete the inorder successor
4. Use our familiar scheme from before to do a key rotation....



# Deletion

Delete 66

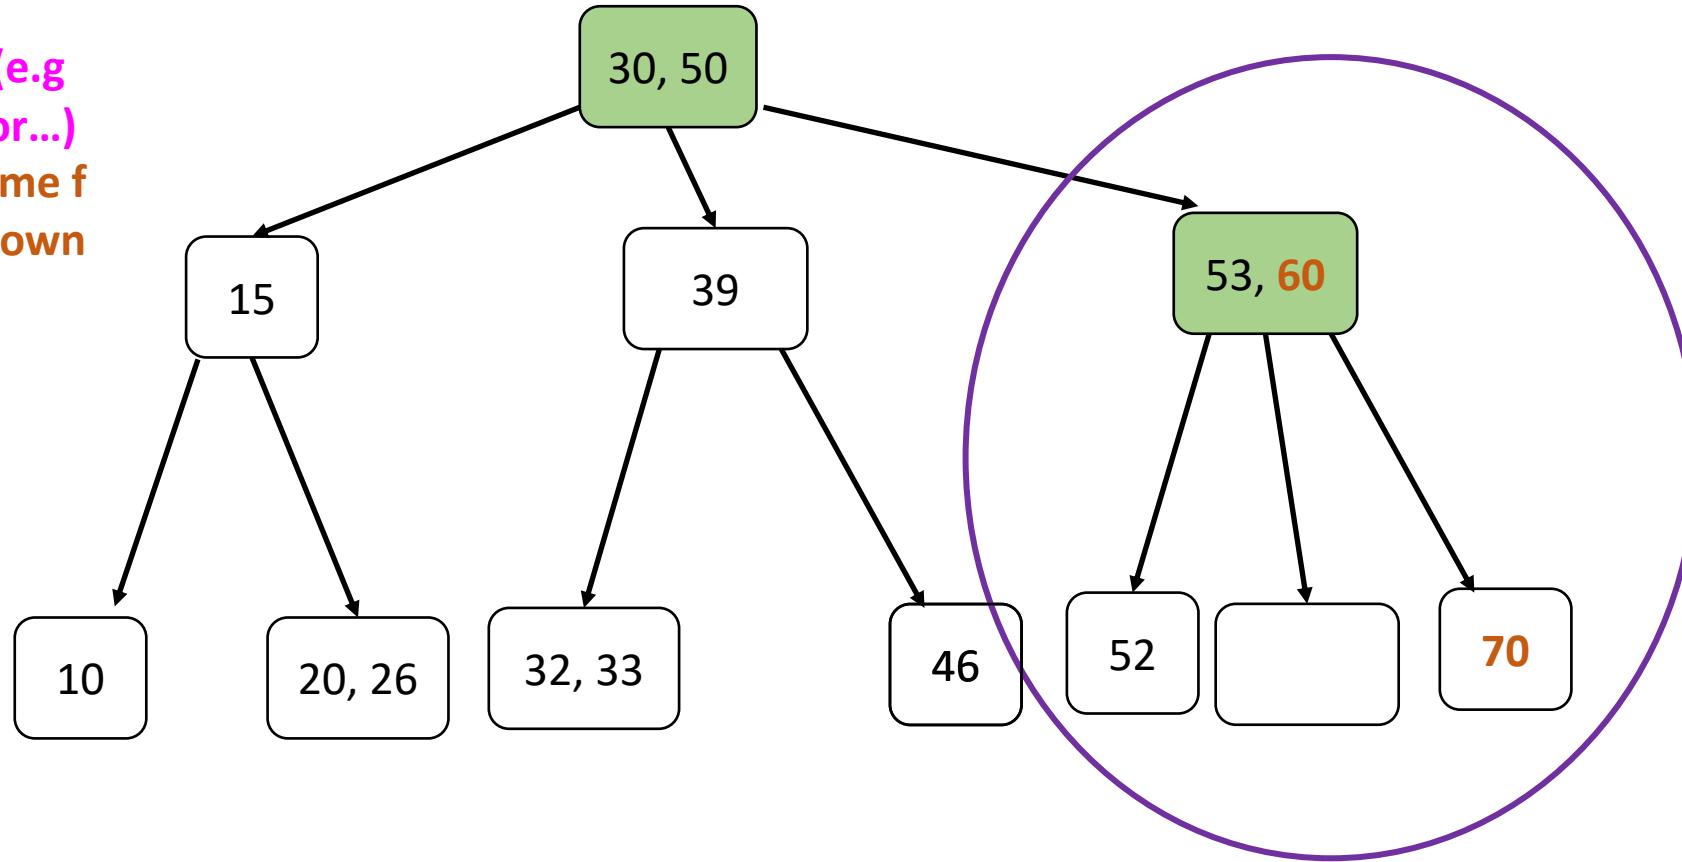
1. Search
2. Have to replace key (e.g. with inorder successor...)
3. Use our familiar scheme from before to push down 70 (again)



# Deletion

Delete 66

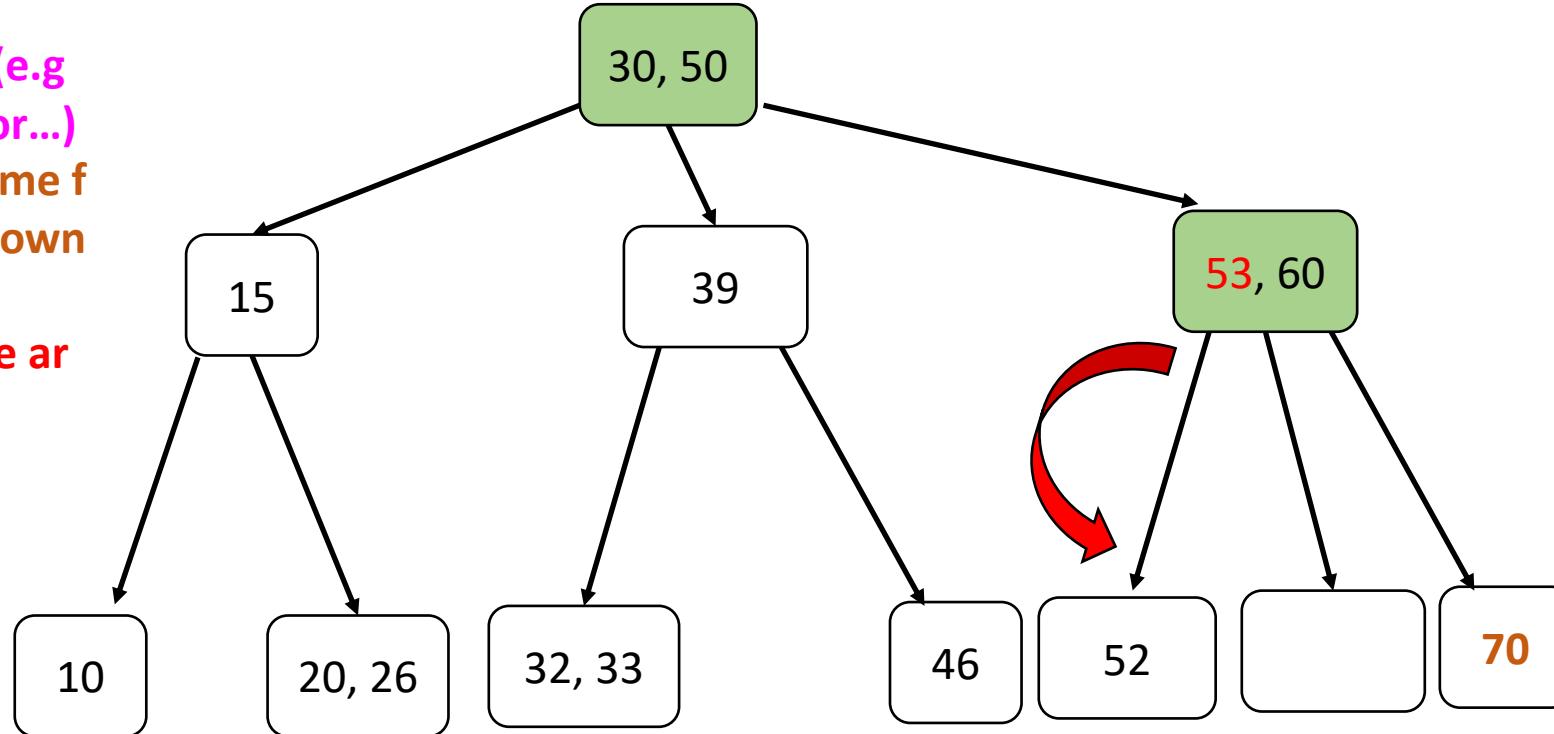
1. Search
2. Have to replace key (e.g. with inorder successor...)
3. Use our familiar scheme from before to push down 70 (again)



# Deletion

Delete 66

1. Search
2. Have to replace key (e.g with inorder successor...)
3. Use our familiar scheme from before to push down 70 (again)
4. Rotate 53 left and we are done.

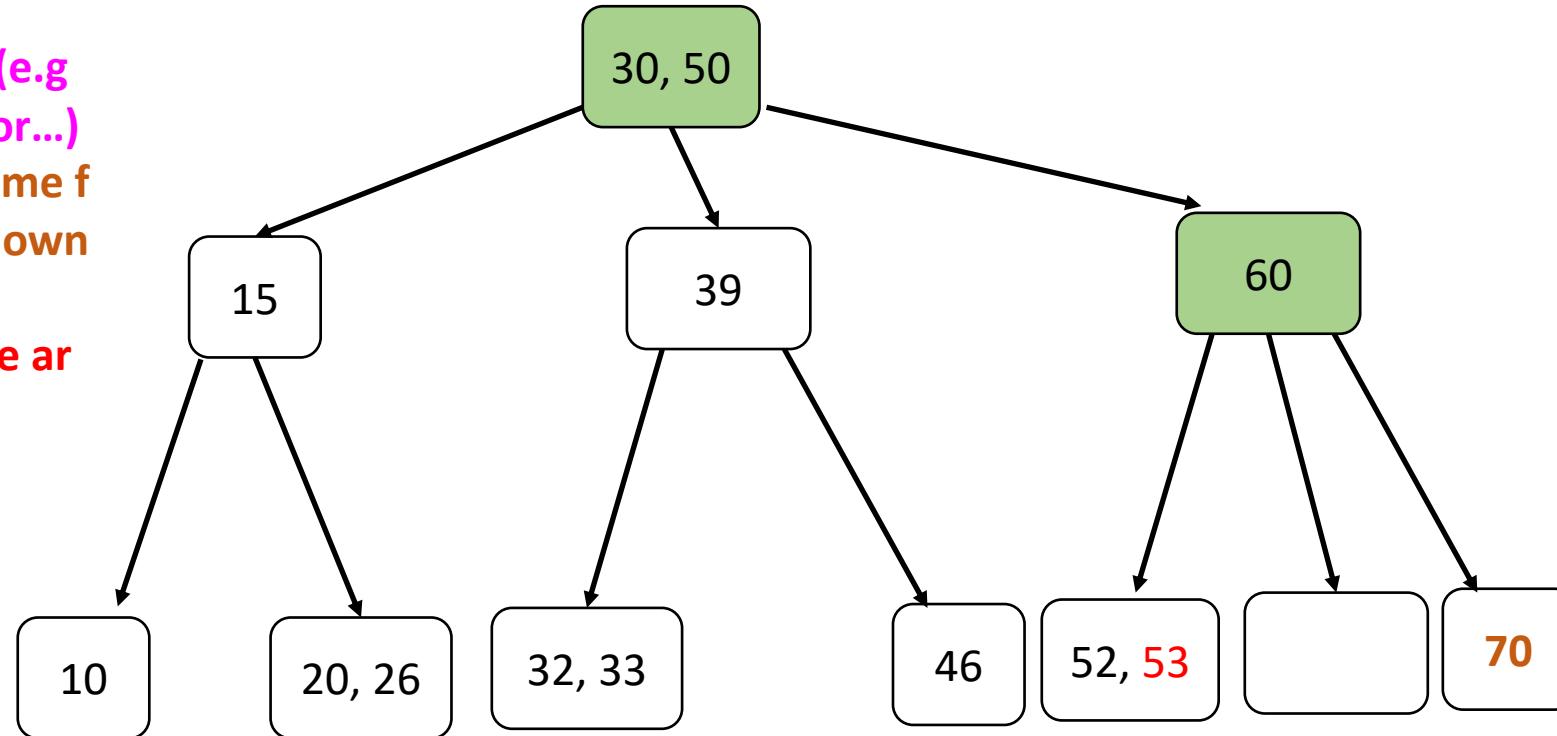


We have to merge the siblings because in this case, no key rotation from the middle

# Deletion

Delete 66

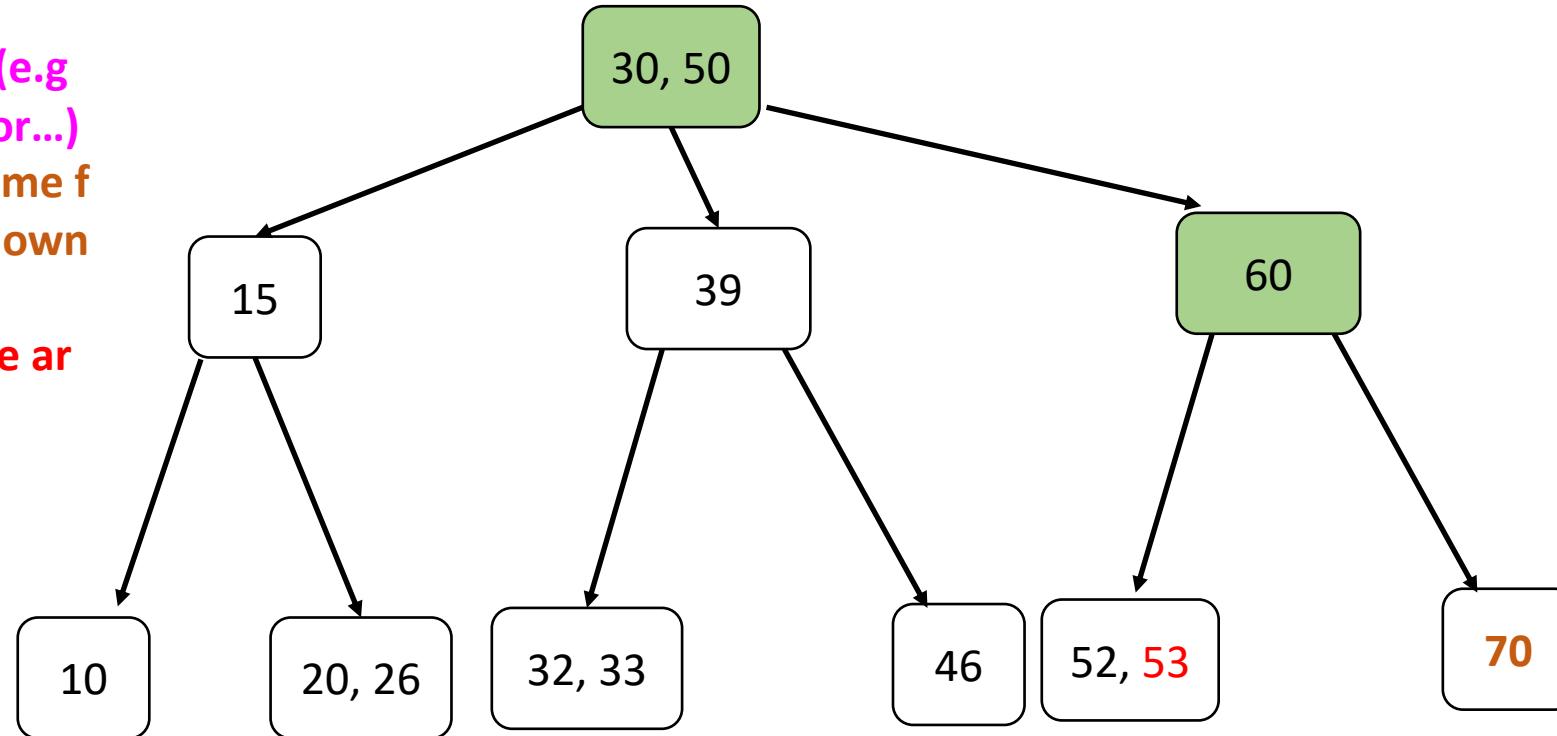
1. Search
2. Have to replace key (e.g with inorder successor...)
3. Use our familiar scheme from before to push down 70 (again)
4. Rotate 53 left and we are done.



# Deletion

Delete 66

1. Search
2. Have to replace key (e.g with inorder successor...)
3. Use our familiar scheme from before to push down 70 (again)
4. Rotate 53 left and we are done.

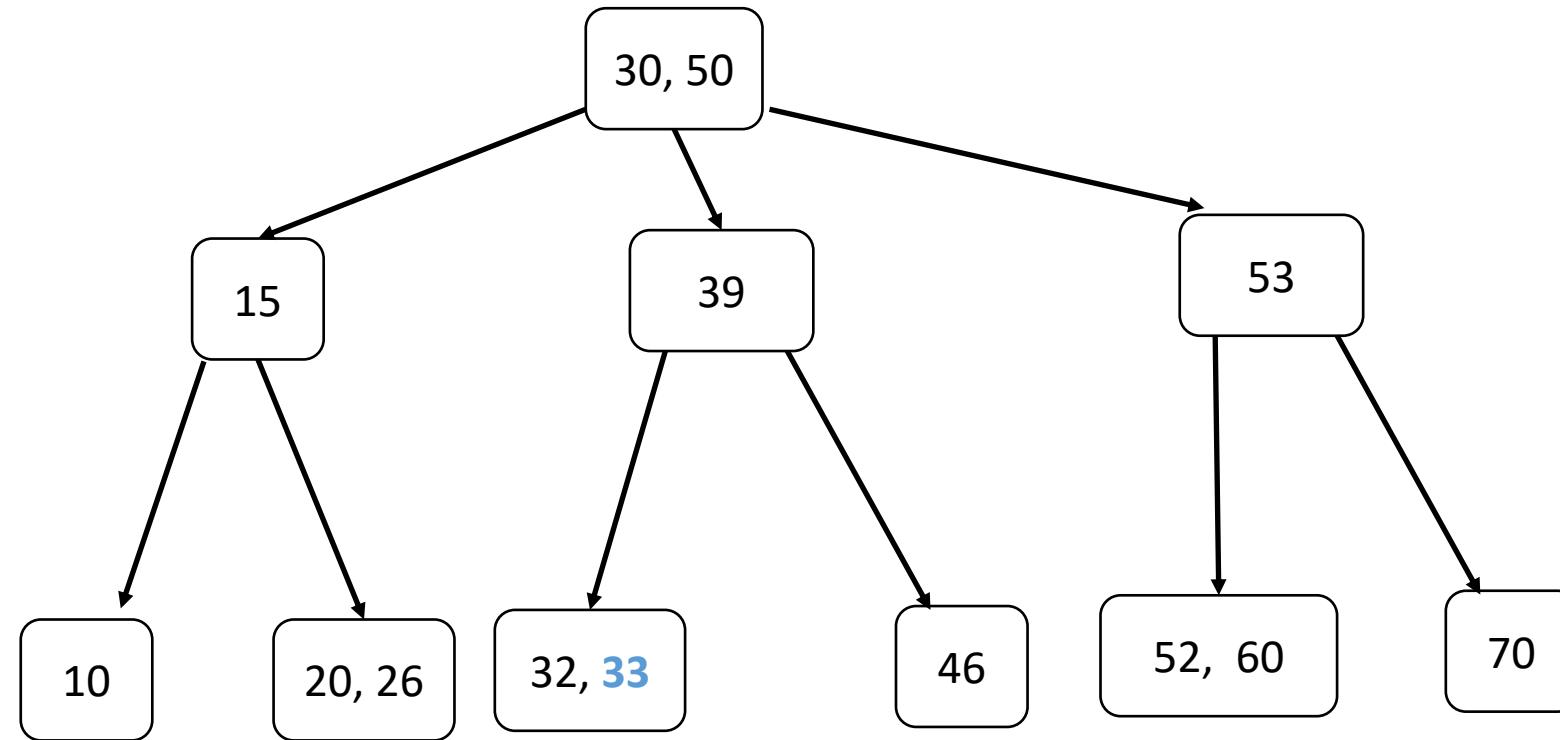


DONE ☺

Scalable Computing Systems Laboratory  
Hanyang University

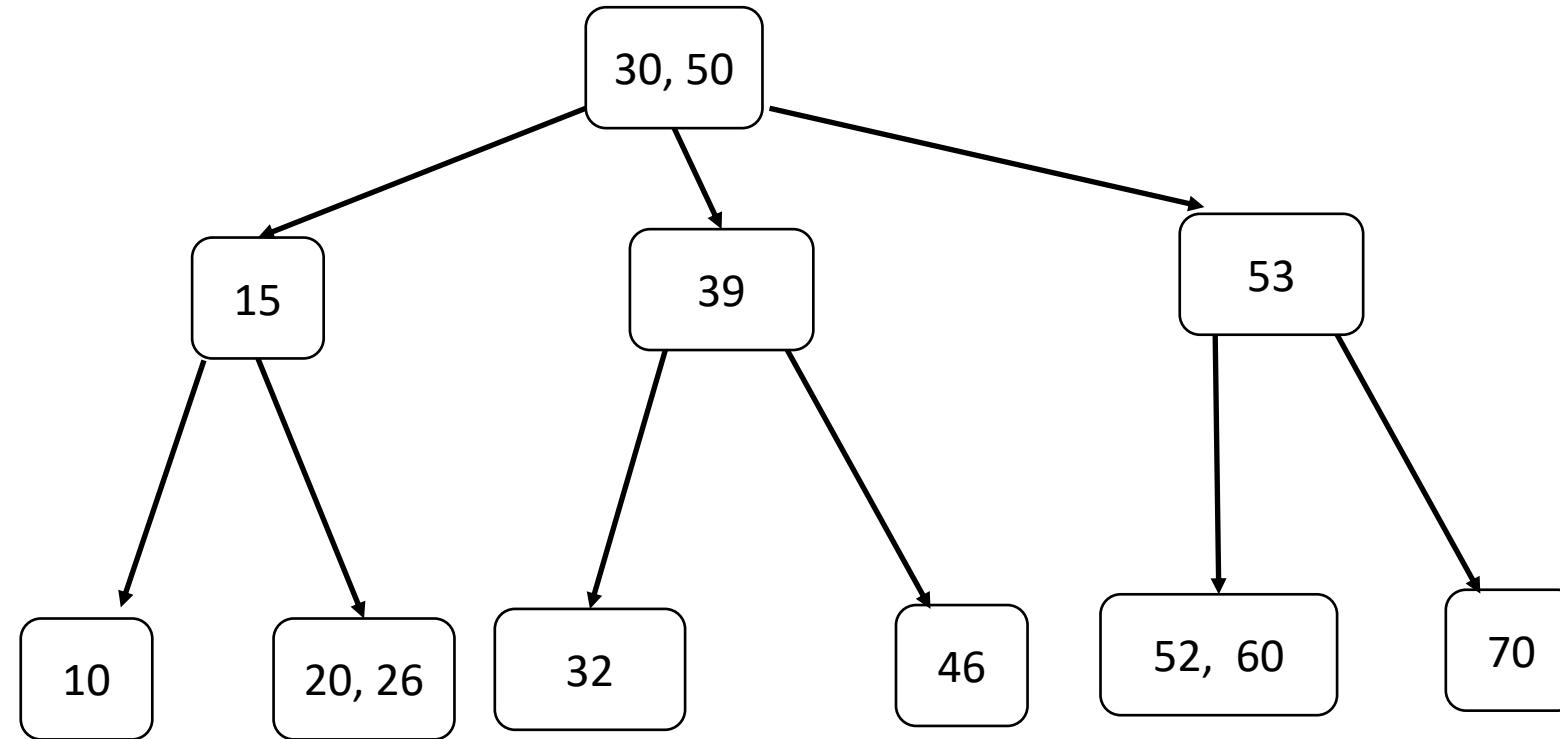
# Deletion

Let's delete 33 real quick to  
make a point...



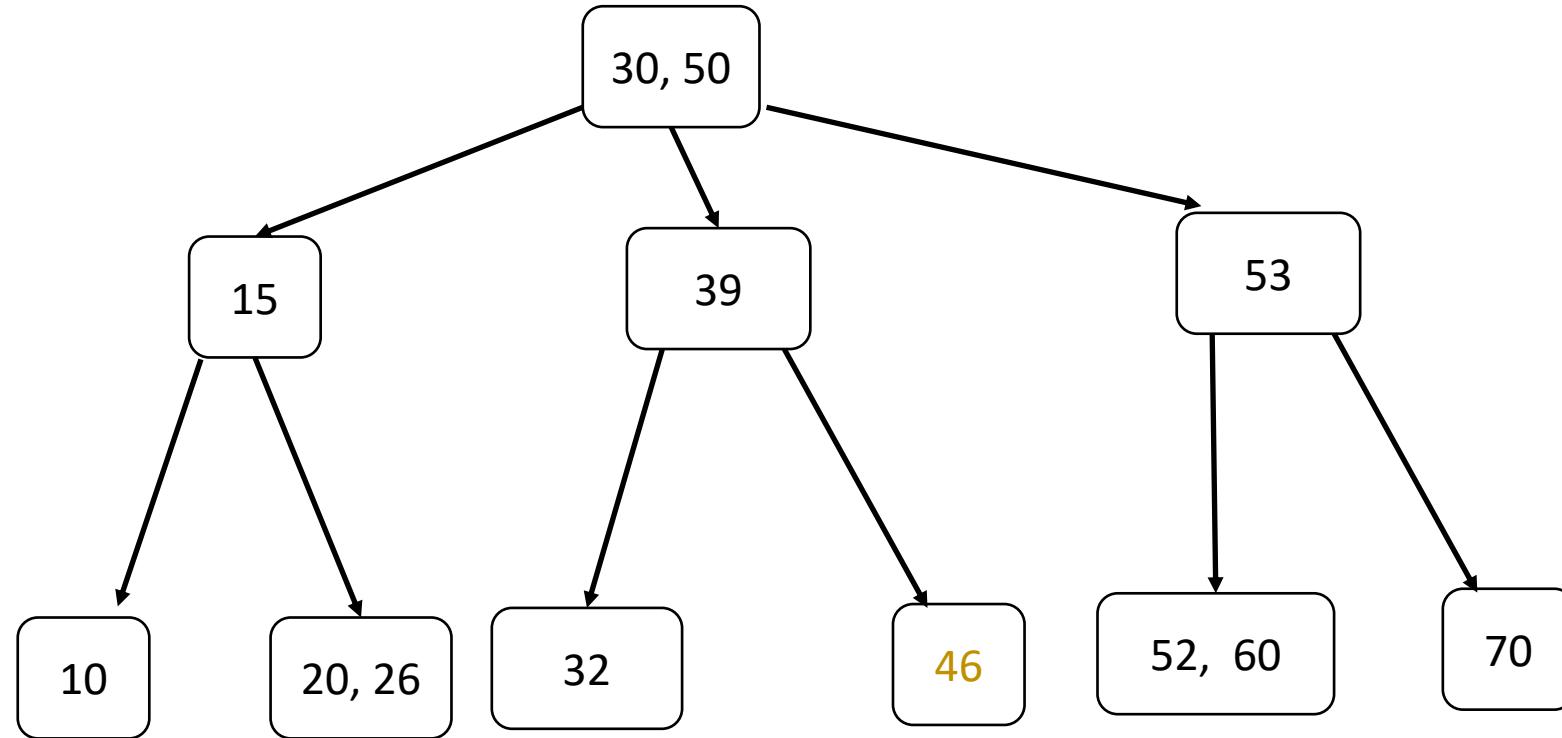
# Deletion

Let's delete 33 real quick to  
make a point...



# Deletion

Let's delete 33 real quick to  
make a point...  
And now let's delete 46.

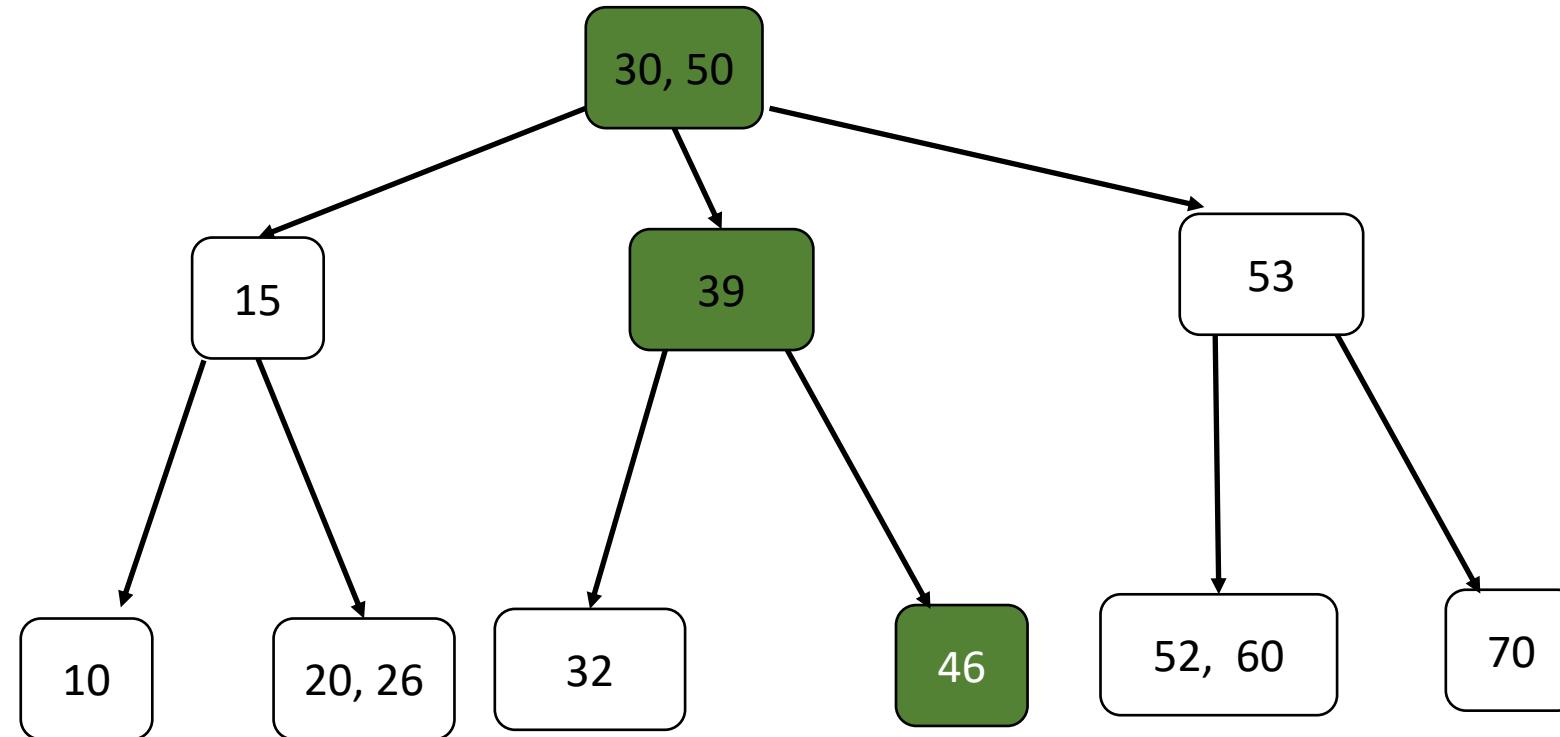


# Deletion

Let's delete 33 real quick to  
make a point...

And now let's delete 46.

1. Search

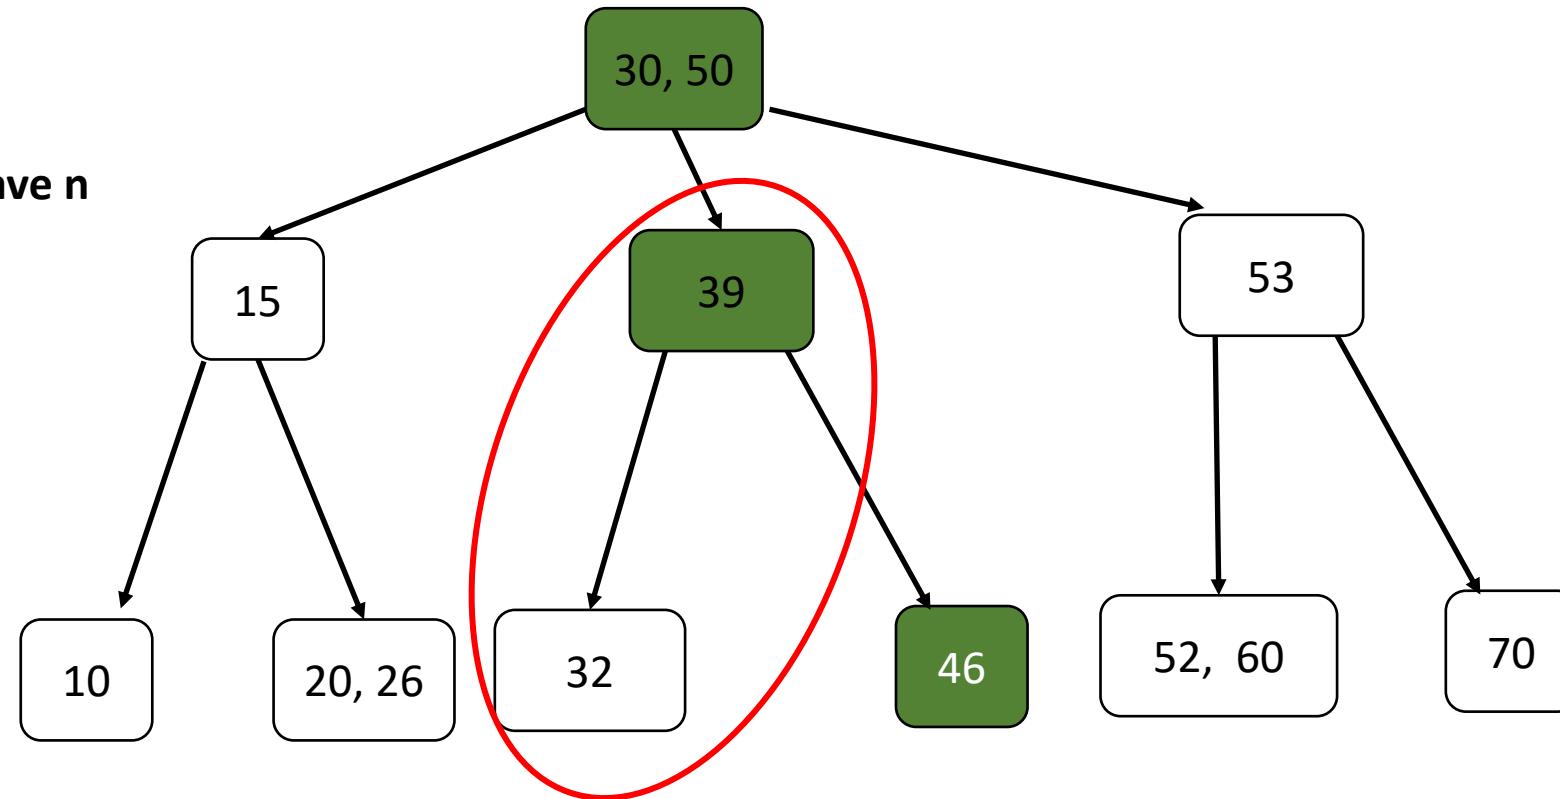


# Deletion

Let's delete 33 real quick to  
make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞

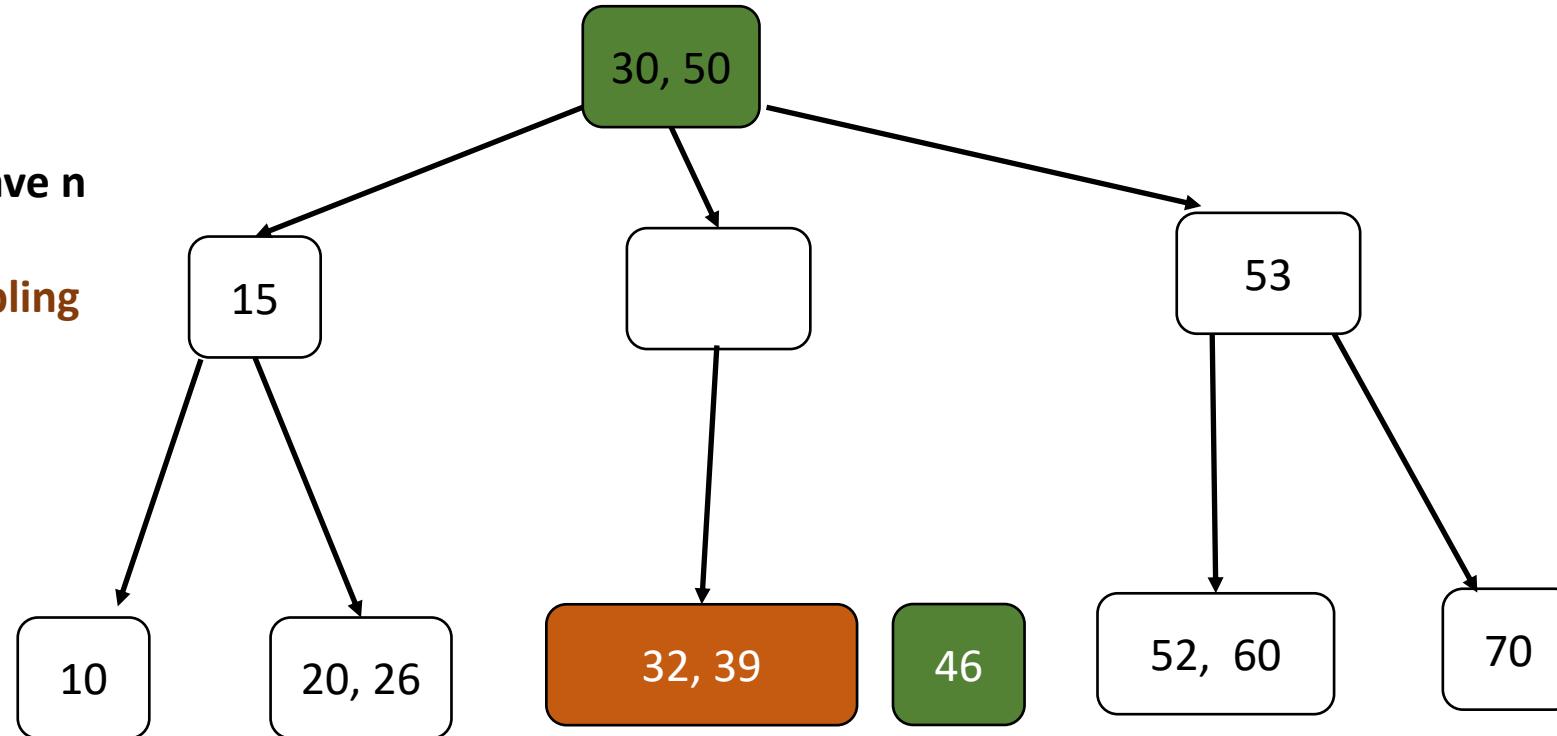


# Deletion

Let's delete 33 real quick to  
make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞
3. Merge parent and sibling into a 3-node...

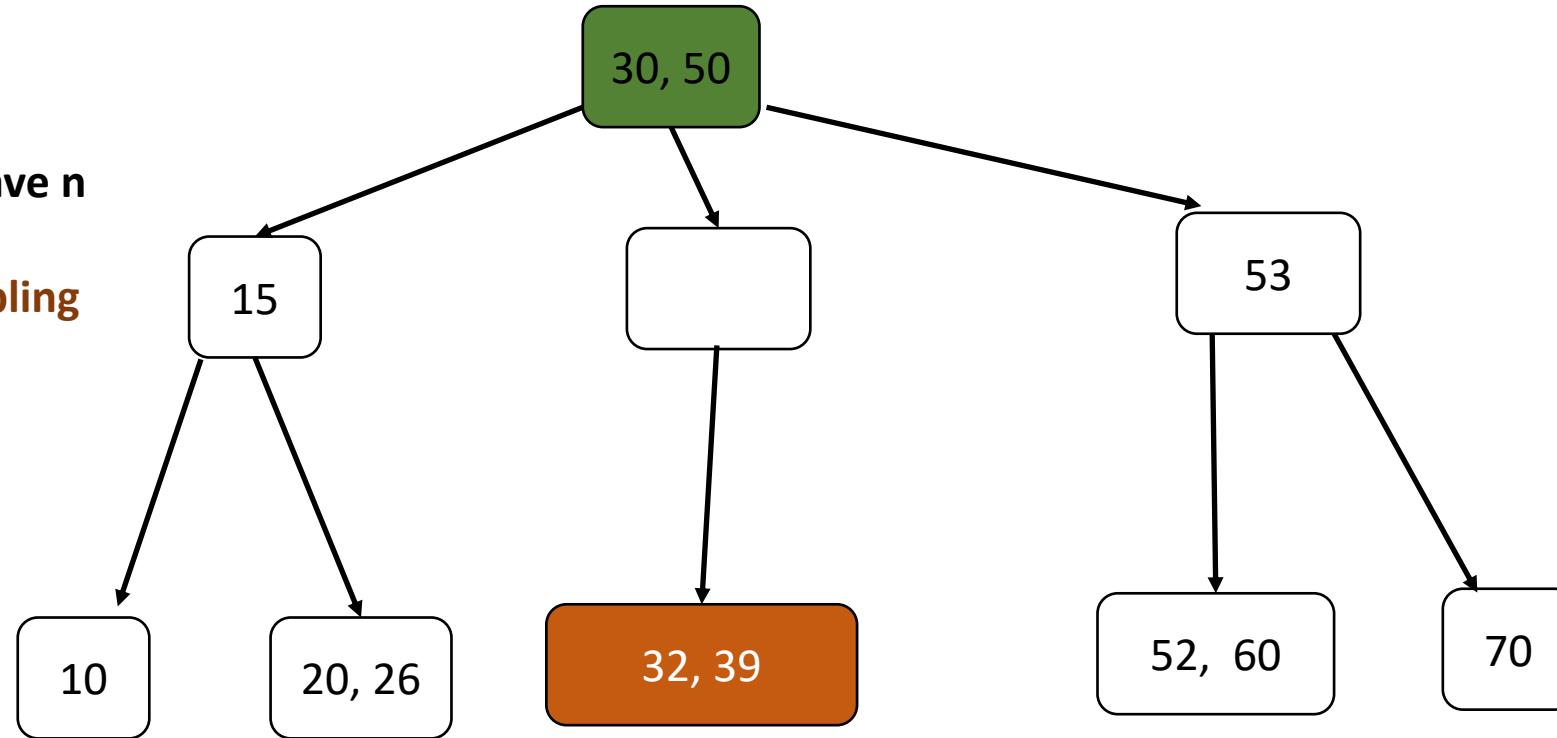


# Deletion

Let's delete 33 real quick to  
make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞
3. Merge parent and sibling into a 3-node...

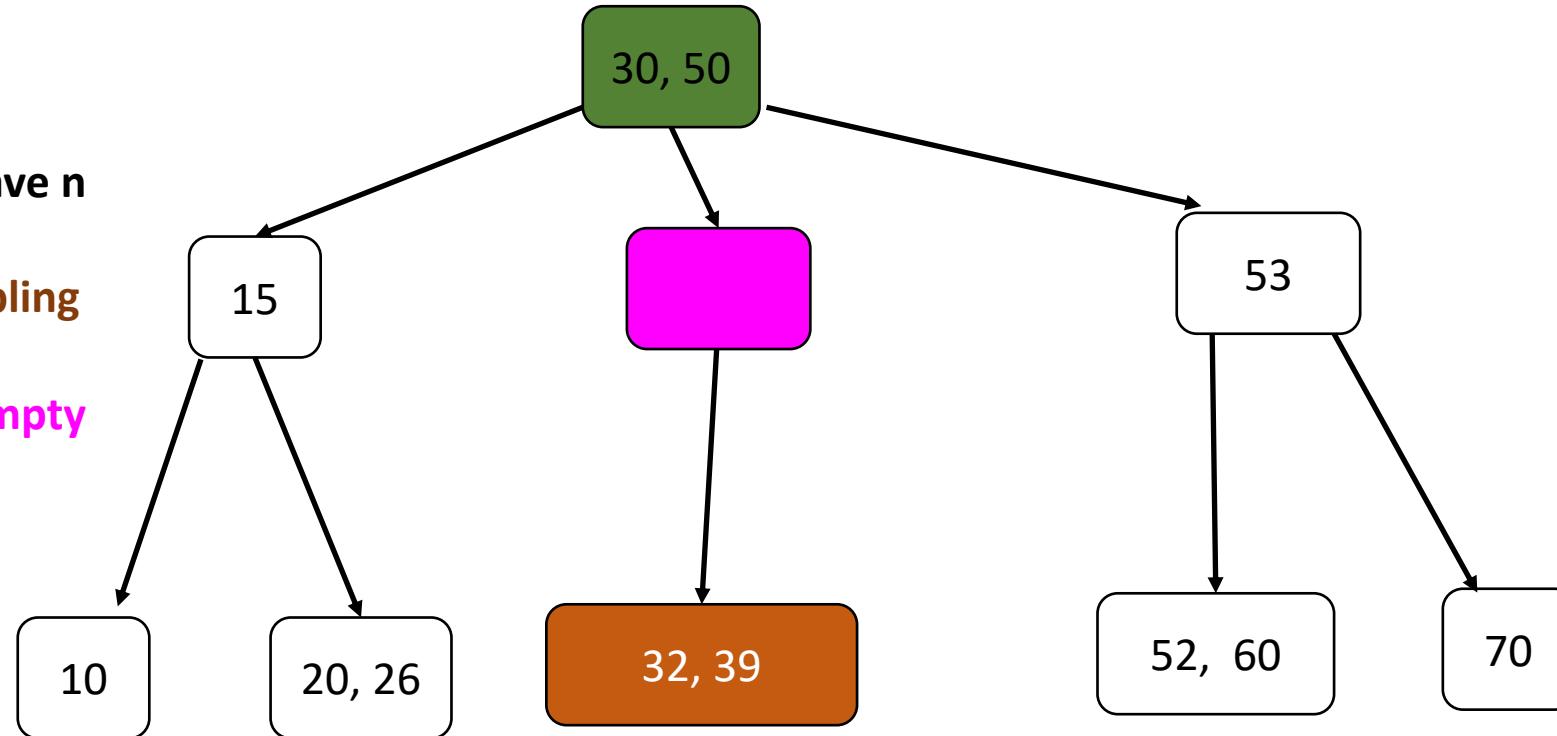


# Deletion

Let's delete 33 real quick to make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞
3. Merge parent and sibling into a 3-node...
4. Deal with the new empty node recursively!

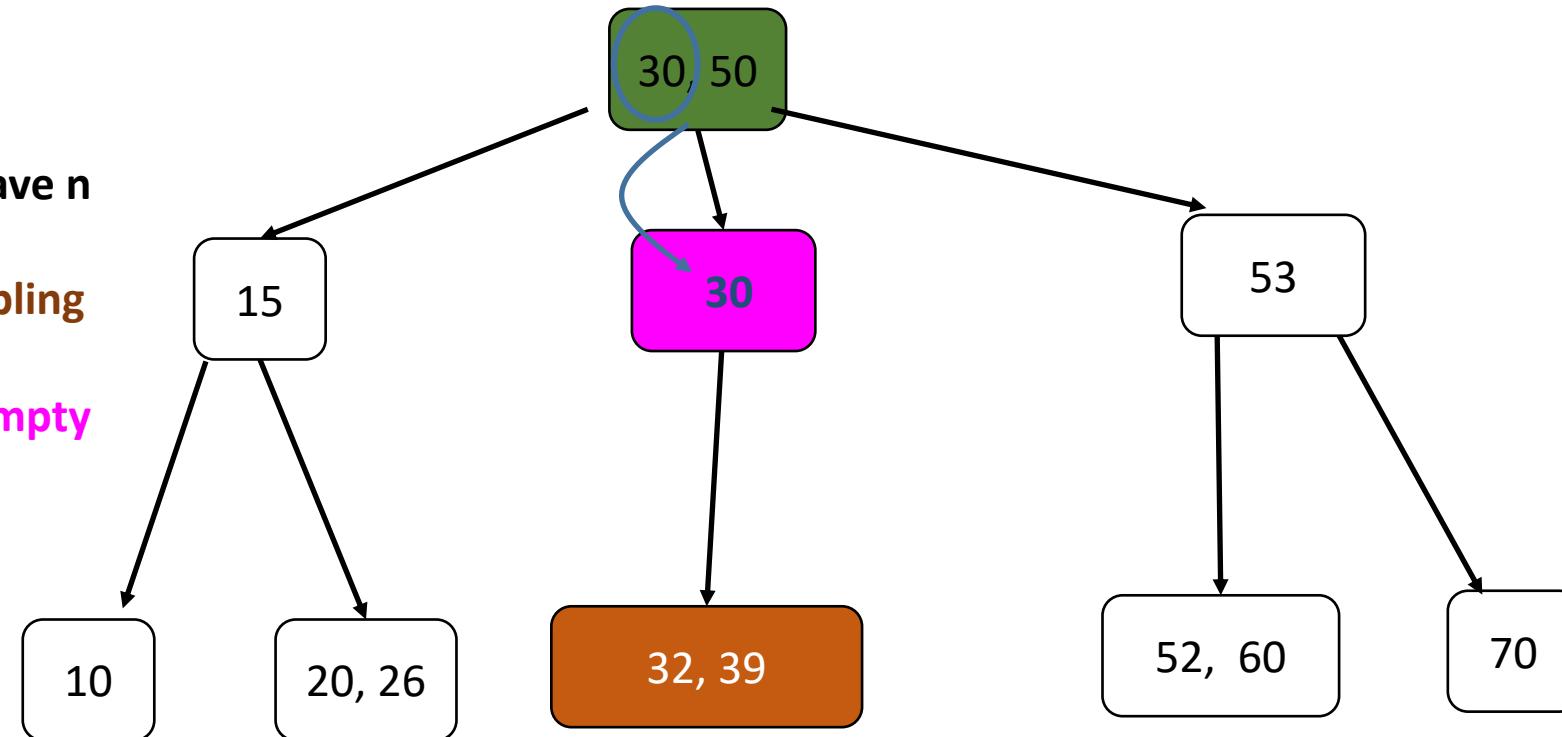


# Deletion

Let's delete 33 real quick to make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞
3. Merge parent and sibling into a 3-node...
4. Deal with the new empty node recursively!



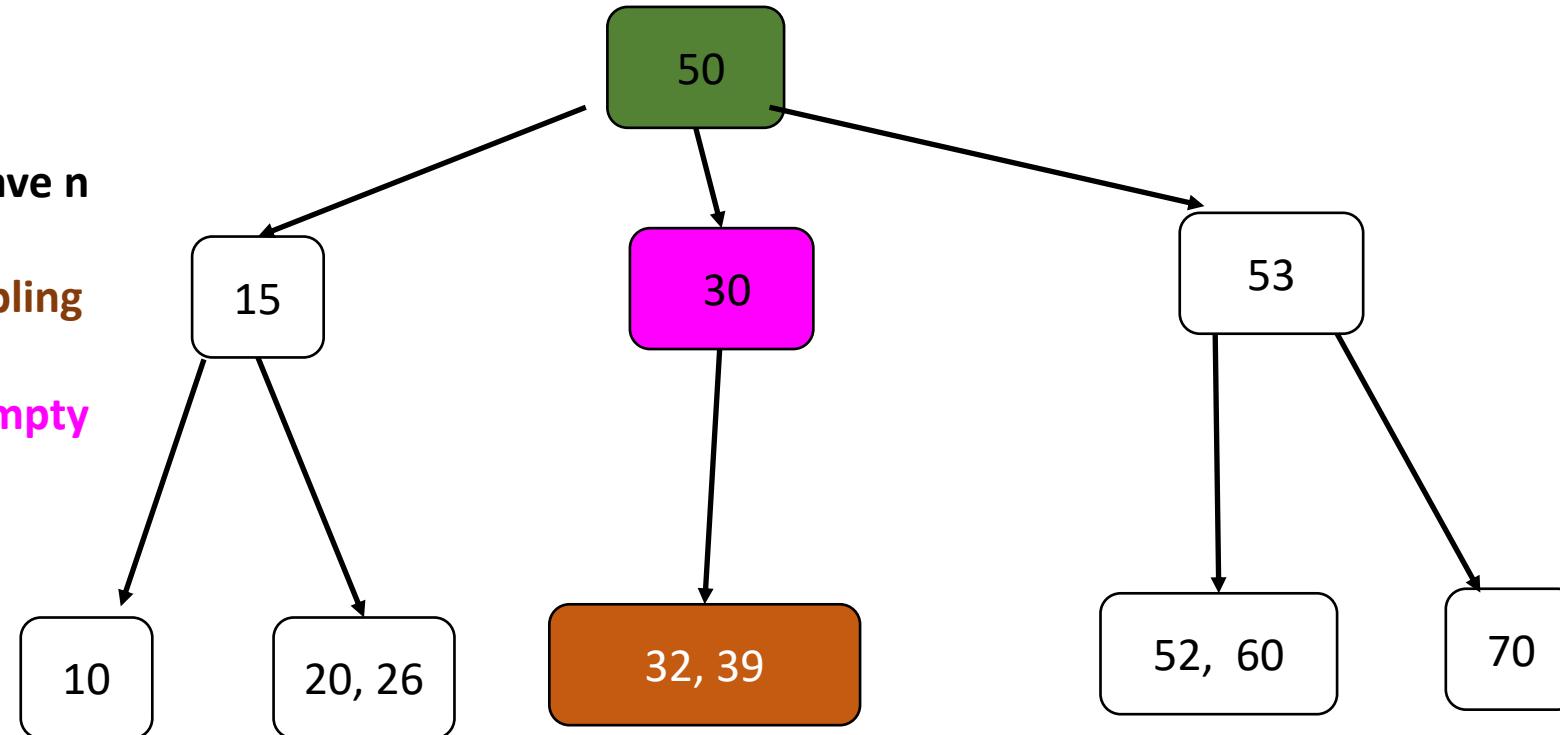
In this case, since only the root is 3-node, borrow appropriate key from parent and merge node with sibling into a three-node!

# Deletion

Let's delete 33 real quick to make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞
3. Merge parent and sibling into a 3-node...
4. Deal with the new empty node recursively!



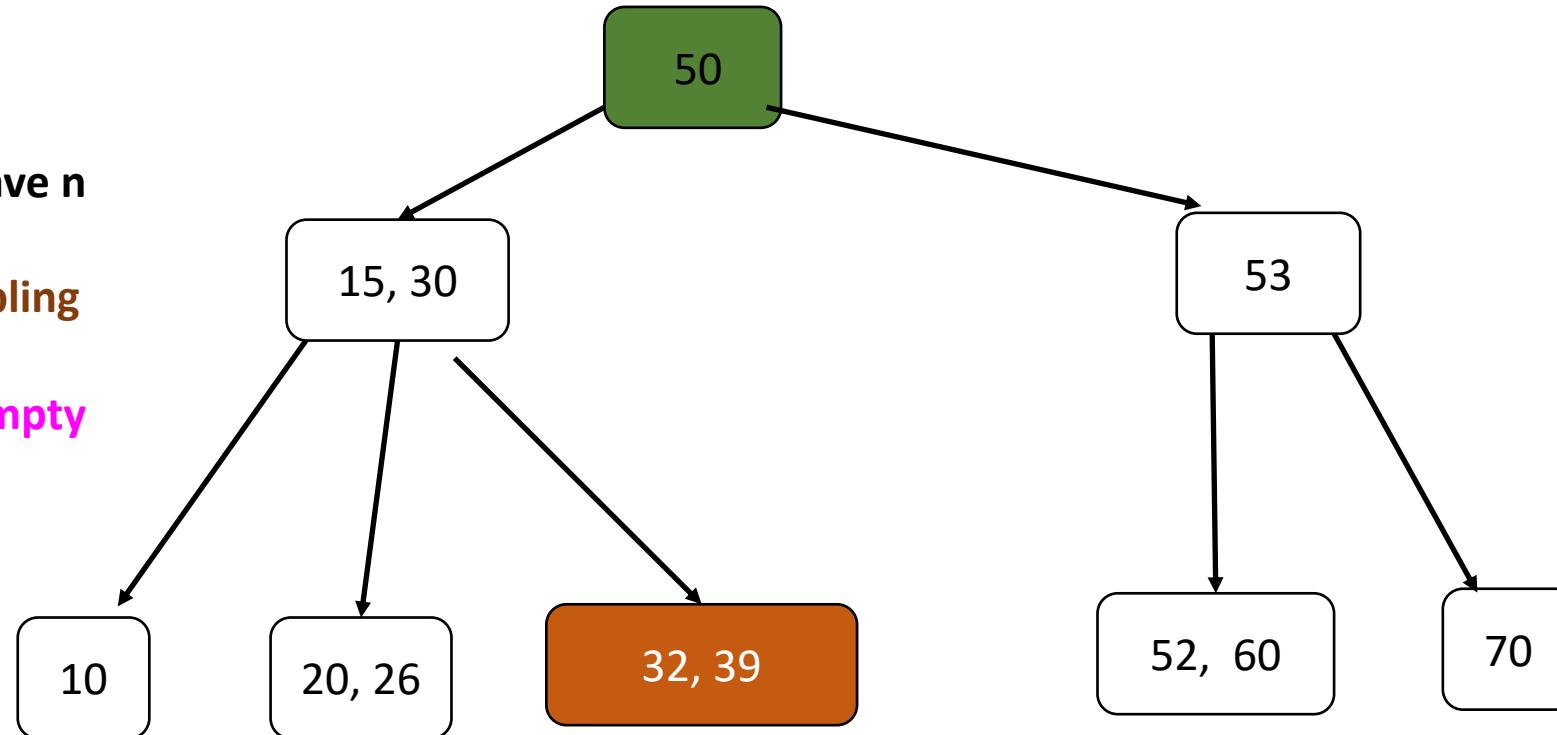
In this case, since only the root is 3-node, borrow appropriate key from parent and merge node with sibling into a three-node!

# Deletion

Let's delete 33 real quick to make a point...

And now let's delete 46.

1. Search
2. Sibling and parent have no keys to spare! 😞
3. Merge parent and sibling into a 3-node...
4. Deal with the new empty node recursively!



In this case, since only the root is 3-node, borrow appropriate key from parent and merge node with sibling into a three-node!

# Red-Black Tree

---

# Red-black trees

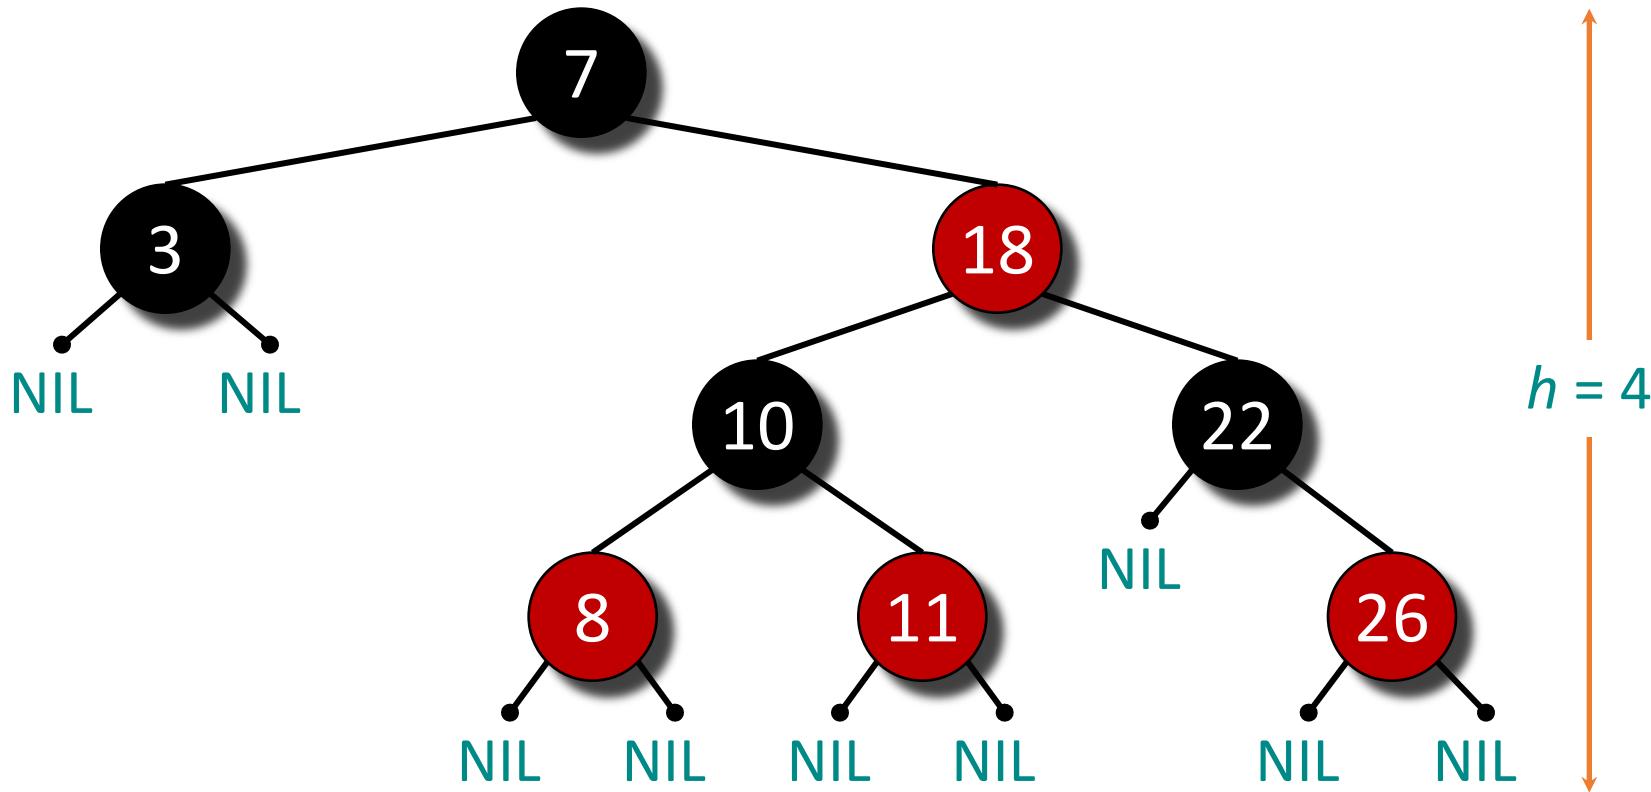
---

This data structure requires an extra one-bit **color** field in each node.

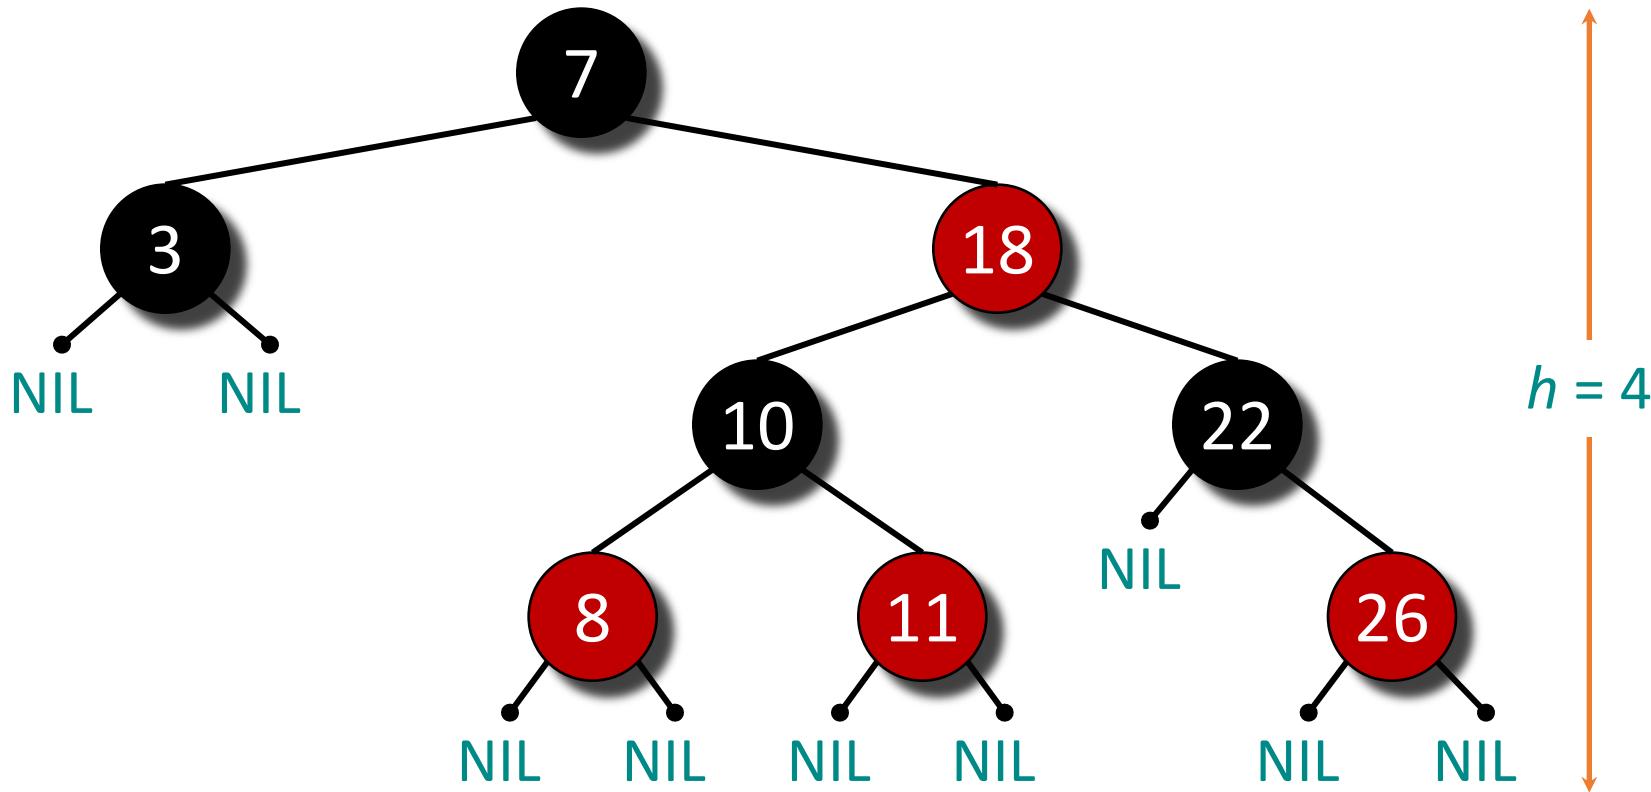
## *Red-black properties:*

1. Every node is either red or black.
2. The root and leaves (**NIL's**) are black.
3. If a node is red, then its parent is black.
4. All simple paths from any node  $x$  to a descendant leaf have the same number of black nodes = **black-height( $x$ )**.

# Example of a red-black tree

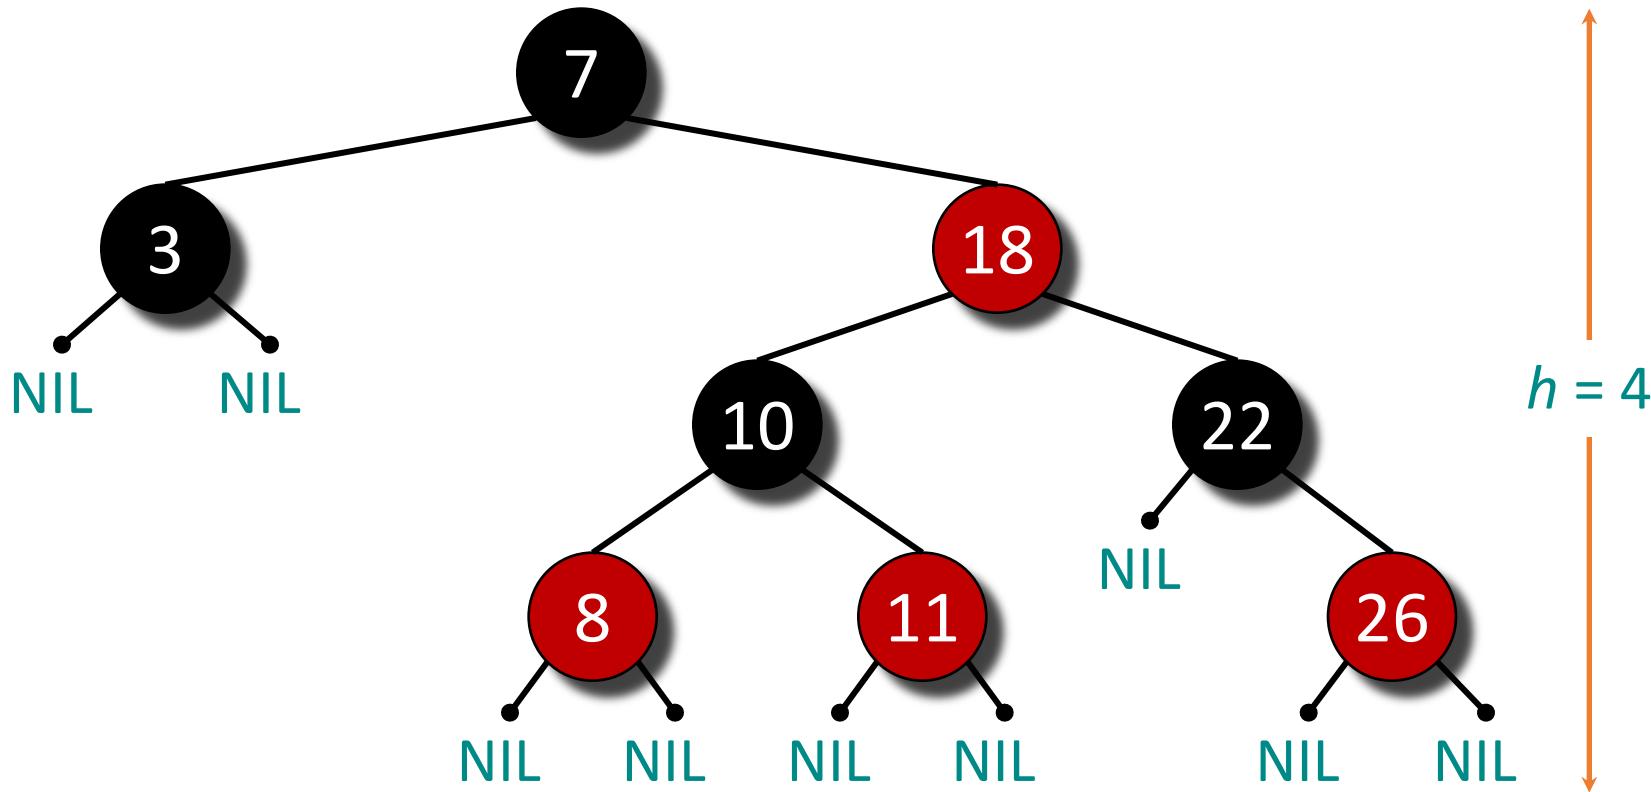


# Example of a red-black tree



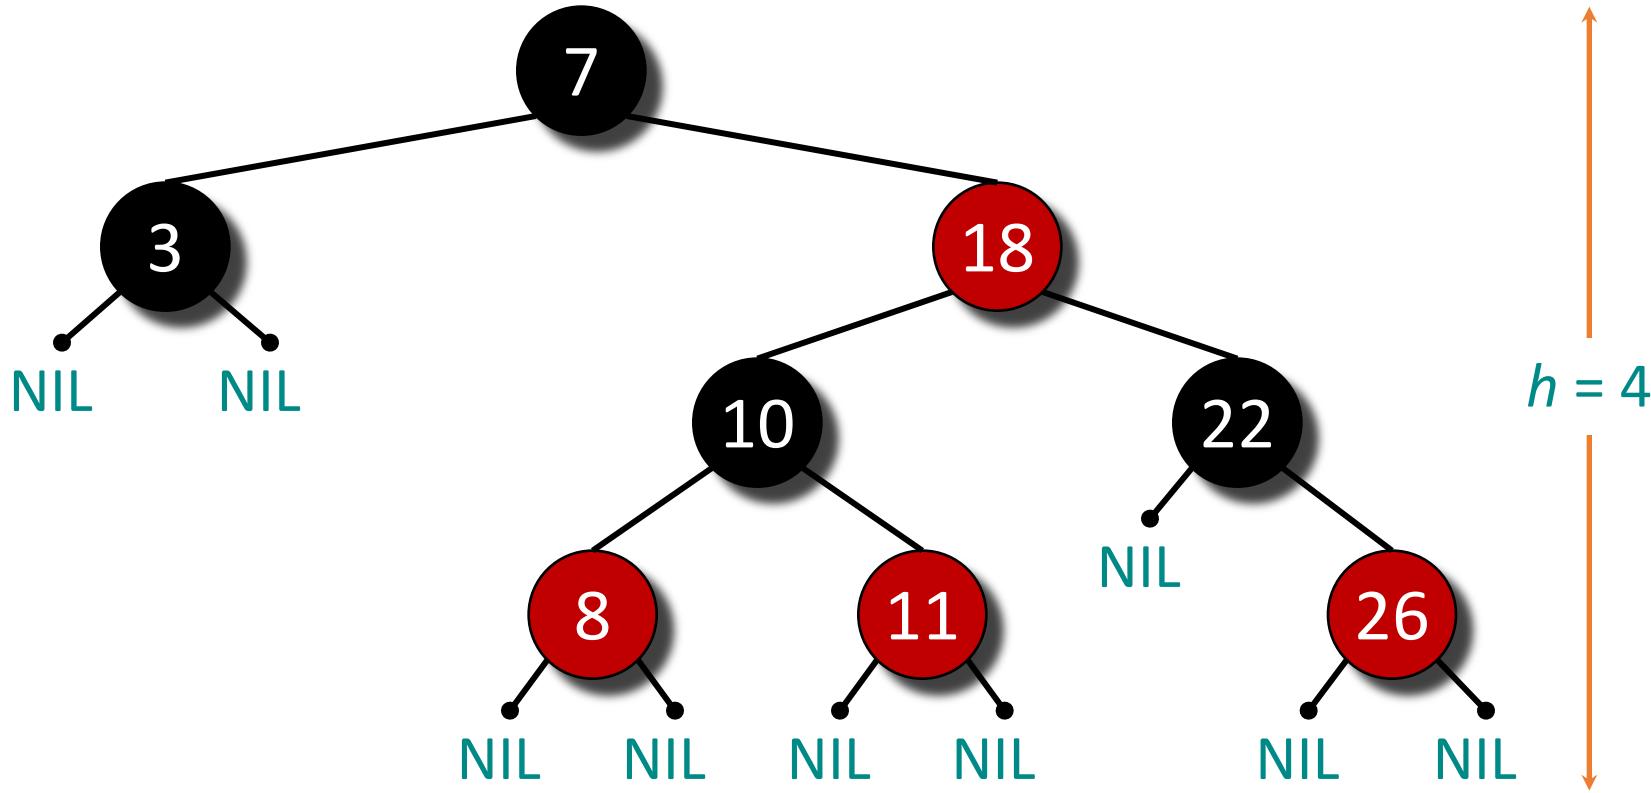
1. Every node is either red or black.

# Example of a red-black tree



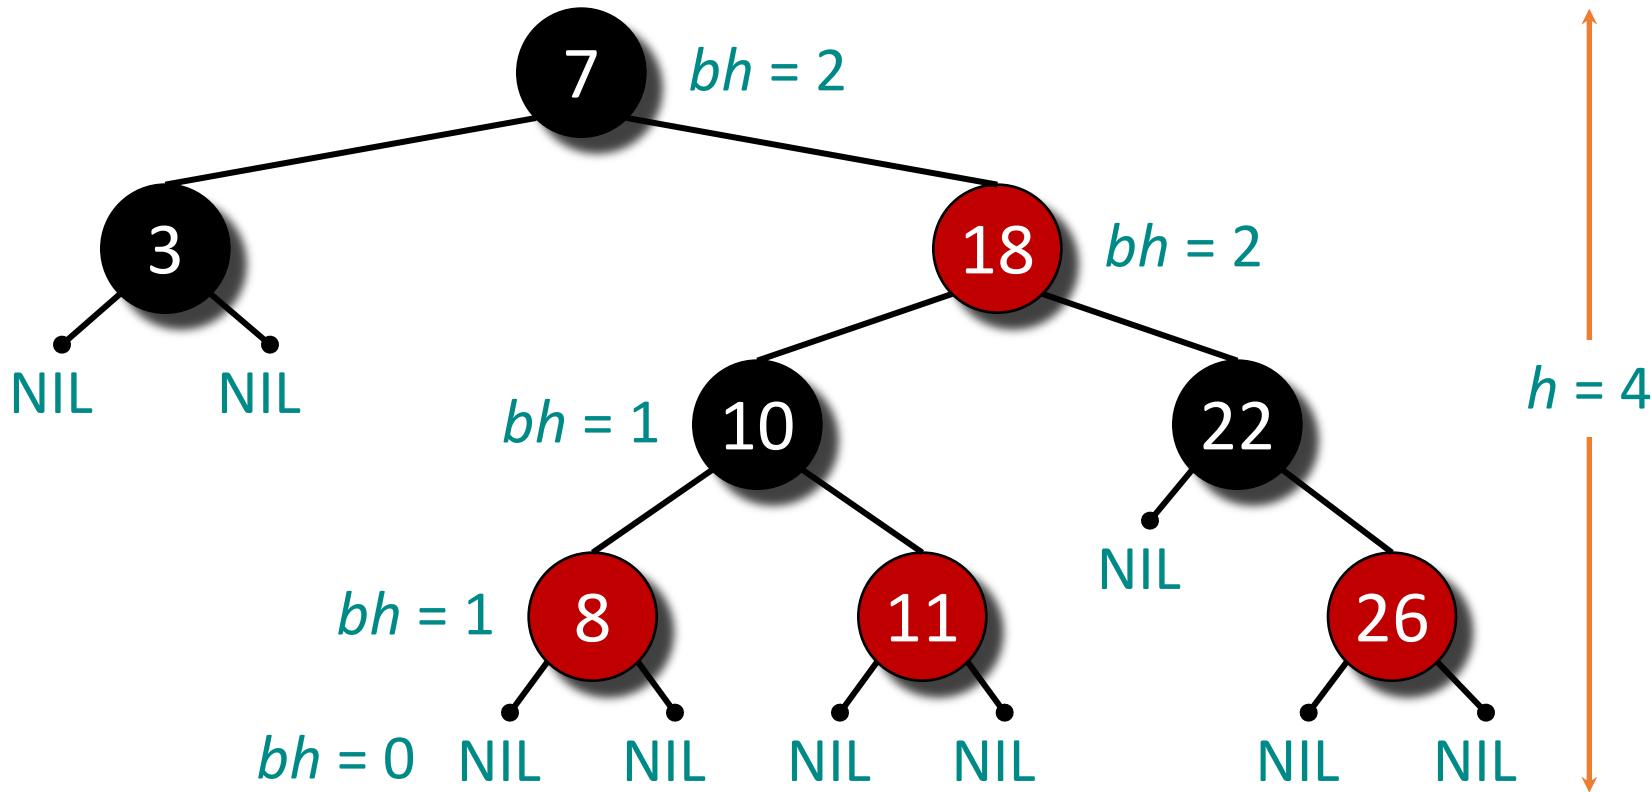
2. The root and leaves (NIL's) are black.

# Example of a red-black tree



3. If a node is red, then its parent is black.

# Example of a red-black tree



4. All simple paths from any node  $x$  to a descendant leaf have the same number of black nodes =  $\text{black-height}(x)$ .

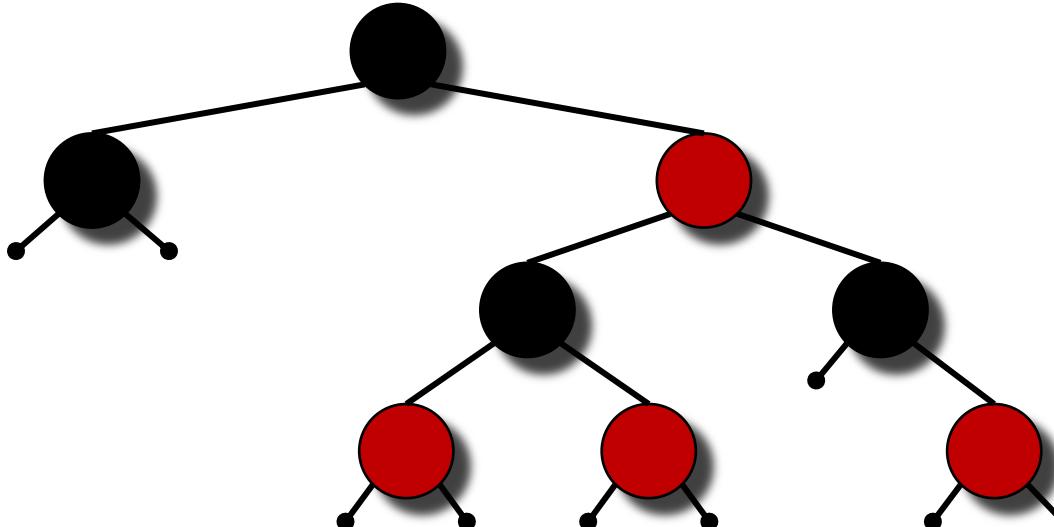
# Height of a red-black tree

**Theorem.** A red-black tree with  $n$  keys has height  
$$h \leq 2 \lg(n + 1).$$

*Proof.* (The book uses induction. Read carefully.)

## INTUITION:

- Merge red nodes into their black parents.



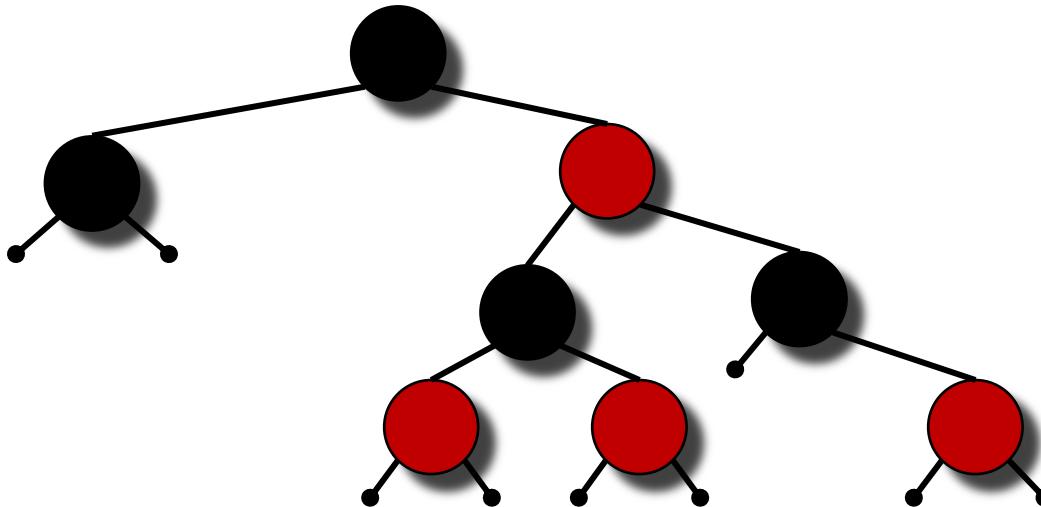
# Height of a red-black tree

**Theorem.** A red-black tree with  $n$  keys has height  
$$h \leq 2 \lg(n + 1).$$

*Proof.* (The book uses induction. Read carefully.)

## INTUITION:

- Merge red nodes into their black parents.



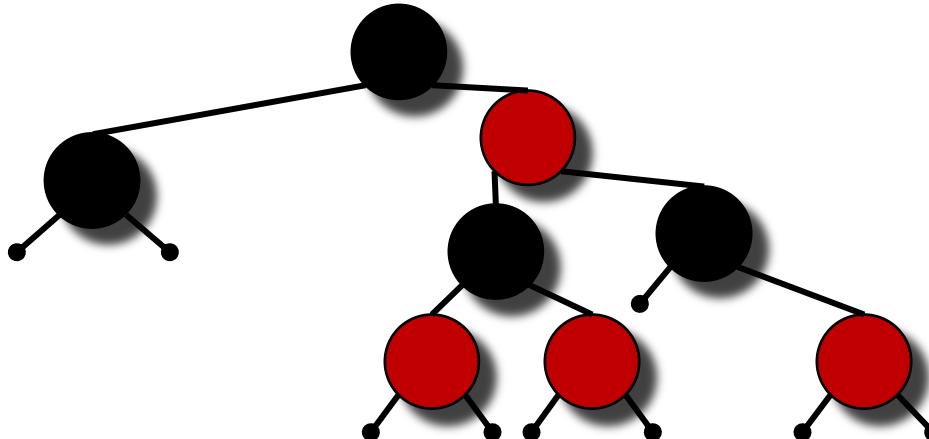
# Height of a red-black tree

**Theorem.** A red-black tree with  $n$  keys has height  
$$h \leq 2 \lg(n + 1).$$

*Proof.* (The book uses induction. Read carefully.)

## INTUITION:

- Merge red nodes into their black parents.



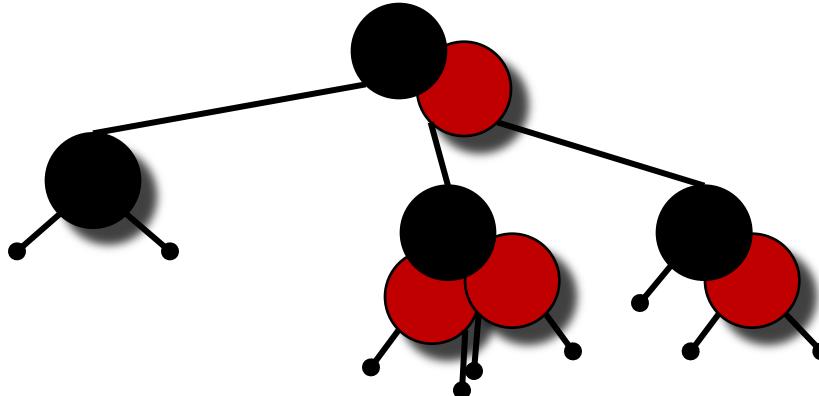
# Height of a red-black tree

**Theorem.** A red-black tree with  $n$  keys has height  
$$h \leq 2 \lg(n + 1).$$

*Proof.* (The book uses induction. Read carefully.)

## INTUITION:

- Merge red nodes into their black parents.



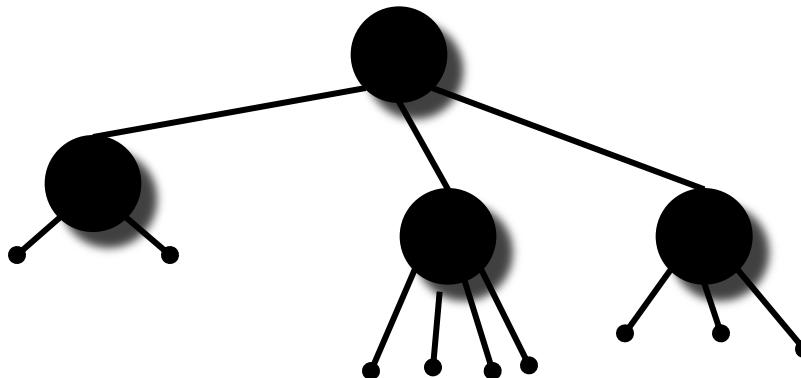
# Height of a red-black tree

**Theorem.** A red-black tree with  $n$  keys has height  
$$h \leq 2 \lg(n + 1).$$

*Proof.* (The book uses induction. Read carefully.)

## INTUITION:

- Merge red nodes into their black parents.



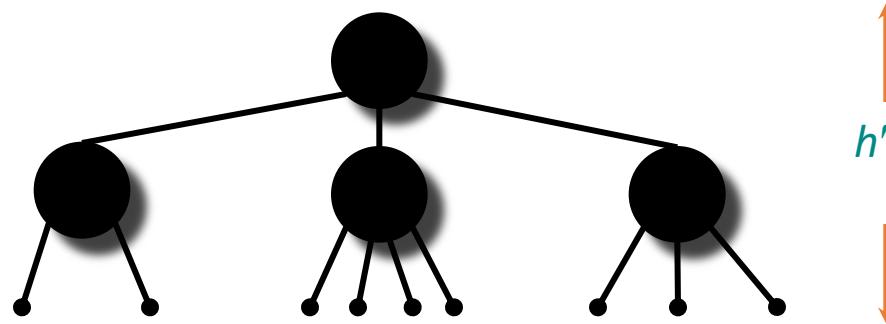
# Height of a red-black tree

**Theorem.** A red-black tree with  $n$  keys has height  
$$h \leq 2 \lg(n + 1).$$

**Proof.** (The book uses induction. Read carefully.)

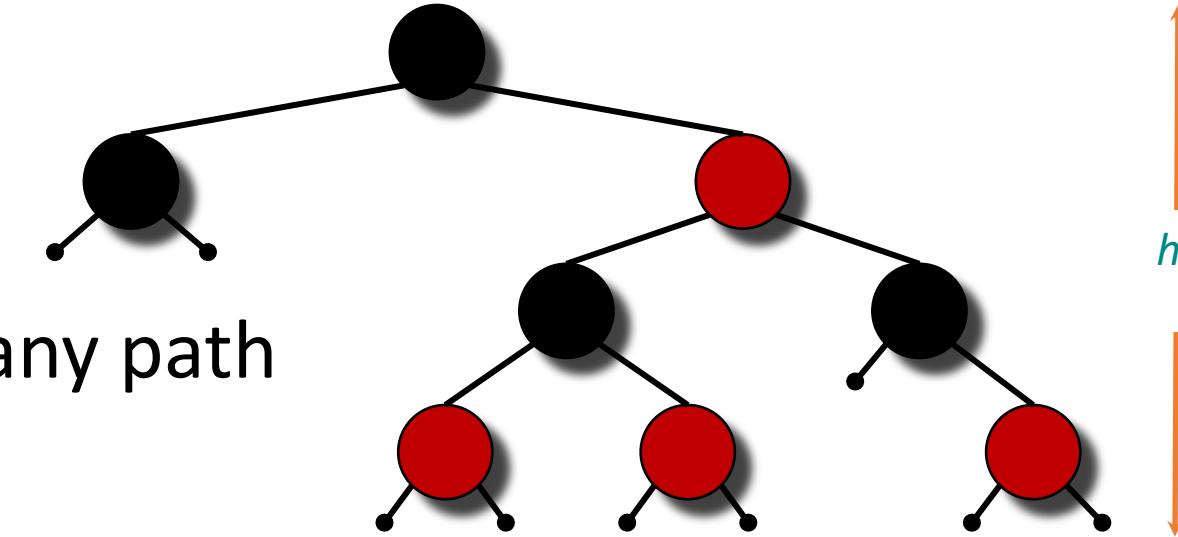
## INTUITION:

- Merge red nodes into their black parents.
- This process produces a tree in which each node has 2, 3, or 4 children.
- The 2-3-4 tree has uniform depth  $h'$  of leaves.

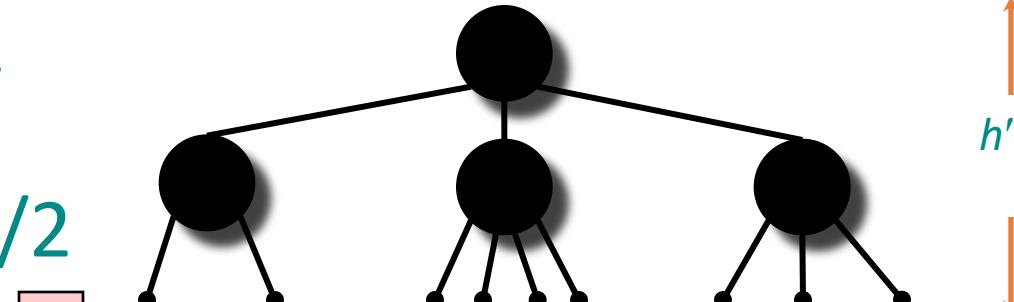


# Proof (continued)

- We have  $h' \geq h/2$ , since at most half the leaves on any path are red.



- The number of leaves in each tree is  $n + 1$   
 $\Rightarrow n + 1 \geq 2^{h'}$   
 $\Rightarrow \lg(n + 1) \geq h' \geq h/2$   
 $\Rightarrow h \leq 2 \lg(n + 1)$ . □



# Query operations

---

**Corollary.** The queries SEARCH, MIN, MAX, SUCCESSOR, and PREDECESSOR all run in  $O(\lg n)$  time on a red-black tree with  $n$  nodes.

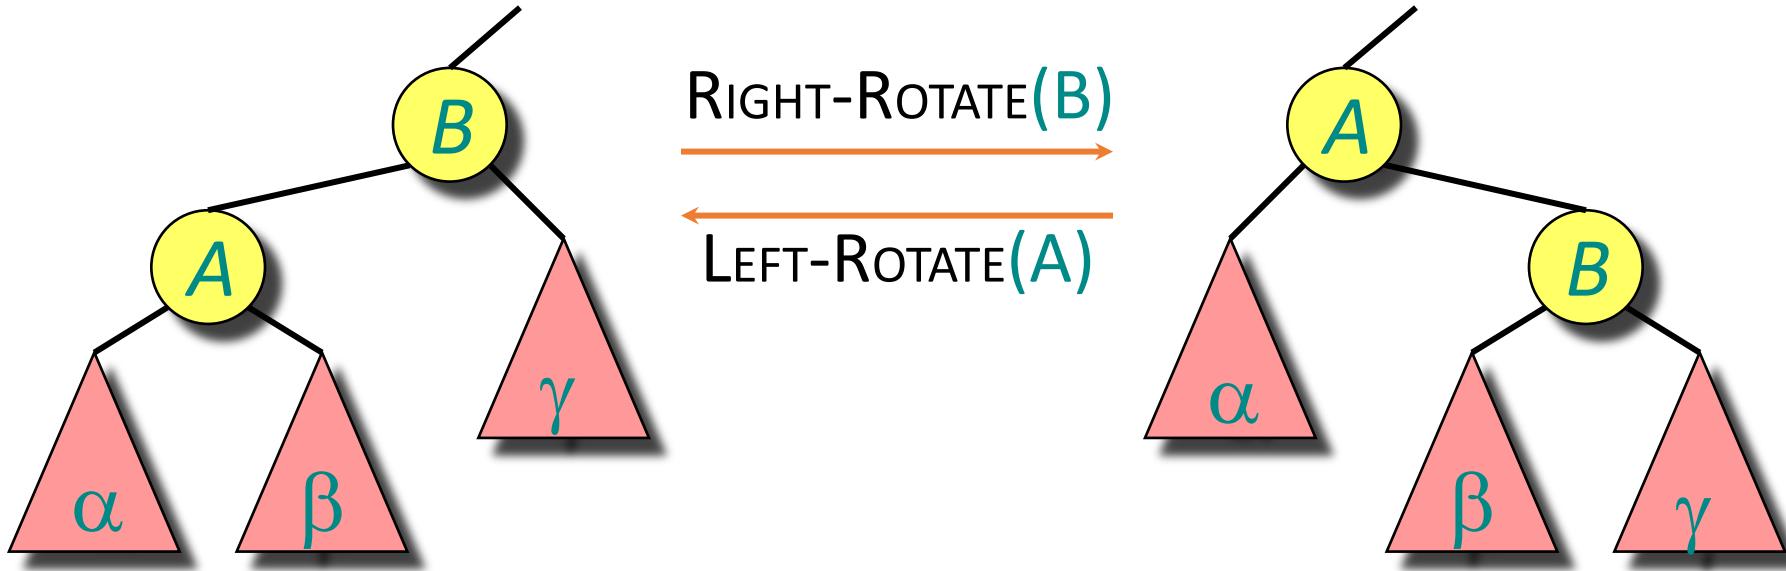
# Modifying operations

---

The operations `INSERT` and `DELETE` cause modifications to the red-black tree:

- the operation itself,
- color changes,
- restructuring the links of the tree via “*rotations*”.

# Rotations



Rotations maintain the inorder ordering of keys:

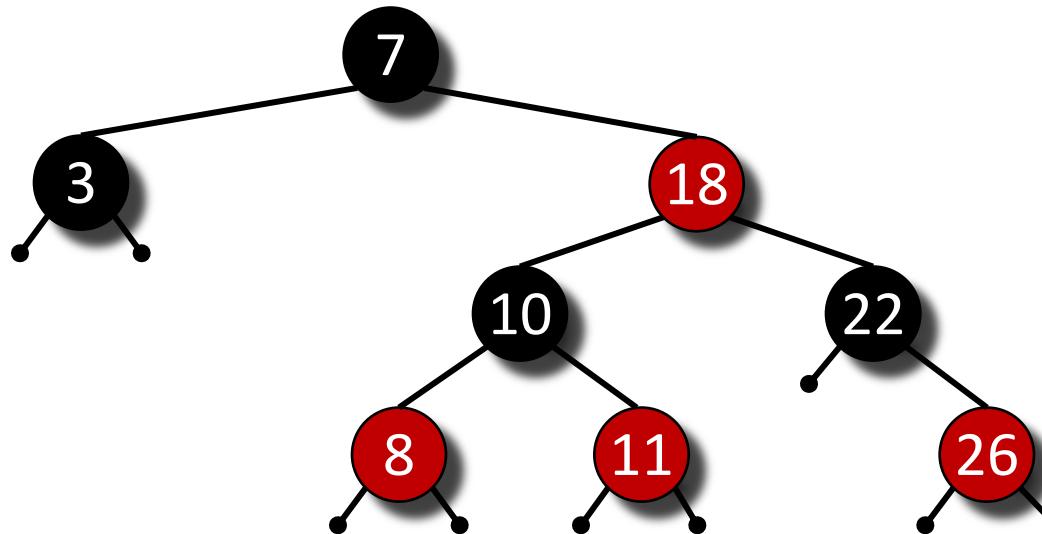
- $a \in \alpha, b \in \beta, c \in \gamma \Rightarrow a \leq A \leq b \leq B \leq c$ .

A rotation can be performed in  $O(1)$  time.

# Insertion into a red-black tree

**IDEA:** Insert  $x$  in tree. Color  $x$  red. Only red-black property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

**Example:**

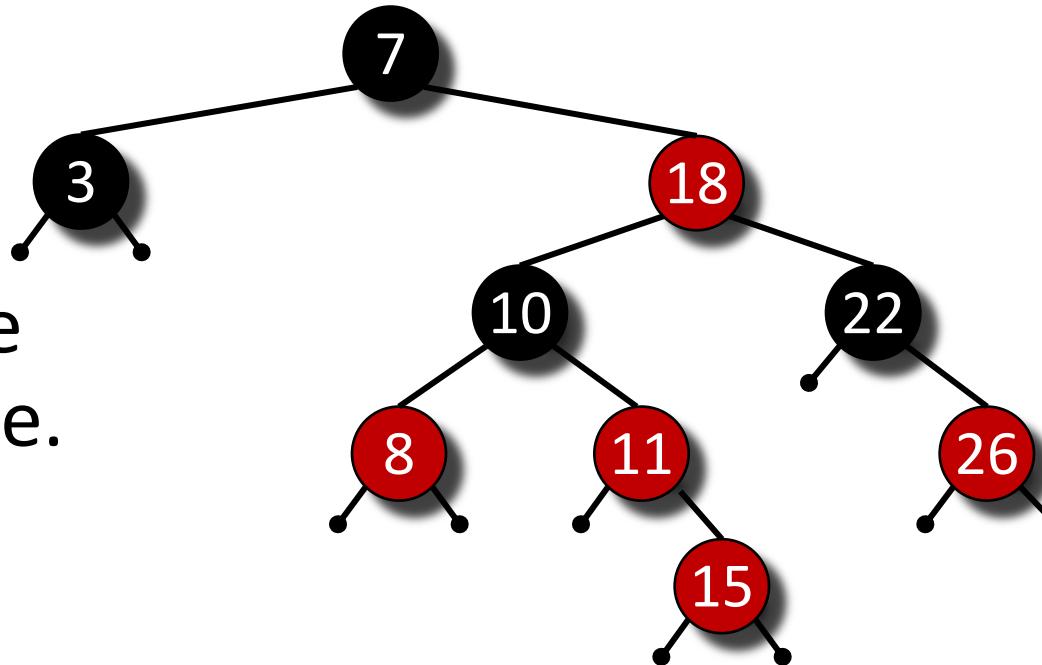


# Insertion into a red-black tree

**IDEA:** Insert  $x$  in tree. Color  $x$  red. Only red-black property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

## Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.

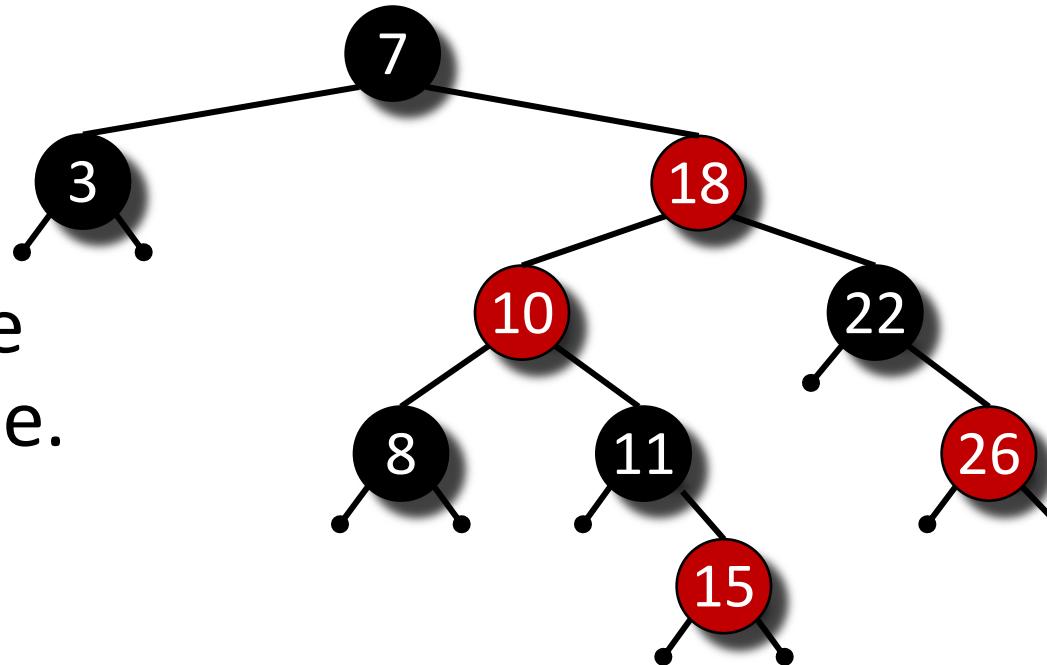


# Insertion into a red-black tree

**IDEA:** Insert  $x$  in tree. Color  $x$  red. Only red-black property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

## Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE(18).

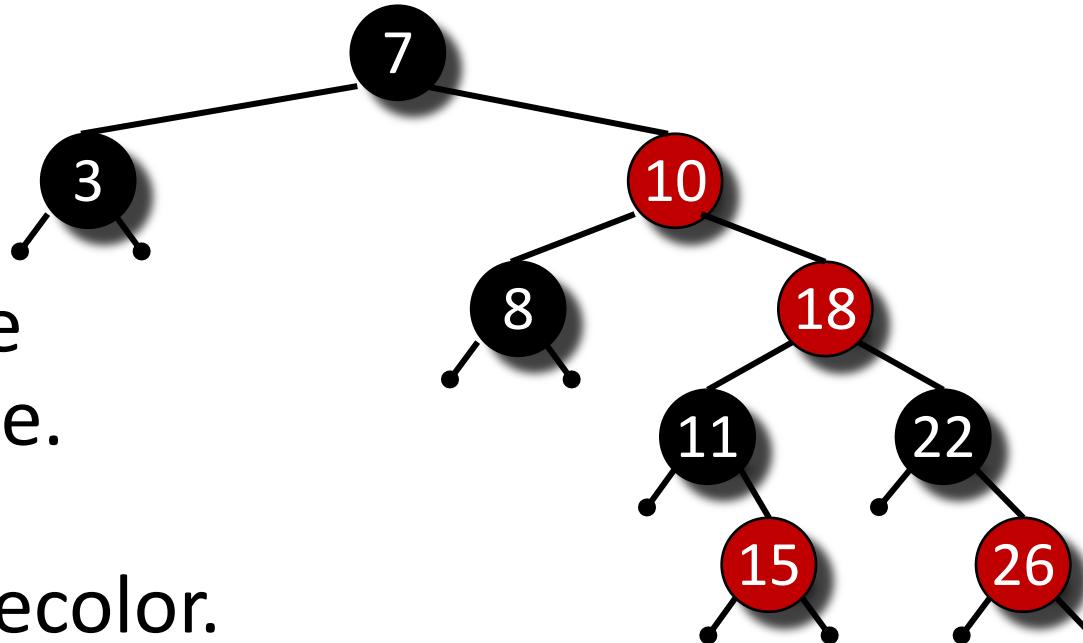


# Insertion into a red-black tree

**IDEA:** Insert  $x$  in tree. Color  $x$  red. Only red-black property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

## Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE(18).
- LEFT-ROTATE(7) and recolor.

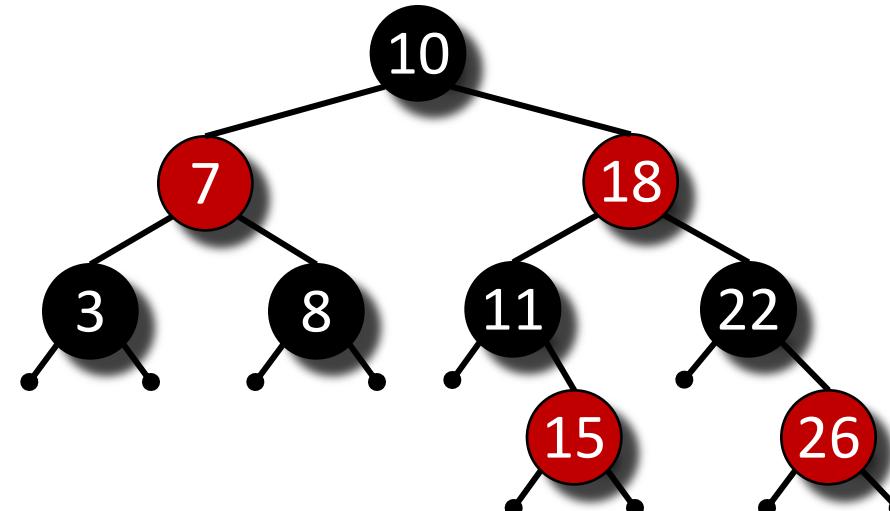


# Insertion into a red-black tree

**IDEA:** Insert  $x$  in tree. Color  $x$  red. Only red-black property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

## Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE(18).
- LEFT-ROTATE(7) and recolor.



# Pseudocode

**RB-INSERT( $T, x$ )**

**TREE-INSERT( $T, x$ )**

$\text{color}[x] \leftarrow \text{RED}$      ▷ only RB property 3 can be violated

**while**  $x \neq \text{root}[T]$  and  $\text{color}[p[x]] = \text{RED}$

**do if**  $p[x] = \text{left}[p[p[x]]]$

**then**  $y \leftarrow \text{right}[p[p[x]]]$      ▷  $y = \text{aunt/uncle of } x$

**if**  $\text{color}[y] = \text{RED}$

**then** ⟨Case 1⟩

**else if**  $x = \text{right}[p[x]]$

**then** ⟨Case 2⟩    ▷ Case 2 falls into Case 3

          ⟨Case 3⟩

**else** ⟨“then” clause with “left” and “right” swapped⟩

**color**[ $\text{root}[T]$ ]  $\leftarrow \text{BLACK}$

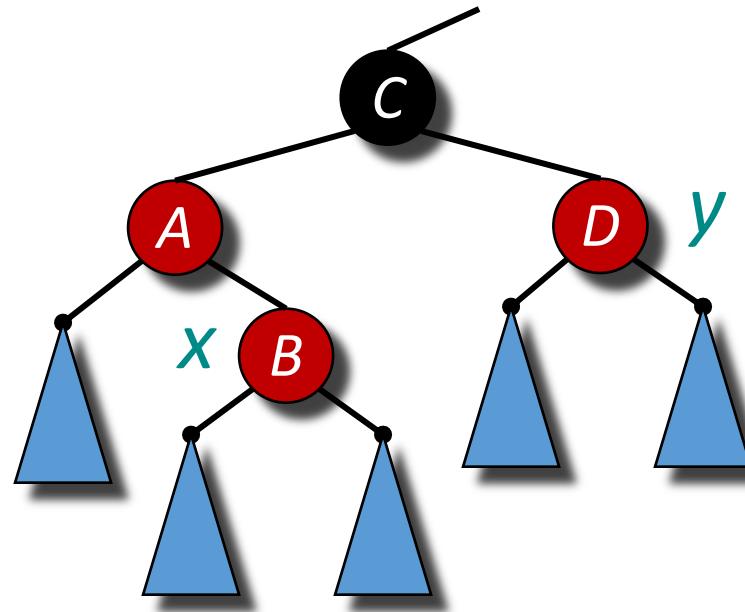
# Graphical notation

---

Let  denote a subtree with a black root.

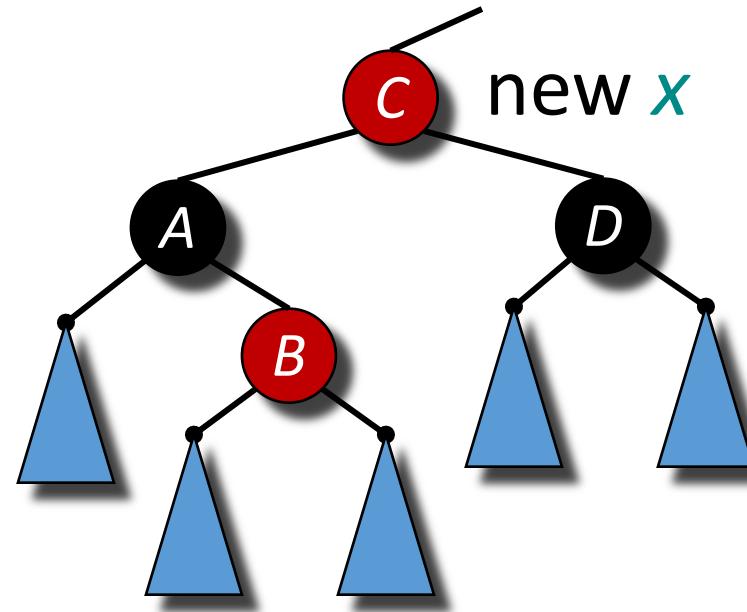
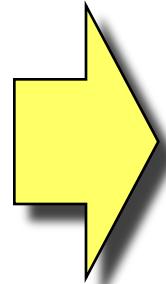
All 's have the same black-height.

# Case 1



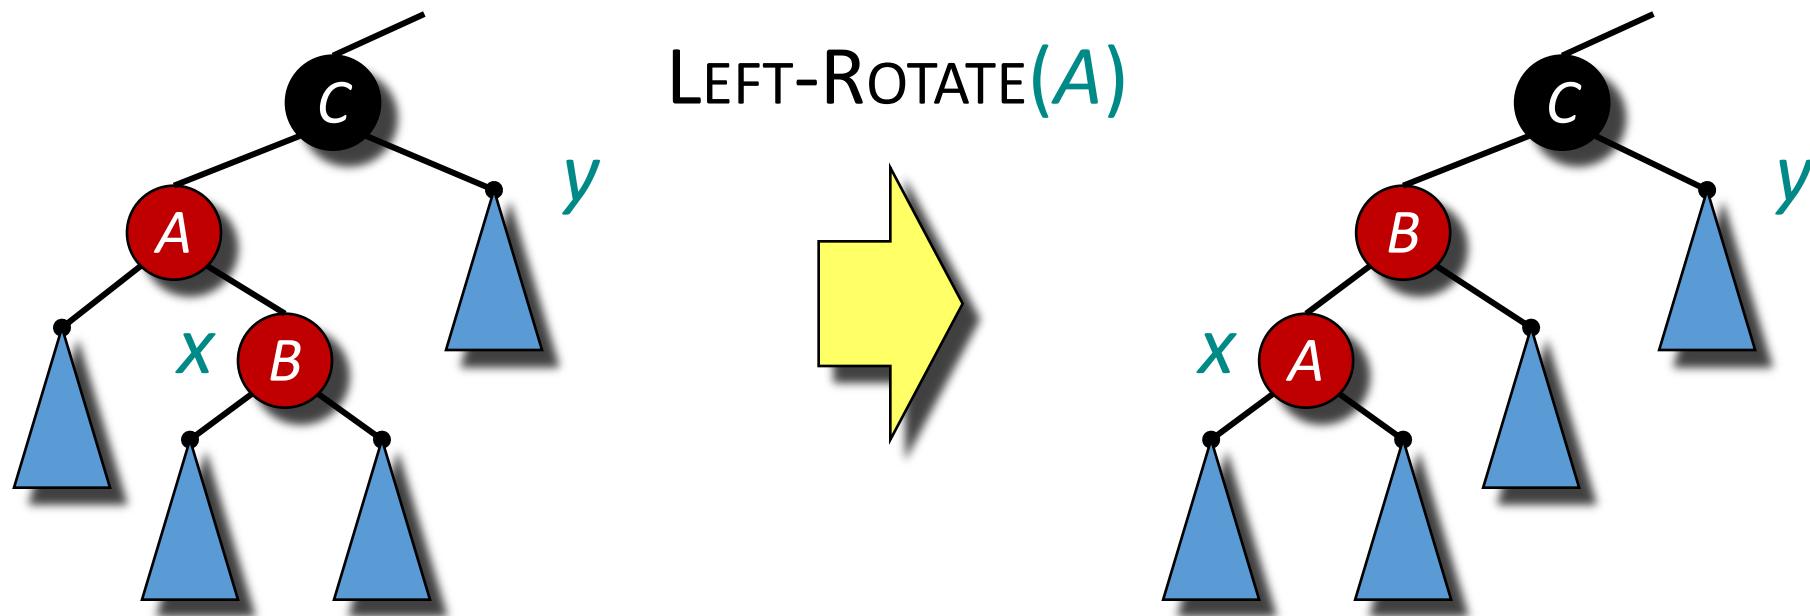
(Or, children of  $A$  are swapped.)

Recolor



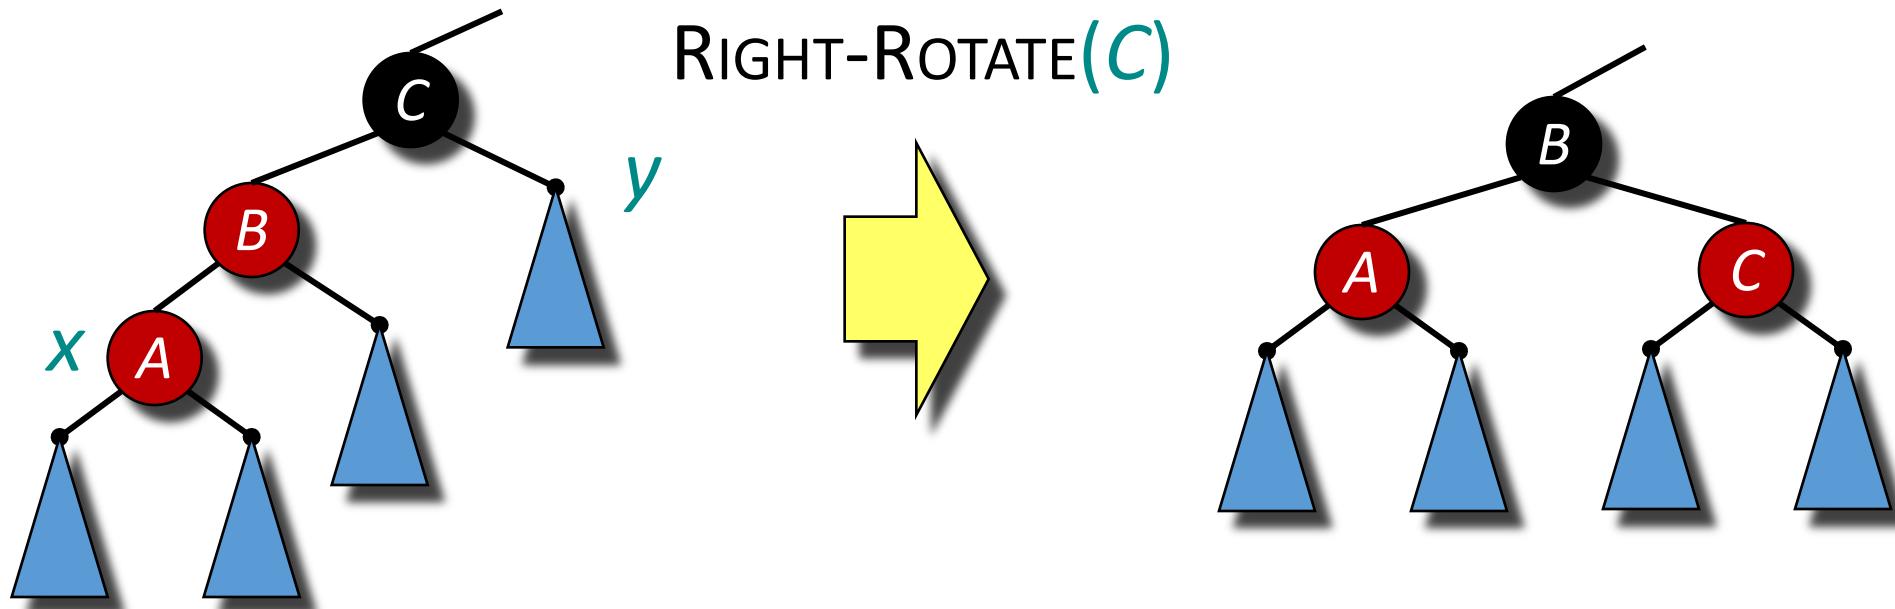
Push  $C$ 's black onto  $A$  and  $D$ , and recurse, since  $C$ 's parent may be red.

# Case 2



Transform to Case 3.

# Case 3



Done! No more violations of RB property 3 are possible.

# Analysis

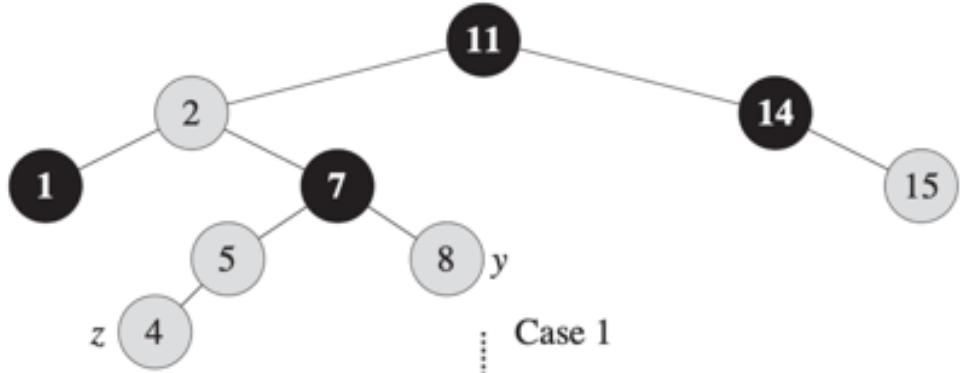
---

- Go up the tree performing Case 1, which only recolors nodes.
- If Case 2 or Case 3 occurs, perform 1 or 2 rotations, and terminate.

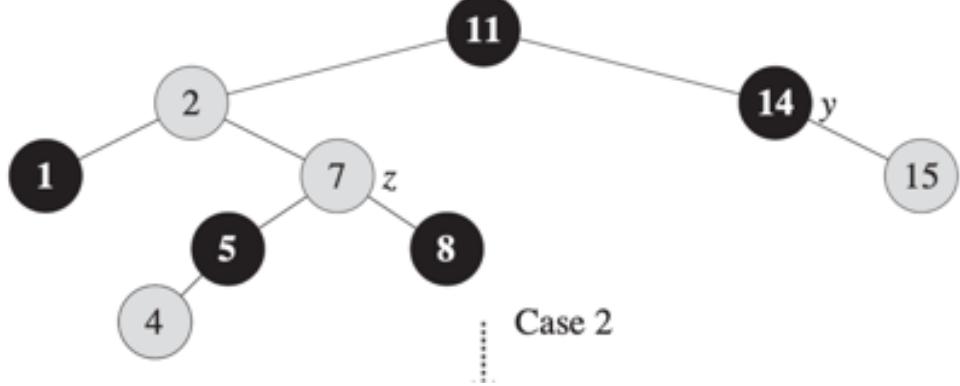
**Running time:**  $O(\lg n)$  with  $O(1)$  rotations.

RB-DELETE — same asymptotic running time and number of rotations as RB-INSERT (see textbook).

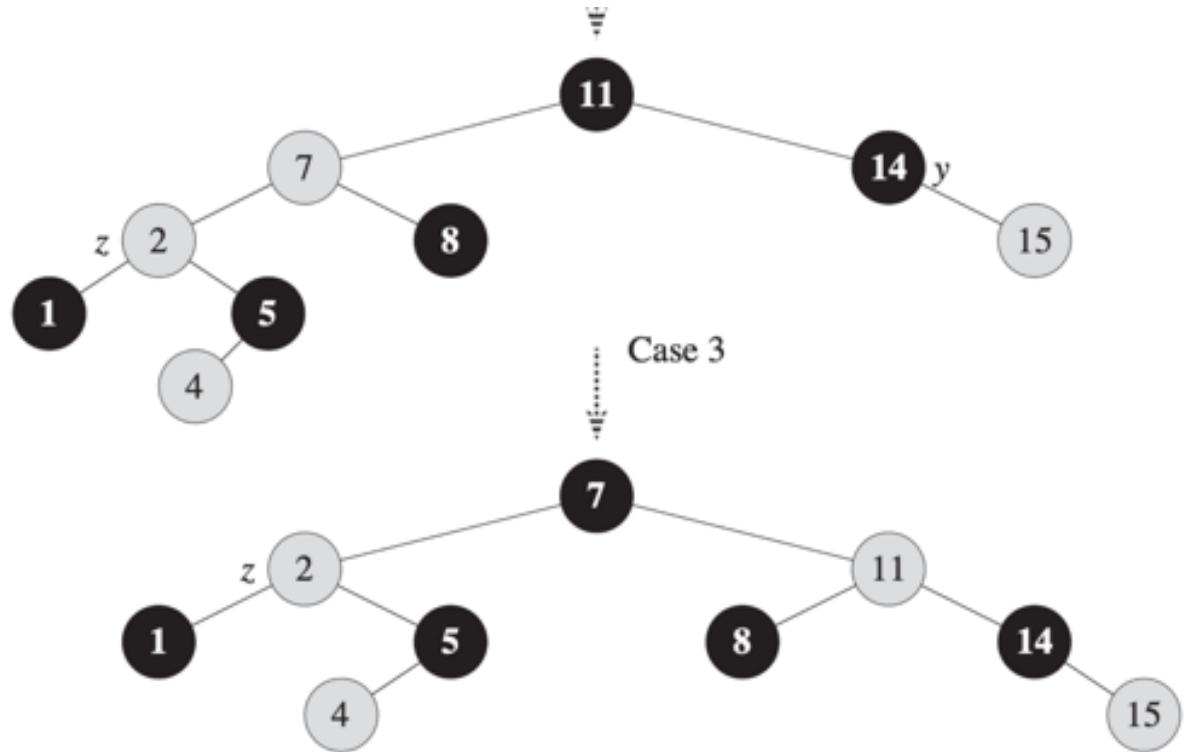
# Insertion Example



Case 1



Case 2



Case 3

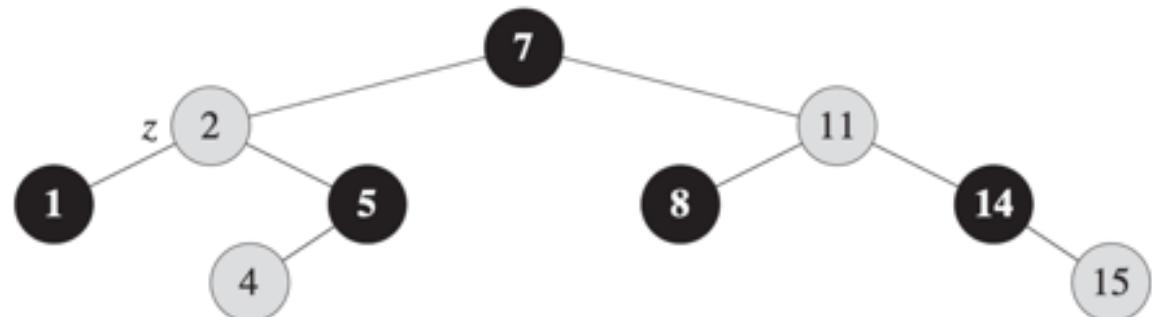
RB-DELETE( $T, z$ )

```
1   $y = z$ 
2   $y\text{-original-color} = y.\text{color}$ 
3  if  $z.\text{left} == T.\text{nil}$ 
4     $x = z.\text{right}$ 
5    RB-TRANSPLANT( $T, z, z.\text{right}$ )
6  elseif  $z.\text{right} == T.\text{nil}$ 
7     $x = z.\text{left}$ 
8    RB-TRANSPLANT( $T, z, z.\text{left}$ )
9  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10    $y\text{-original-color} = y.\text{color}$ 
11    $x = y.\text{right}$ 
12   if  $y.p == z$ 
13      $x.p = y$ 
14   else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15      $y.\text{right} = z.\text{right}$ 
16      $y.\text{right}.p = y$ 
17   RB-TRANSPLANT( $T, z, y$ )
18    $y.\text{left} = z.\text{left}$ 
19    $y.\text{left}.p = y$ 
20    $y.\text{color} = z.\text{color}$ 
21 if  $y\text{-original-color} == \text{BLACK}$ 
22   RB-DELETE-FIXUP( $T, x$ )
```

# Deletion

RB-TRANSPLANT( $T, u, v$ )

```
1  if  $u.p == T.\text{nil}$ 
2     $T.\text{root} = v$ 
3  elseif  $u == u.p.\text{left}$ 
4     $u.p.\text{left} = v$ 
5  else  $u.p.\text{right} = v$ 
6   $v.p = u.p$ 
```

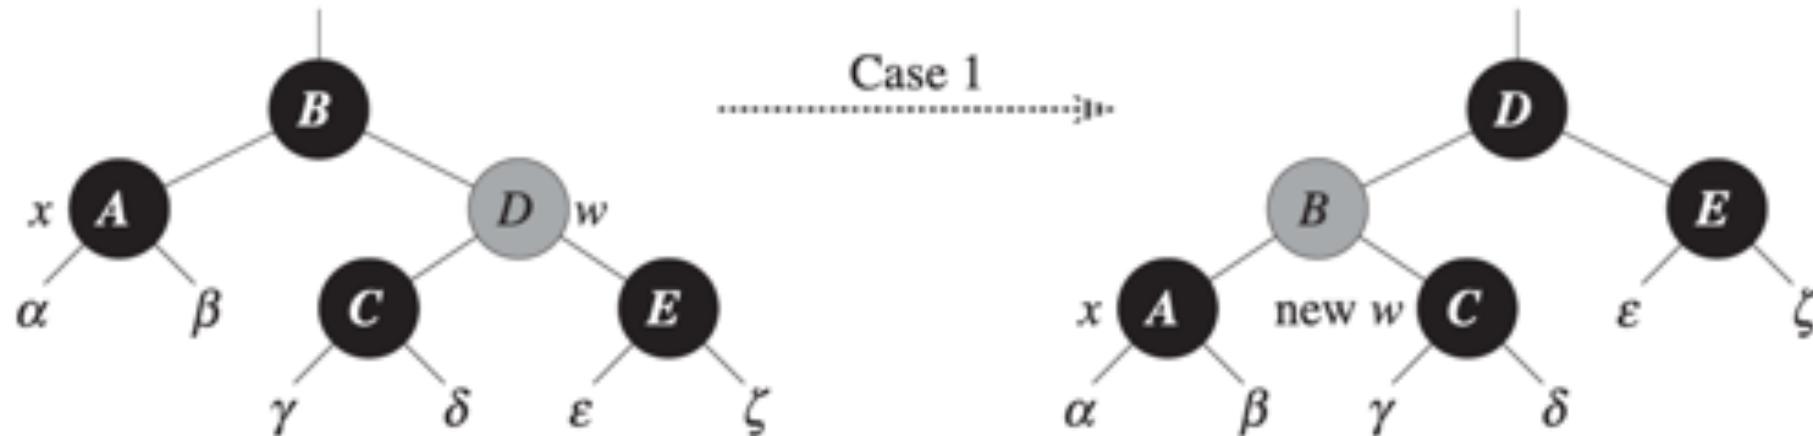


RB-DELETE-FIXUP( $T, x$ )

```
1  while  $x \neq T.root$  and  $x.color == \text{BLACK}$ 
2      if  $x == x.p.left$ 
3           $w = x.p.right$ 
4          if  $w.color == \text{RED}$ 
5               $w.color = \text{BLACK}$                                 // case 1
6               $x.p.color = \text{RED}$                             // case 1
7              LEFT-ROTATE( $T, x.p$ )                      // case 1
8               $w = x.p.right$                             // case 1
9          if  $w.left.color == \text{BLACK}$  and  $w.right.color == \text{BLACK}$ 
10          $w.color = \text{RED}$                                 // case 2
11          $x = x.p$                                     // case 2
12     else if  $w.right.color == \text{BLACK}$ 
13          $w.left.color = \text{BLACK}$                           // case 3
14          $w.color = \text{RED}$                                 // case 3
15         RIGHT-ROTATE( $T, w$ )                        // case 3
16          $w = x.p.right$                             // case 3
17          $w.color = x.p.color$                           // case 4
18          $x.p.color = \text{BLACK}$                         // case 4
19          $w.right.color = \text{BLACK}$                       // case 4
20         LEFT-ROTATE( $T, x.p$ )                      // case 4
21          $x = T.root$                                 // case 4
22     else (same as then clause with "right" and "left" exchanged)
23      $x.color = \text{BLACK}$ 
```

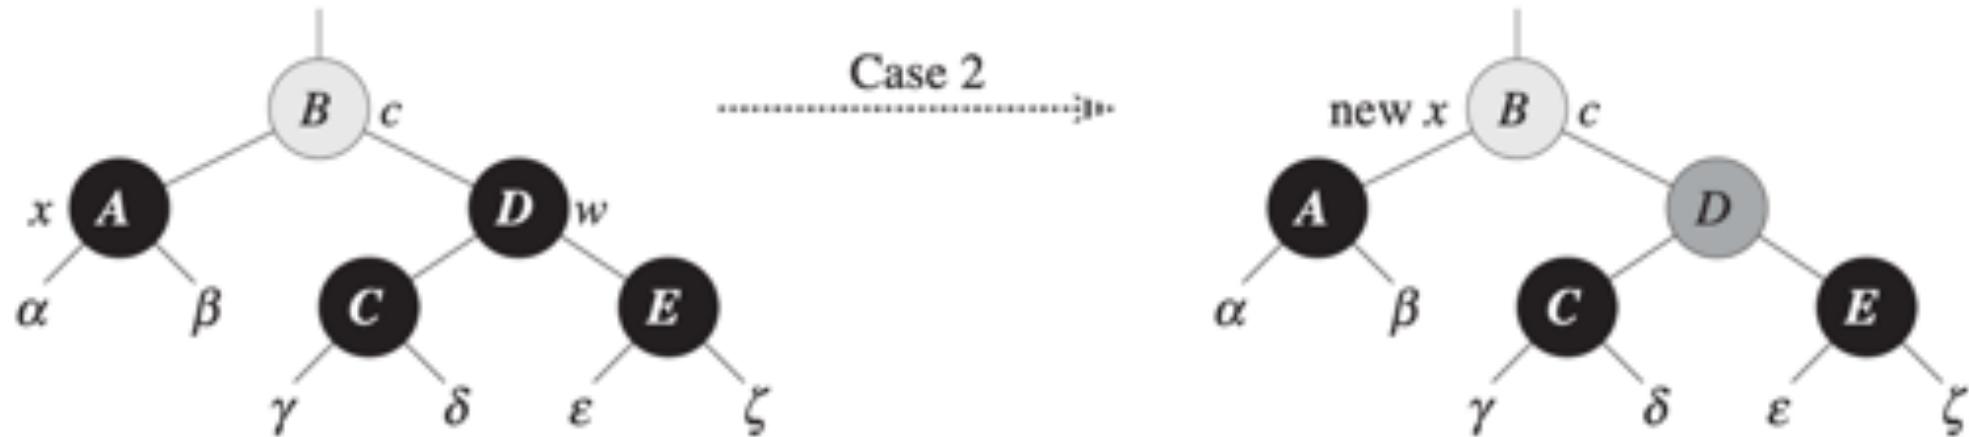
# Deletion

- Case 1:  $x$ 's sibling  $w$  is red



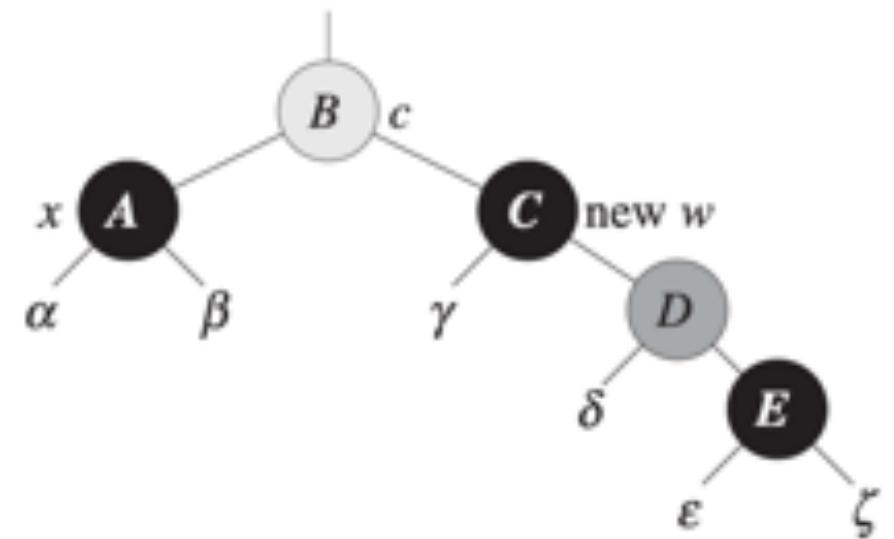
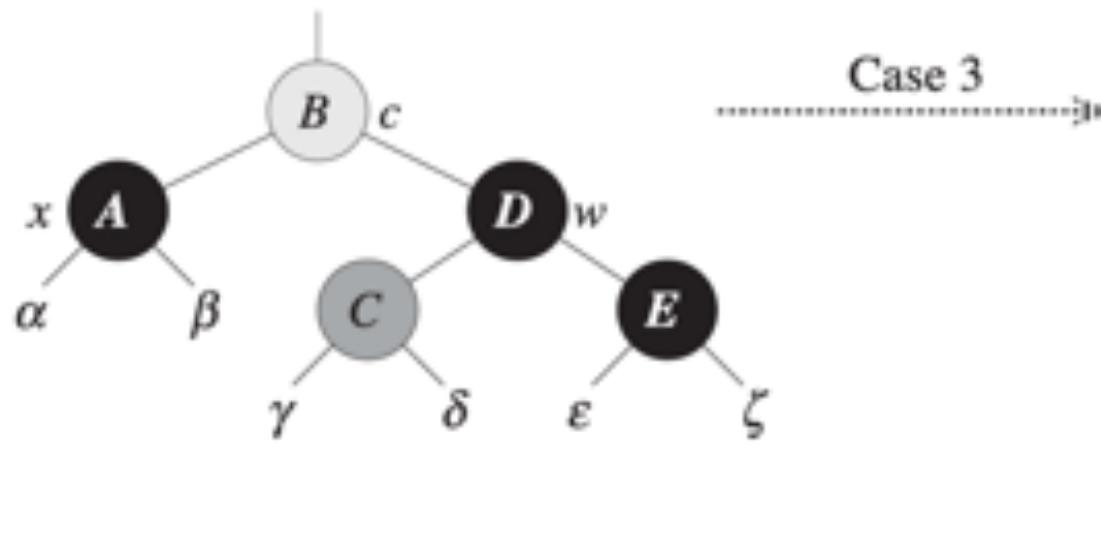
# Deletion

- Case 2:  $x$ 's sibling  $w$  is black, and both of  $w$ 's children are black



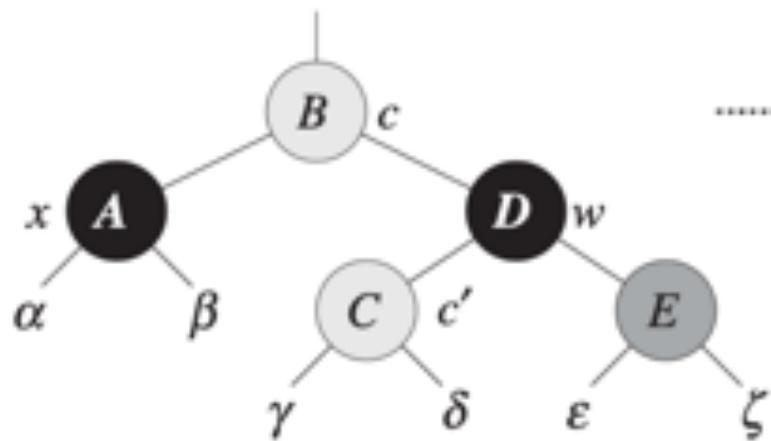
# Deletion

- Case 3:  $x$ 's sibling  $w$  is black,  $w$ 's left child is red, and  $w$ 's right child is black

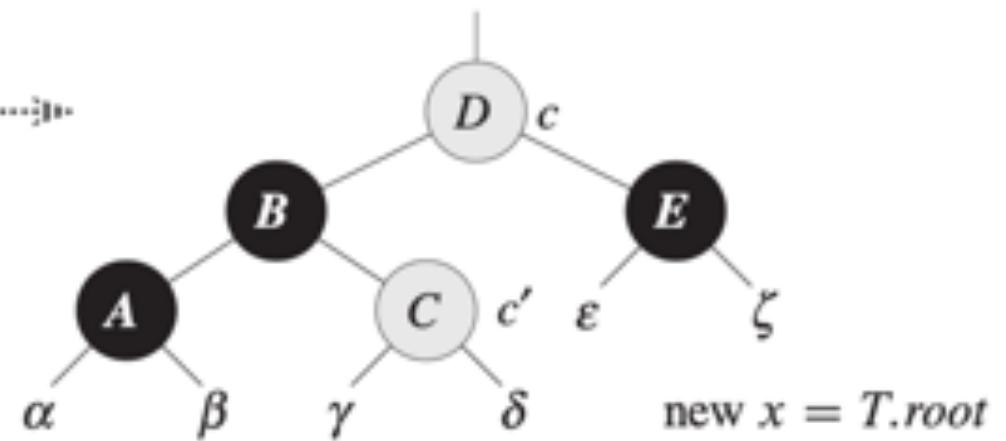


# Deletion

- Case 4:  $x$ 's sibling  $w$  is black, and  $w$ 's right child is red



Case 4



# Thank You